

CMSC740

Advanced Computer Graphics

Fall 2025

Matthias Zwicker

Today

- Course overview
- Organization
- Introduction to rendering

Computer graphics

“Algorithms to create and manipulate images”



Challenges

- Modeling, content creation requires intense manual work
 - 3D modeling & animation, image & video editing & processing, etc.
- Realistic rendering computationally expensive
 - Real-time rendering in AR/VR limited to lower quality
- Currently
 - Feature length movies cost hundreds of millions
 - Limited applications of AR/VR



Research vision

- AI-based digital media creation
 - User provides simple, high-level input
 - AI system fills in details
- Benefits
 - Faster, cheaper media creation
 - Democratize visual story-telling: broaden access to low-budget producers, like small studios and individuals
 - Support and enable new media formats (AR, VR, ...)

Core areas

- Rendering: computing images given 3D scene descriptions (geometry, light sources, material properties)
- Modeling: creating 3D scene descriptions (manual, algorithmic, reconstruction from measurements/images)
- Animation: making things move (physics simulation)

This course

- First half: rendering (photo-realistic)
- Second half: modeling
- Focus on recent neural-network/AI based techniques

Other graphics courses at UMD

- CMSC425, Game Programming
- CMSC427, Computer Graphics (real-time rendering)
- CMSC498F/838C: Advances in XR (AR/VR)
- CMSC828X, Physically-based Modeling, Simulation, and Animation
- CMSC848B, Computational Imaging

Today

- Course overview
- Organization
- Introduction to rendering

Organization

- Class material
 - Course page on UMD Canvas, log in with directory ID
 - Piazza for course questions
 - Virtual Study Assistant AI chatbot, configured with course materials
- Teaching assistants
 - Haiyang Yin, Yun He, CS PhD students

Organization

- Lectures
Tuesdays and Thursdays, 2pm-3:15pm,
IRB 2107
- Office hours
 - TAs: Mondays, 3-5pm, location TBA
 - Instructor: By appointment (send me email)

Grading

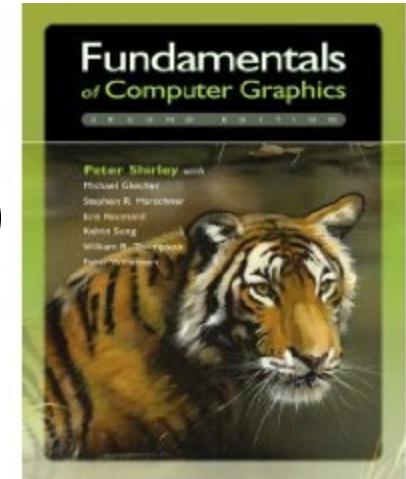
- Programming assignments
 - 40% of final grade
 - Individual projects
 - During first half of course
- Final project
 - 40% of final grade
 - Groups of two
 - During second half of course
- Final exam
 - 15% of final grade
- Class participation
 - 5% of final grade

Policies

- Standard UMD academic integrity policies
<https://www.studentconduct.umd.edu/academic-dishonesty>
- Programming assignments
 - Each student submits own solution (no copying, no direct collaboration on code, no AI generated code)
 - Penalty for late submission: 50% of original score
 - Public posting (github etc.) of solutions prohibited

What you should know

- Programming experience
in Java, Python, C++, data structures
(CMSC420)
- Linear algebra (MATH240)
 - Vectors, matrices, systems of linear equations, coordinate transformations
- (Basic background in 3D graphics)
 - E.g., Shirley, Fundamentals of Computer Graphics



What you should know

- Mathematics
 - Vectors in 3D (dot product, vector product)
 - Matrix algebra (addition, multiplication, inverse, determinant, transpose, orthogonal matrices)
 - Coordinate systems (homogeneous coordinates, change of coordinates)
 - Transformation matrices (scaling, rotation, translation, shearing)
 - Plane equations
 - Systems of linear equations (matrix representation, elementary solution procedures)
- Graphics
 - Color
 - 3D scene representation (triangle meshes, scene graphs)

Learning objectives

- Understand and apply mathematical concepts and advanced algorithms in computer graphics
 - Rendering (rendering equation, Monte Carlo integration, path tracing, 3D geometry processing, numerical techniques, etc.)
 - Deep learning/AI in computer graphics
- Sharpen your programming skills (python, neural network techniques)

Python programming experience

- A. Never used
- B. Beginner (worked with it less than 16 hours/2 days)
- C. Intermediate (worked with it 16-80 hours/2-10 days)
- D. Advanced

I have trained a neural network

- A. Yes
- B. No

I can explain matrix multiplication

- A. Strongly Agree
- B. Agree
- C. Somewhat Agree
- D. Neutral
- E. Somewhat Disagree
- F. Disagree
- G. Strongly Disagree

I can explain the concept of the basis of a vector space

- A. Strongly Agree
- B. Agree
- C. Somewhat Agree
- D. Neutral
- E. Somewhat Disagree
- F. Disagree
- G. Strongly Disagree

Today

- Course overview
- Organization
- Introduction to rendering

First half of course

- 7-week crash course “theory and practice of photo-realistic rendering for computer graphics”

Rendering algorithms

- Problem statement
 - Given computer representation of 3D scene, generate image of the scene that would be captured by a virtual camera at a given location
- Applications
 - Movies, games, design, architecture, virtual/augmented reality, visualization, tele-collaboration, etc.



Challenges

- „Photo-realism“
 - Realistic simulation of light transport
- Efficiency
 - Thousands of CPU years for movie production

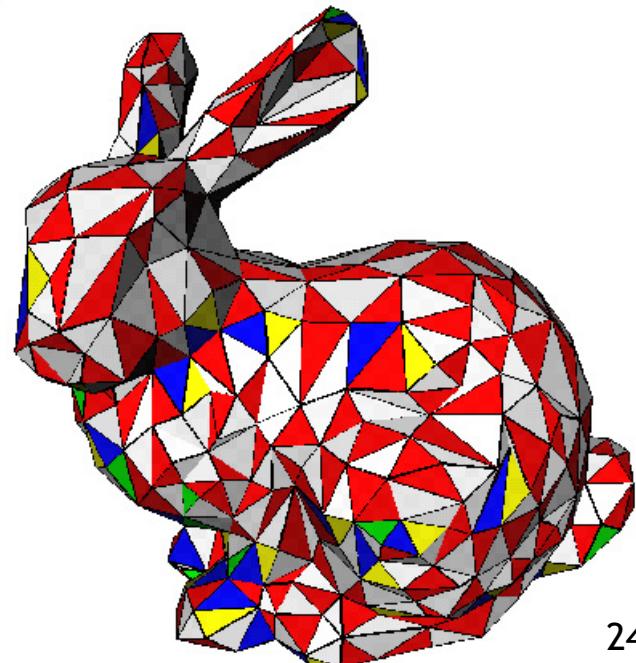


Challenges

- Real-time rendering
 - Games, VR, AR (90 frames per second (fps), 7-15ms motion-to-photon latency)
 - Highly limited computation budget for light transport simulation

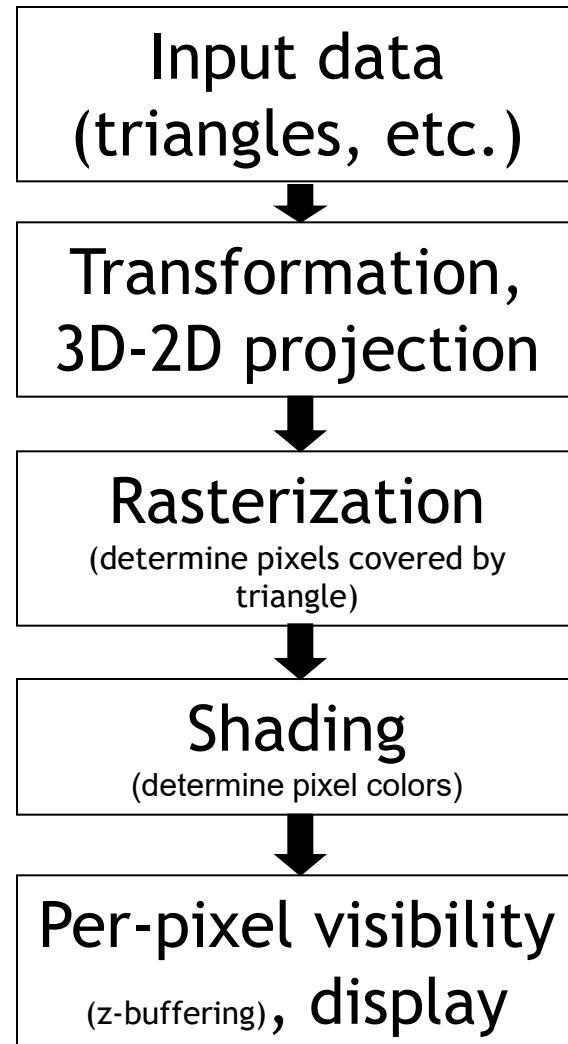
Rendering strategies

- Two fundamentally different strategies for image generation (rendering)
 1. Rendering pipeline (Z-buffering, rasterization)
 2. Ray tracing
- In both cases, most popular object representation are triangle meshes
 - Data structure: array of x,y,z vertex positions; index array indicating triplets of vertices forming triangles



Rendering pipeline

https://en.wikipedia.org/wiki/Graphics_pipeline

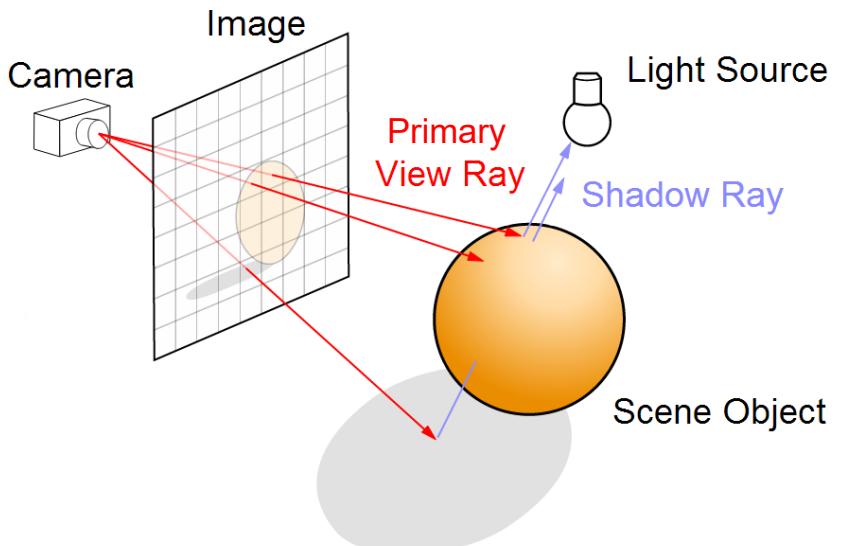
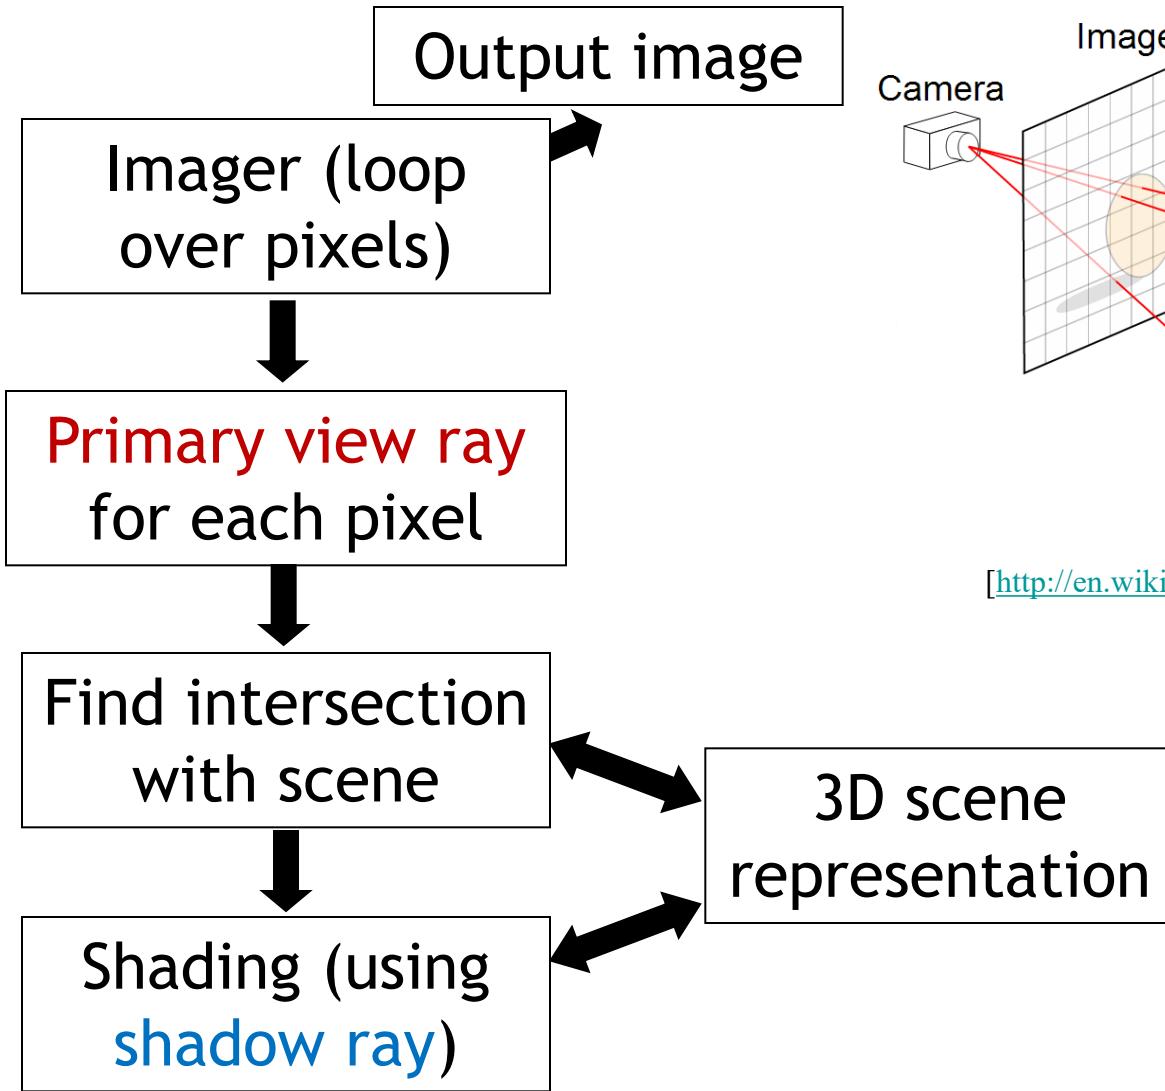


Object order algorithm,
paint (transform,
project, rasterize) one
primitive (usually
triangle) after another,
visibility via z-buffer
(store/compare per-
pixel depth)

Z-buffering

- Implemented in graphics hardware (ATI, NVidia GPUs)
- Standardized APIs (OpenGL, Direct3D, Vulcan)
- Interactive rendering, AR/VR, games
- Limited photo-realism

Ray tracing



[[http://en.wikipedia.org/wiki/Ray_tracing_\(graphics\)](http://en.wikipedia.org/wiki/Ray_tracing_(graphics))]

Image order algorithm,
paint one pixel after another

Ray tracing vs. rasterization

- Rasterization
 - Desirable memory access pattern („streaming“, data locality, avoids random scene access)
 - Suitable for **real time rendering** (OpenGL, DirectX)
 - Full global illumination not possible with purely object order algorithm
- Ray tracing
 - Undesirable memory access pattern (random scene access)
 - Requires sophisticated data structures for fast scene access
 - Full global illumination possible
 - Most popular for **photo-realistic rendering**
 - Recently with hardware support

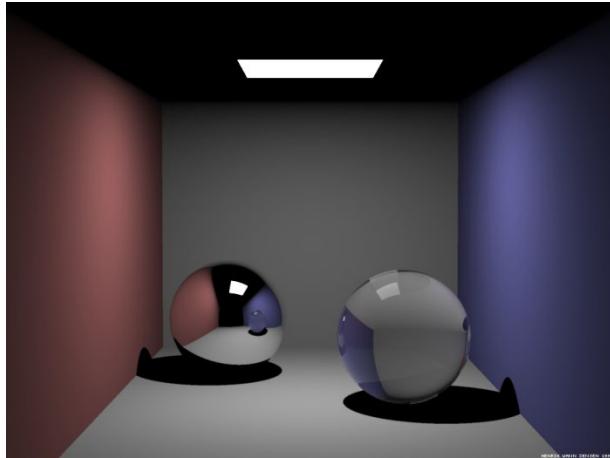
<https://www.nvidia.com/content/dam/en-zz/Solutions/design-visualization/technologies/turing-architecture/NVIDIA-Turing-Architecture-Whitepaper.pdf>

Goal of first half of class

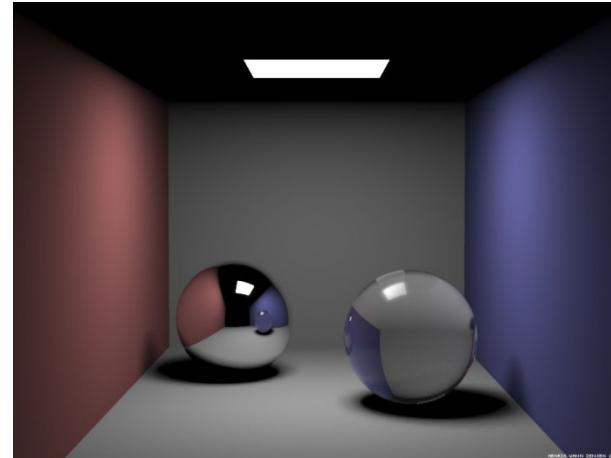
- Realistic rendering: learn theory and practice of light transport simulation for computer graphics
 - “Rendering algorithms” based on ray tracing
 - Spatial data structures for ray tracing, physics background, rendering equation, (quasi) Monte Carlo integration, (multiple) importance sampling, Markov Chain Monte Carlo sampling, leveraging neural networks, etc.

Realistic rendering

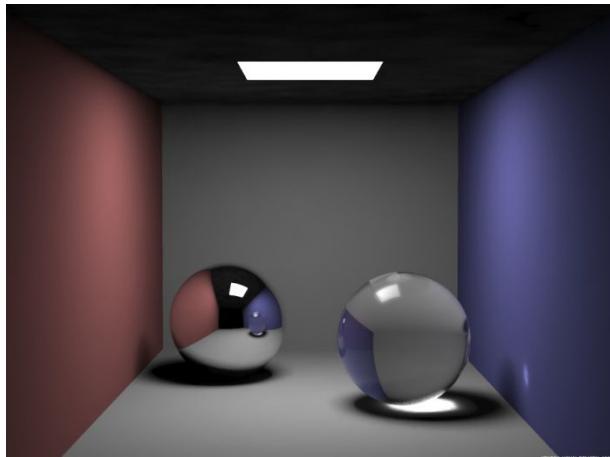
Ray
tracing



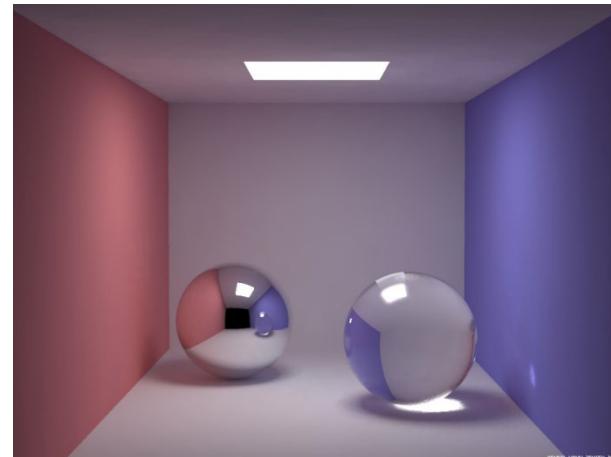
Soft
shadows



Caustics

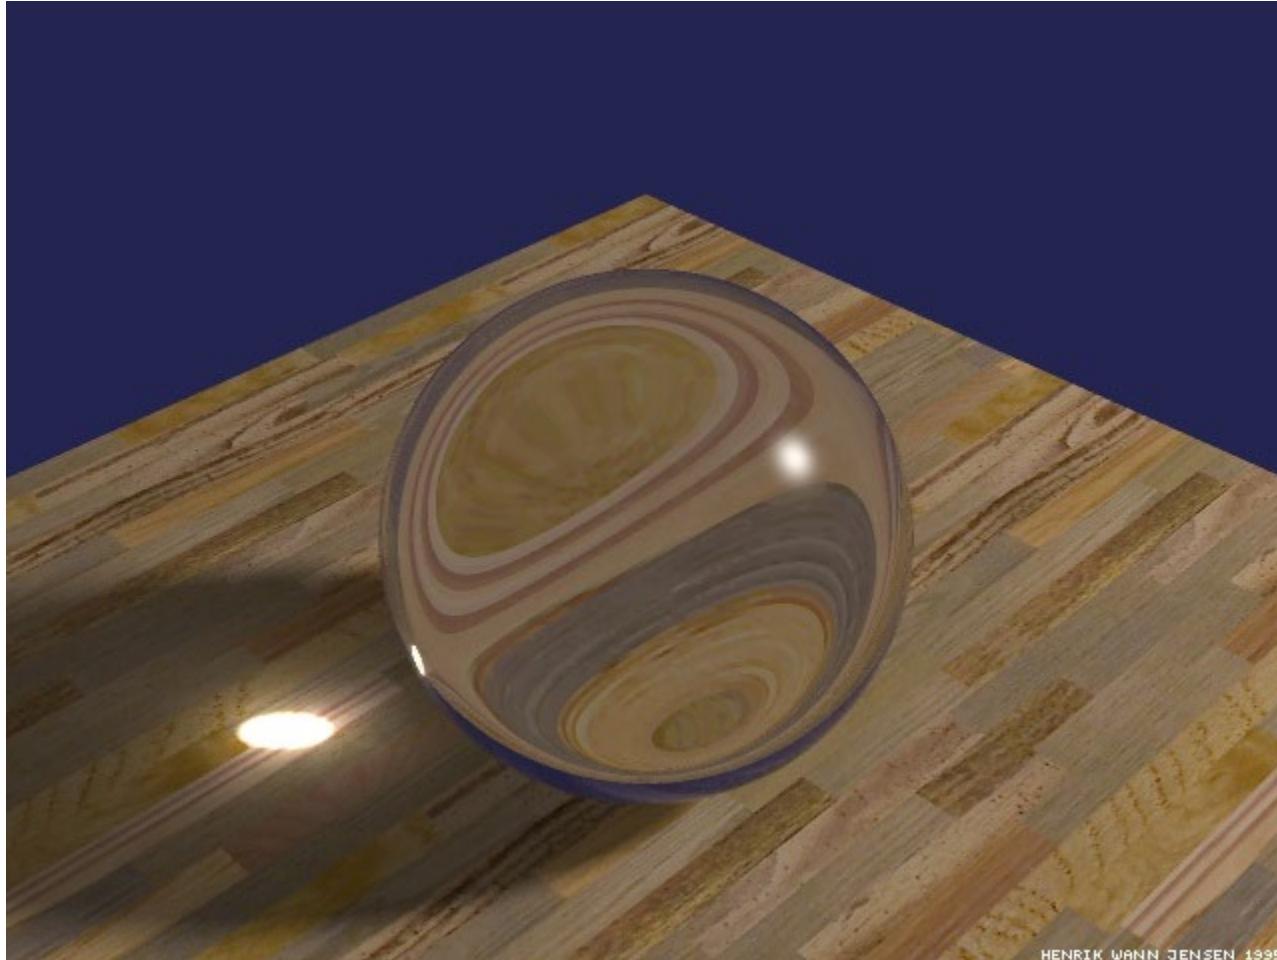


Global
illum.



[Wann Jensen]

Caustics



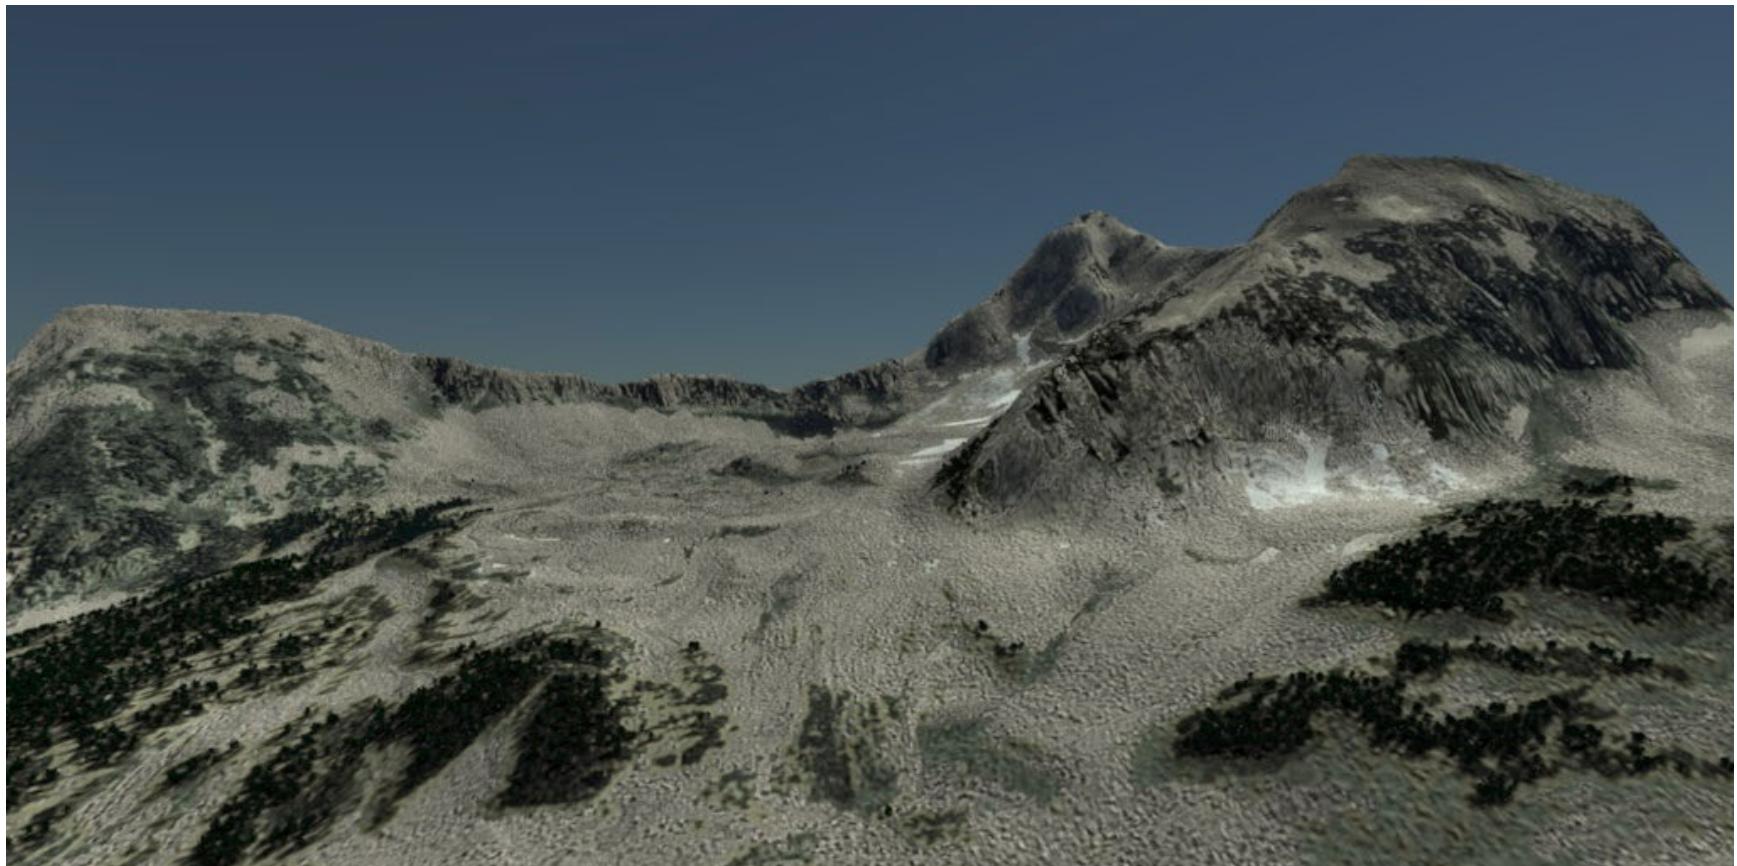
[Wann Jensen]
<http://graphics.ucsd.edu/~henrik/papers/book/>

Indirect illumination



[Wann Jensen]

Atmospheric effects



[Wann Jensen]
<http://graphics.ucsd.edu/~henrik/papers/nightsky/>

Participating media



[Novak et al.]

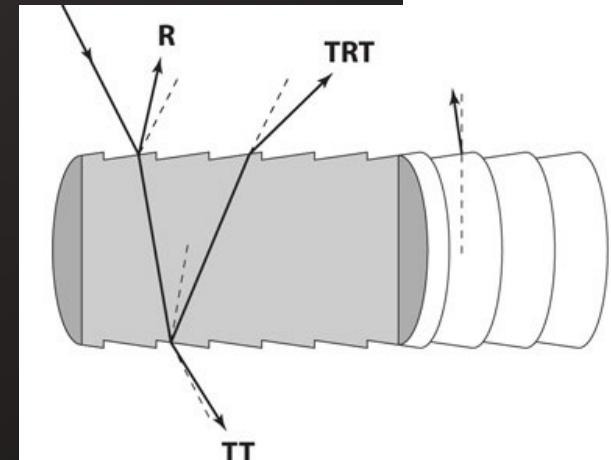
<http://zurich.disneyresearch.com/~wjarosz/publications/novak12vrls.html>

Complex materials



Maxwell renderer

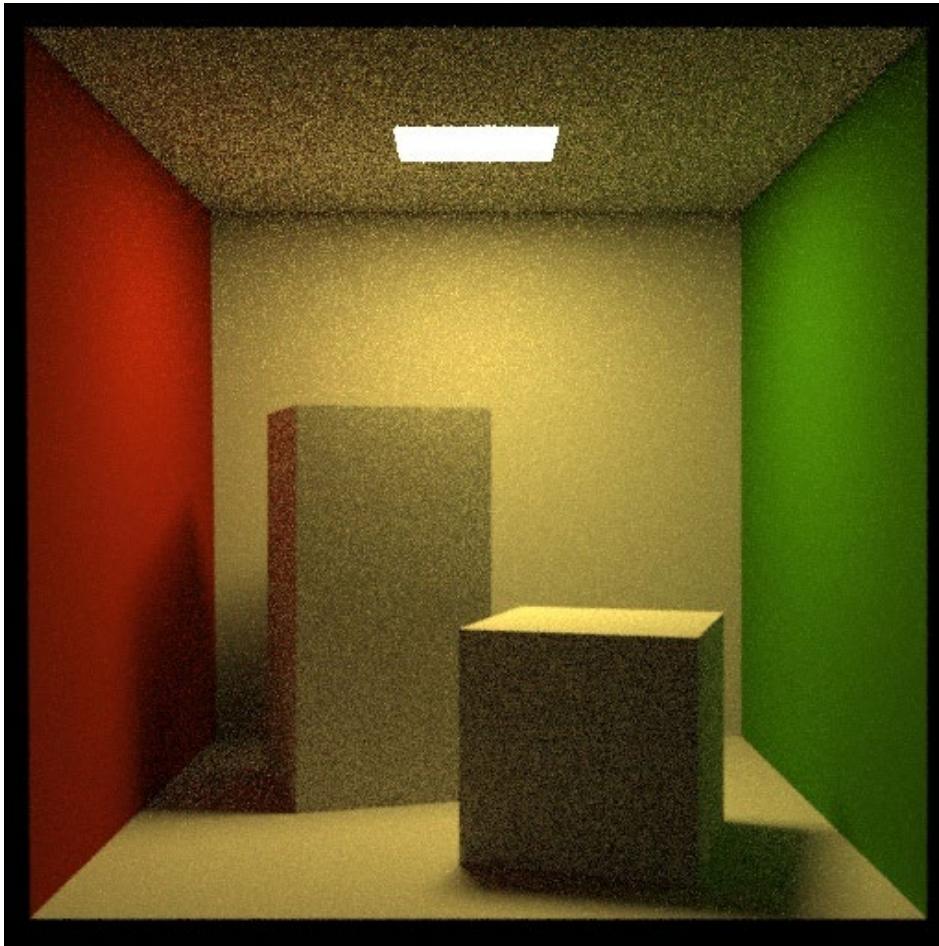
<http://www.maxwellrender.com/index.php/features>



Reflection and transmission
in a hair fiber

[<http://www.cs.cornell.edu/~srm/publications/SG03-hair-abstract.html>]

Rendering artifacts, noise



Cornell box

http://en.wikipedia.org/wiki/Cornell_Box

Big challenge,
address with more
computation time,
and sophisticatd
algorithms

Recommended book

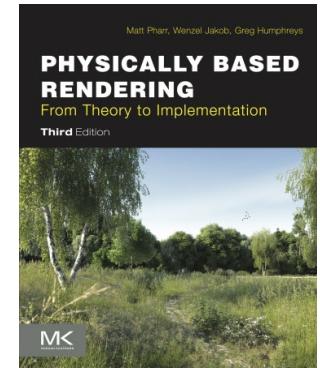
- „Physically based rendering“, by Pharr, Jakob, Humphreys

- „Bible“ of realistic rendering
 - Available freely online

<https://www.pbr-book.org/>

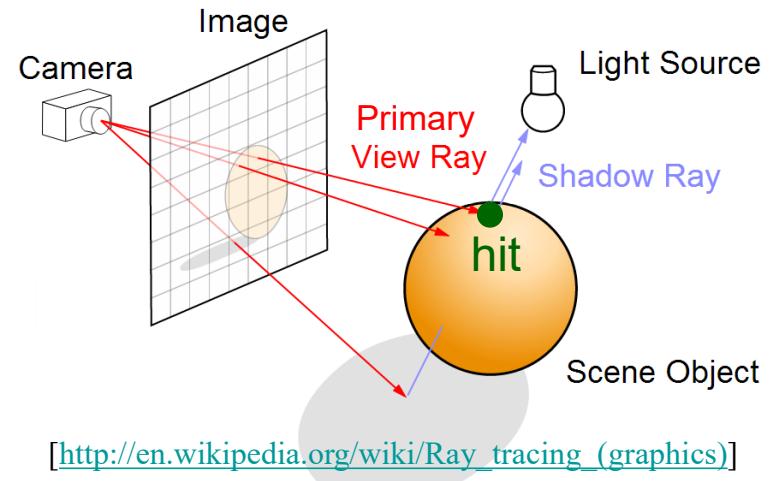
<http://pbrt.org/>

- Detailed, useful as reference
 - Includes many advanced concepts
 - With C++ code



Ray tracing pseudocode

```
rayTrace() {  
    construct scene representation  
  
    for each pixel  
        ray = computePrimaryViewRay( pixel )  
        hit = first intersection of ray with scene  
        color = shade( hit ) // using shadow ray  
        set pixel color  
}
```



[[http://en.wikipedia.org/wiki/Ray_tracing_\(graphics\)](http://en.wikipedia.org/wiki/Ray_tracing_(graphics))]

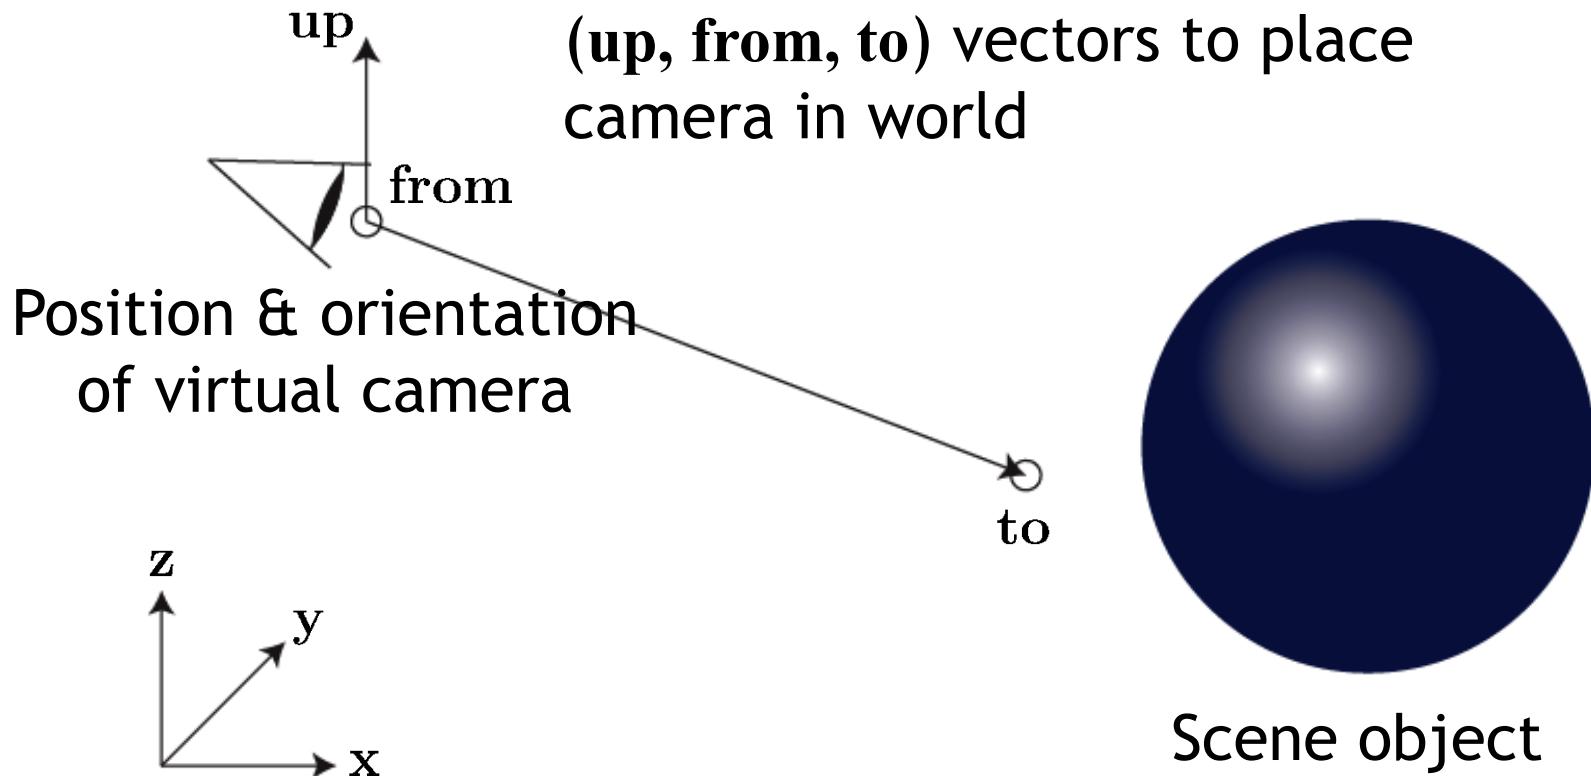
Computational complexity of ray tracing depends on

- A. Number of pixels in image
- B. Number of triangles in 3D triangle meshes
- C. Product of number of pixels and number of triangles

Computing primary rays

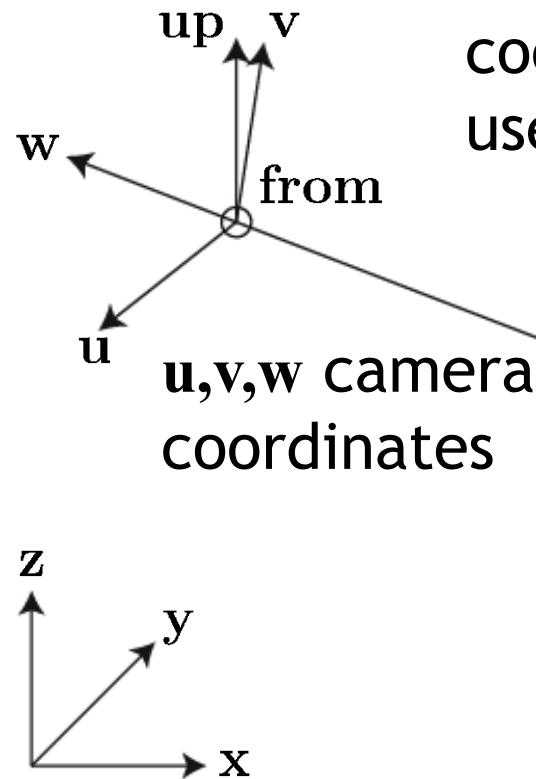
- Primary rays sometimes called “camera rays”, “view rays”
- Need to define **camera coordinate system**
 - Specifies position and orientation of virtual camera in 3D space
- Need to define **viewing frustum**
 - Pyramid shaped volume that contains 3D volume seen by camera
- Given an image pixel
 1. Determine ray in camera coordinates, **intrinsic camera matrix**
 2. Transform ray to world coordinates (sometimes also called “canonic” coordinates), **extrinsic camera matrix**

Camera coordinate system



World coordinates given by
x,y,z basis vectors

Camera coordinate system



Need to construct basis of camera coordinate system (**u,v,w**) based on user parameters (**up, from, to**)

x,y,z world coordinates

Right handed coordinate systems!

https://en.wikipedia.org/wiki/Right-hand_rule

Camera coordinate system

- Given from, to, up

$$w = \frac{\text{from} - \text{to}}{\|\text{from} - \text{to}\|}$$

$$u = \frac{\text{up} \times w}{\|\text{up} \times w\|}$$

$$v = w \times u$$

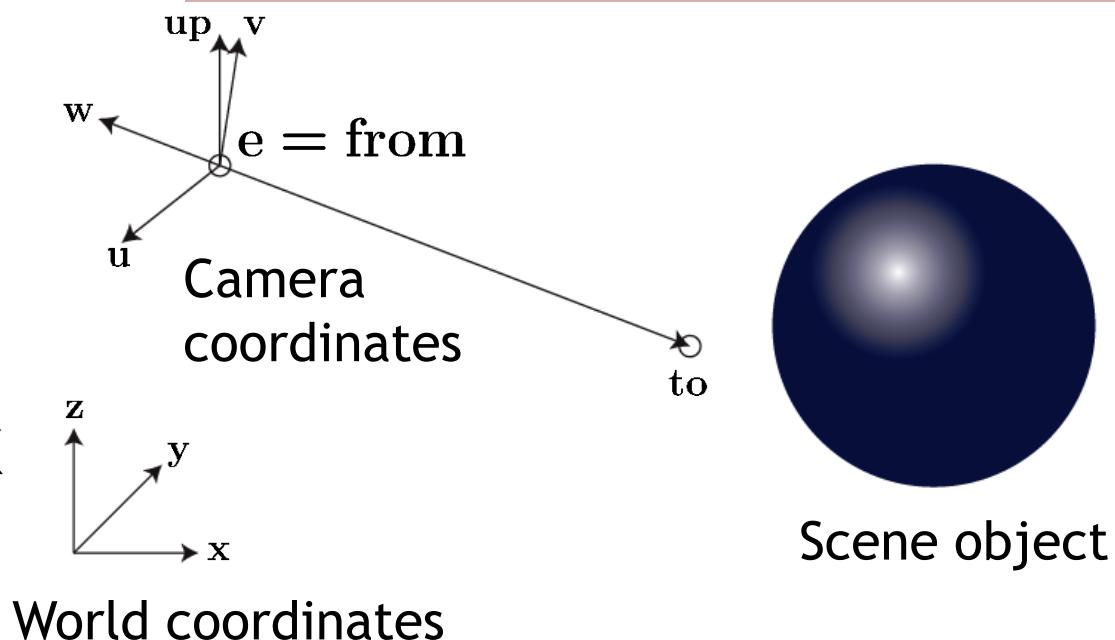
$$e = \text{from}$$

- Extrinsic 4x4 camera matrix

$$[u \ v \ w \ e]$$

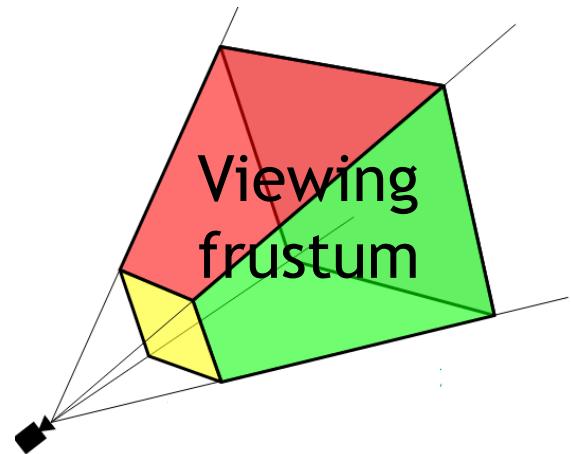
Use homogeneous coordinates to enable affine transformations (including rotation, translation) with matrix multiplications

https://en.wikipedia.org/wiki/Transformation_matrix#Affine_transformations

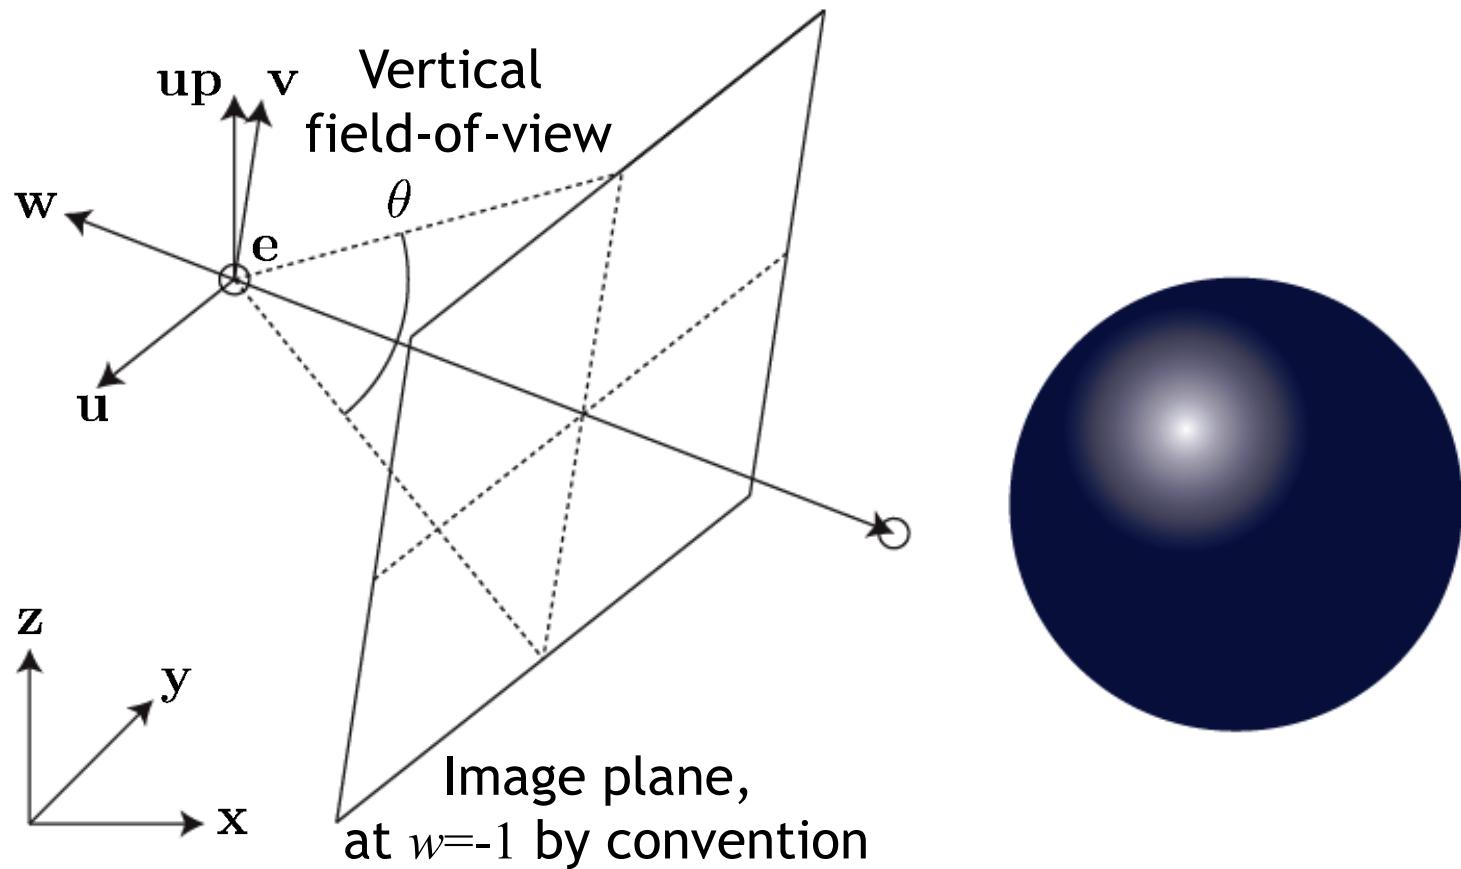


Viewing frustum

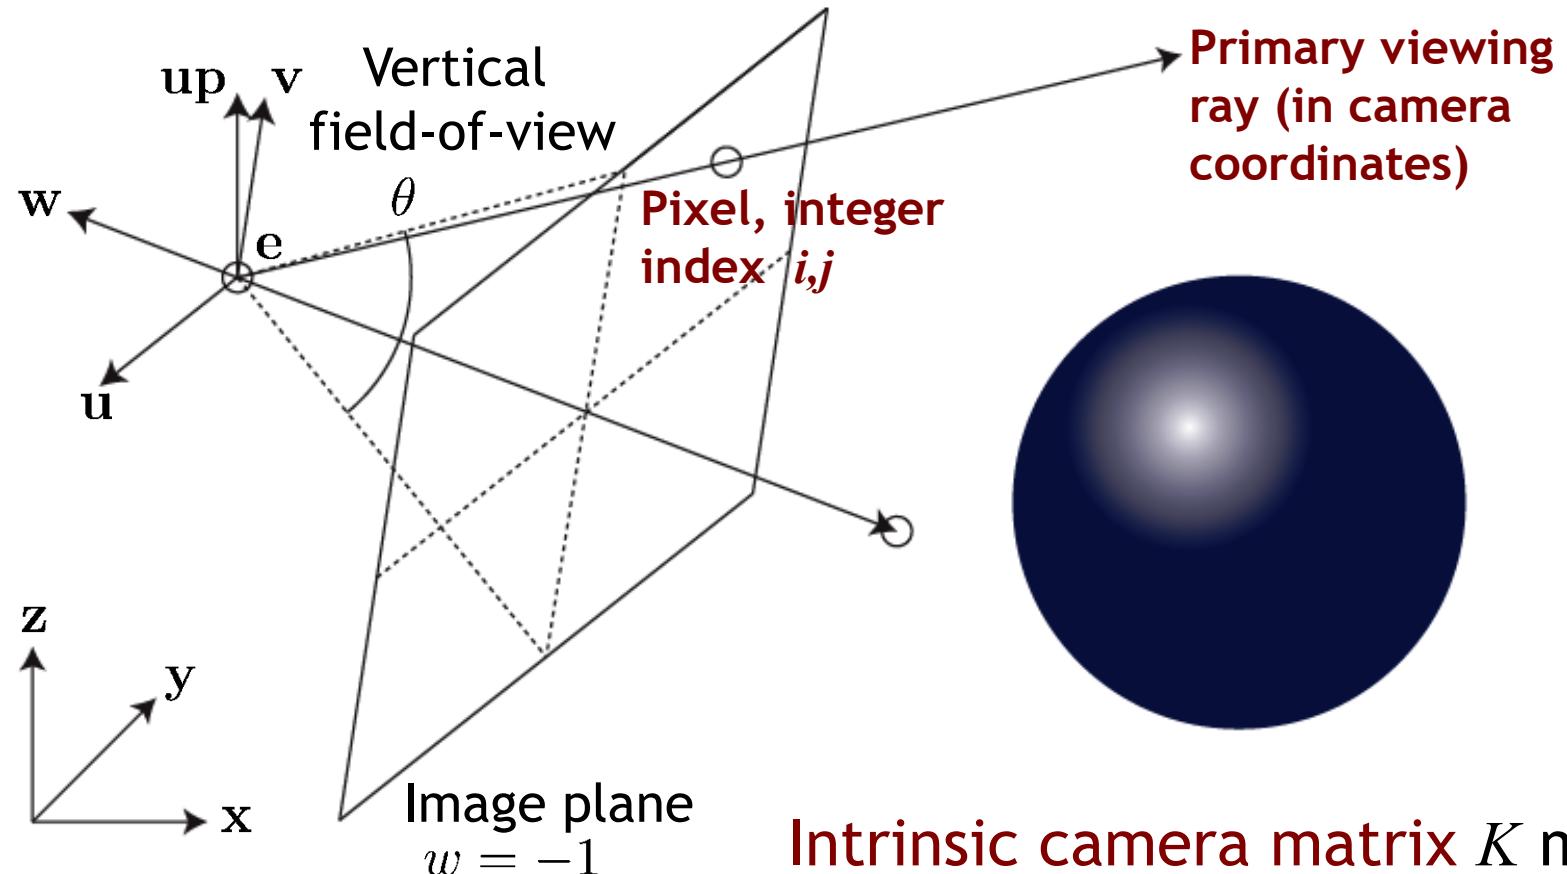
- 3D volume that is projected within image boundaries
 - “Pyramidal cone” that contains all camera rays
- Defined by
 - Vertical field-of-view θ
 - Aspect ratio $aspect = width/height$
 - By convention, image plane at $w=-1$



Viewing frustum

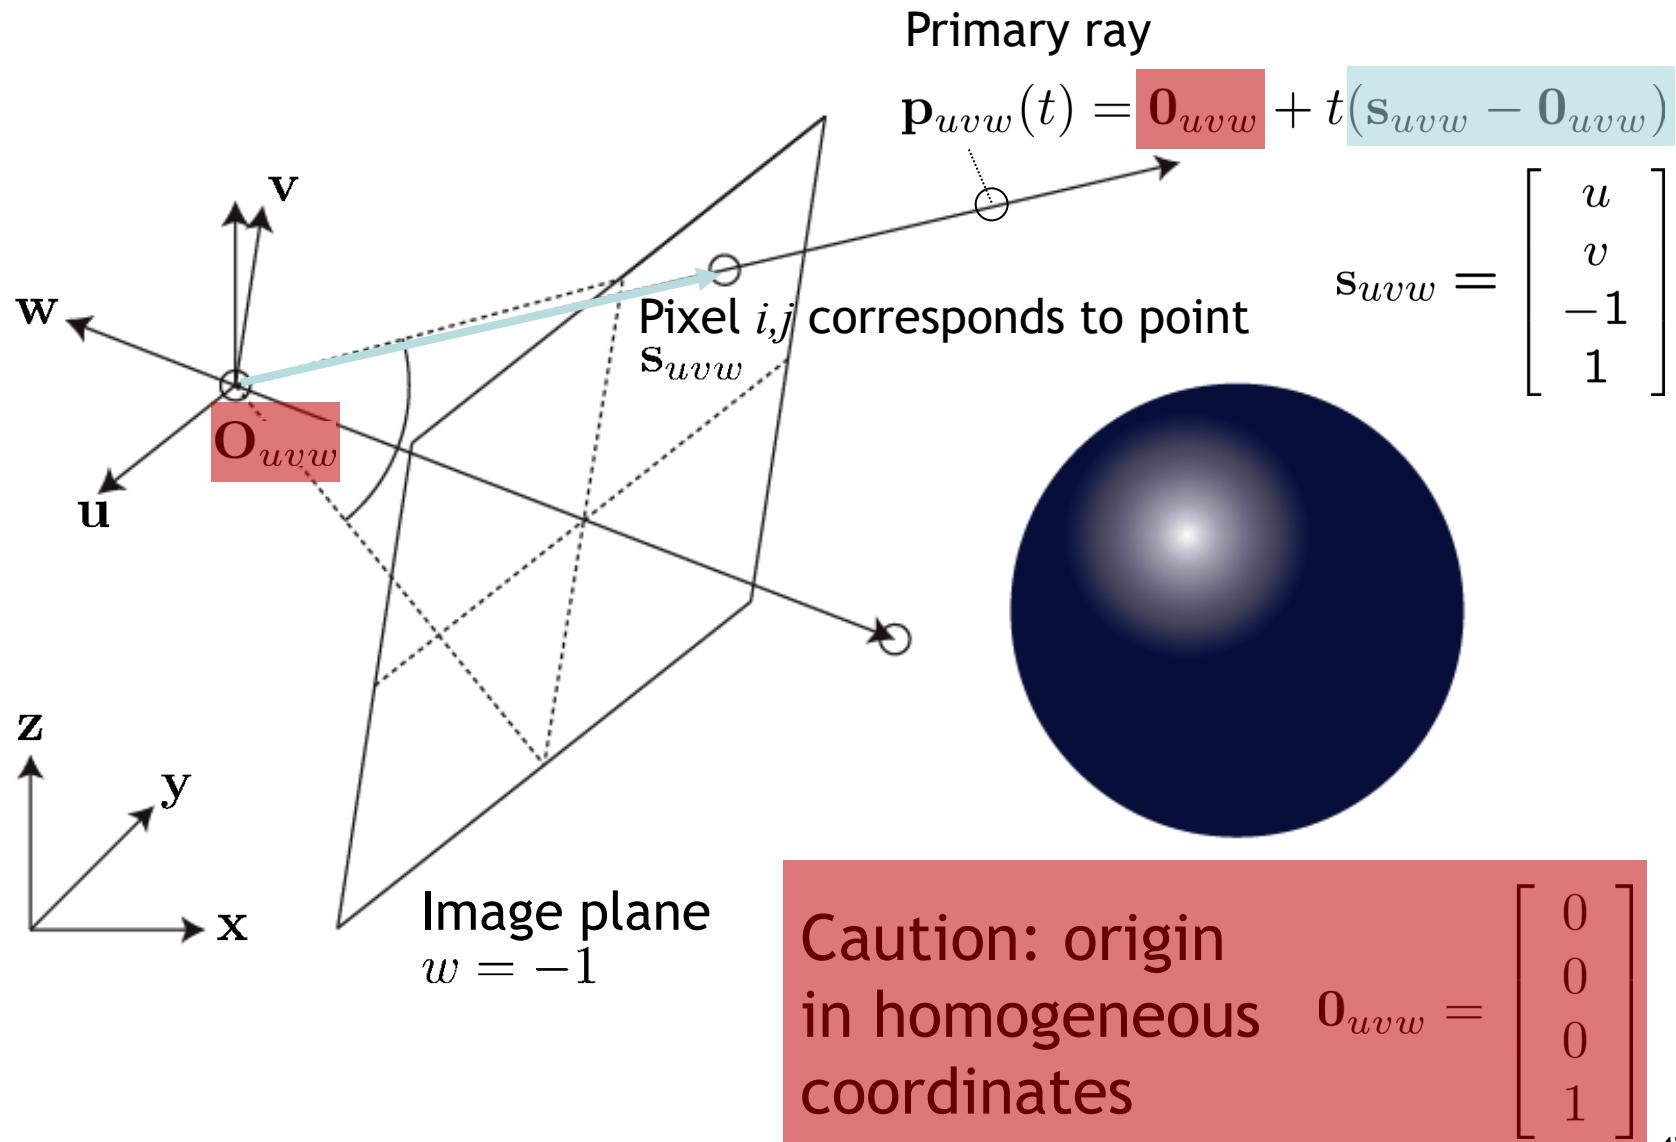


From pixel i,j to primary viewing ray



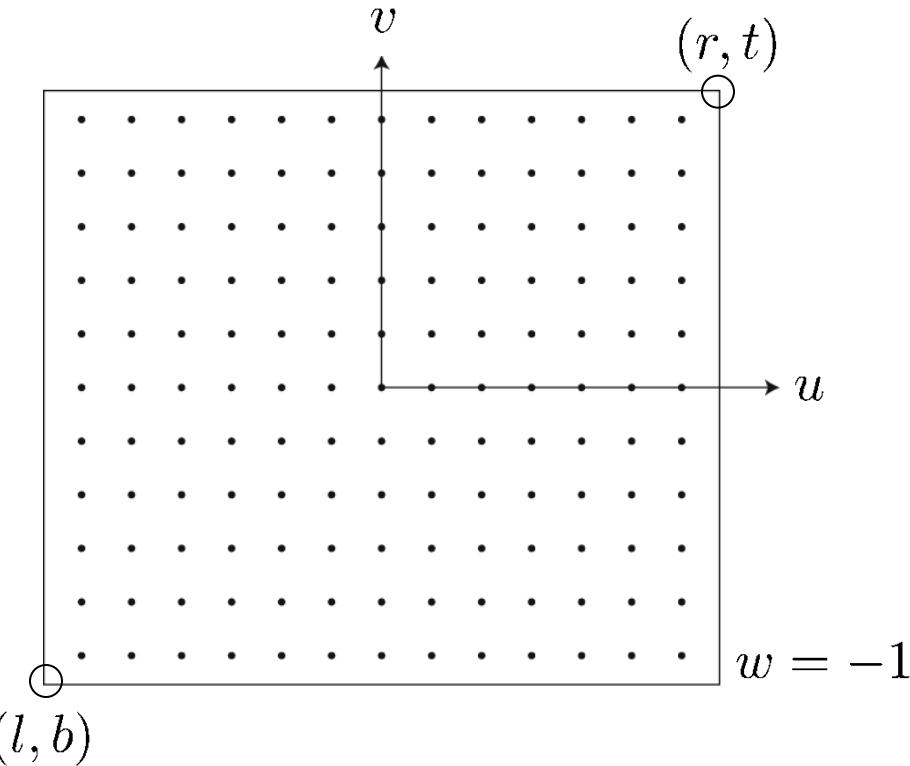
Intrinsic camera matrix K maps pixel coordinates (i, j) to ray in camera coordinates

Primary rays in camera coords.



Bounds of image plane t, b, l, r

- Vertical field-of-view θ
- Aspect ratio width/height $aspect$
- Image coordinates u, v



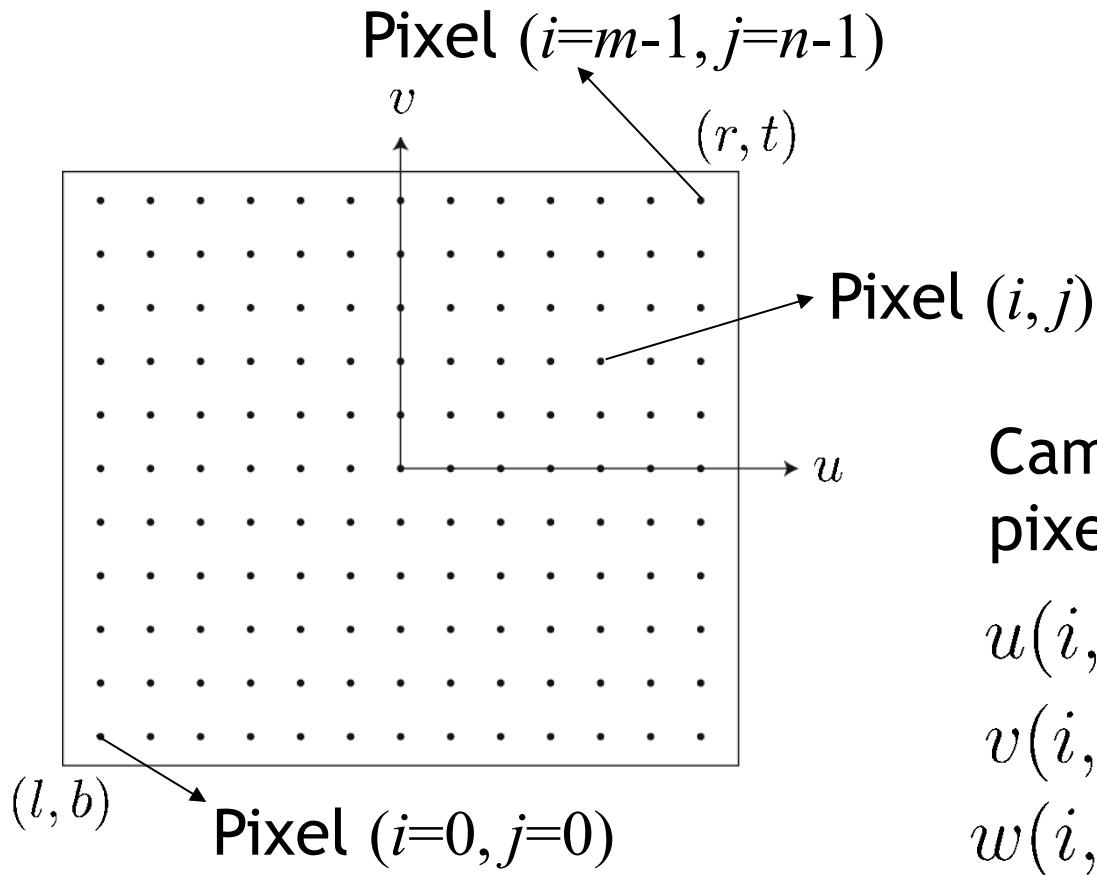
Top, bottom, left, right
coords. of image corners

$$t = -b = \tan(\theta/2)$$

$$r = -l = aspect \cdot t$$

Image plane position u, v given pixel index i, j

- Image resolution $m \times n$ pixels



Camera coords. of **center** of pixel i, j

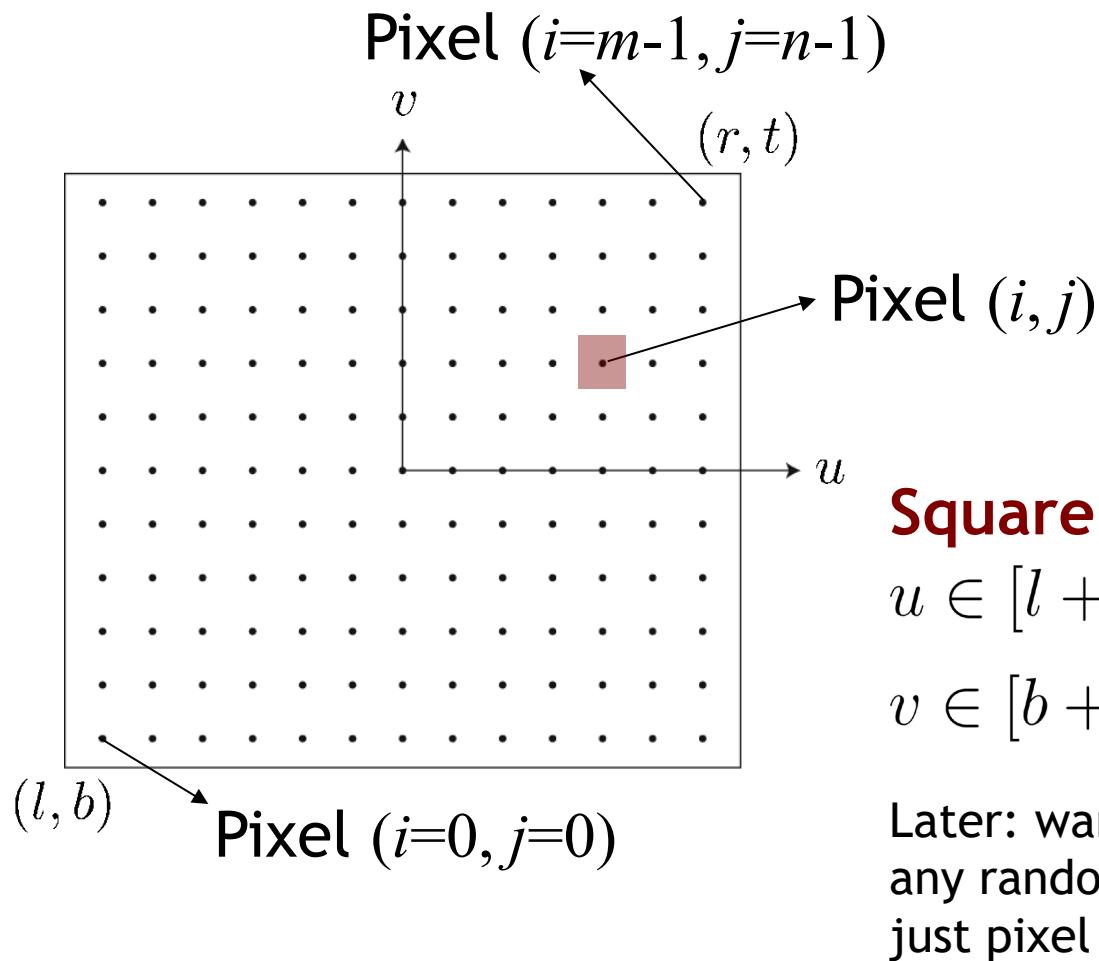
$$u(i, j) = l + (r - l) \frac{i+0.5}{n}$$

$$v(i, j) = b + (t - b) \frac{j+0.5}{n}$$

$$w(i, j) = -1 \quad \text{by convention}$$

Area of pixel i,j in u,v cords.

- Image resolution $m \times n$ pixels



Square area around pixel i, j

$$u \in [l + (r - l)\frac{i}{m}, l + (r - l)\frac{i+1}{m}]$$

$$v \in [b + (t - b)\frac{j}{n}, b + (t - b)\frac{j+1}{n}]$$

Later: want primary ray for pixel through any random position within pixel, not just pixel center

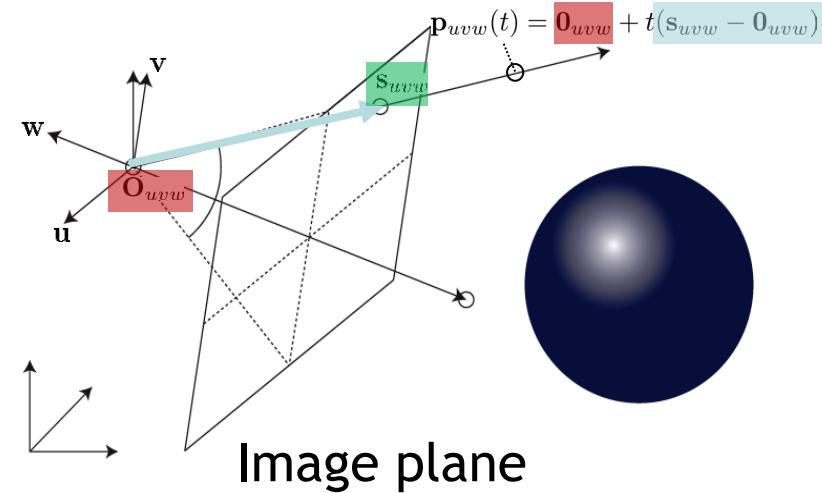
Primary rays in camera coords.

- Given location (u, v) on image plane, get ray $\mathbf{p}_{uvw}(t)$ in camera coords. through (u, v)

$$\mathbf{p}_{uvw}(t) = \mathbf{0}_{uvw} + t(\mathbf{s}_{uvw} - \mathbf{0}_{uvw})$$

$$\begin{aligned}\mathbf{s}_{uvw} - \mathbf{0}_{uvw} &= \begin{bmatrix} u(i, j) \\ v(i, j) \\ -1 \\ 1 \end{bmatrix} - \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix} \\ &= K \begin{bmatrix} i \\ j \\ 0 \\ 1 \end{bmatrix}\end{aligned}$$

Intrinsic camera matrix K



Transformation to world coords.

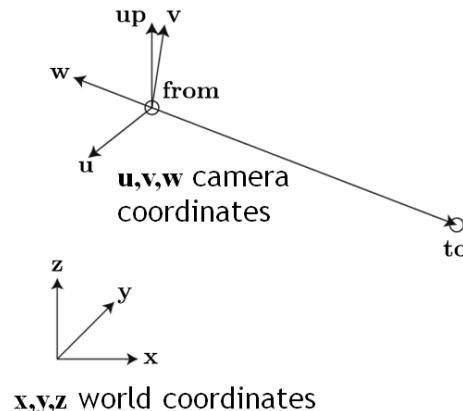
- World coordinates also called “canonic” coordinates
- Basis vectors of camera coords: $\mathbf{u}, \mathbf{v}, \mathbf{w}, \mathbf{e}$ column vectors
- Change of basis matrix (**camera to world, extrinsic camera matrix**)

https://en.wikipedia.org/wiki/Change_of_basis

$$\mathbf{M} = [\mathbf{u} \ \mathbf{v} \ \mathbf{w} \ \mathbf{e}]$$

- Need ray $\mathbf{p}_{xyz}(t)$ in world coordinates

$$\begin{aligned}\mathbf{p}_{xyz}(t) &= \mathbf{M}\mathbf{p}_{uvw}(t) \\ &= \mathbf{M}\mathbf{0}_{uvw} + t(\mathbf{M}\mathbf{s}_{uvw} - \mathbf{M}\mathbf{0}_{uvw})\end{aligned}$$



Ray tracing pseudocode

```
rayTrace() {  
    construct scene representation  
  
    for each pixel  
        ray = computePrimaryViewRay( pixel )  
        hit = first intersection with scene  
        color = shade( hit ) // using shadow ray  
        set pixel color  
}
```

Python renderer notes

- Instead of using loops, code operates on tensors
 - Tensors represent entire sets of rays
- Advantages
 - Underlying implementation hides complexity of parallelization on multi-core CPUs, GPUs

Ray-surface intersection

- Implicit surfaces
- Parametric surfaces

Implicit surfaces

http://en.wikipedia.org/wiki/Surface#Extrinsically_defined_surfaces_and_embeddings

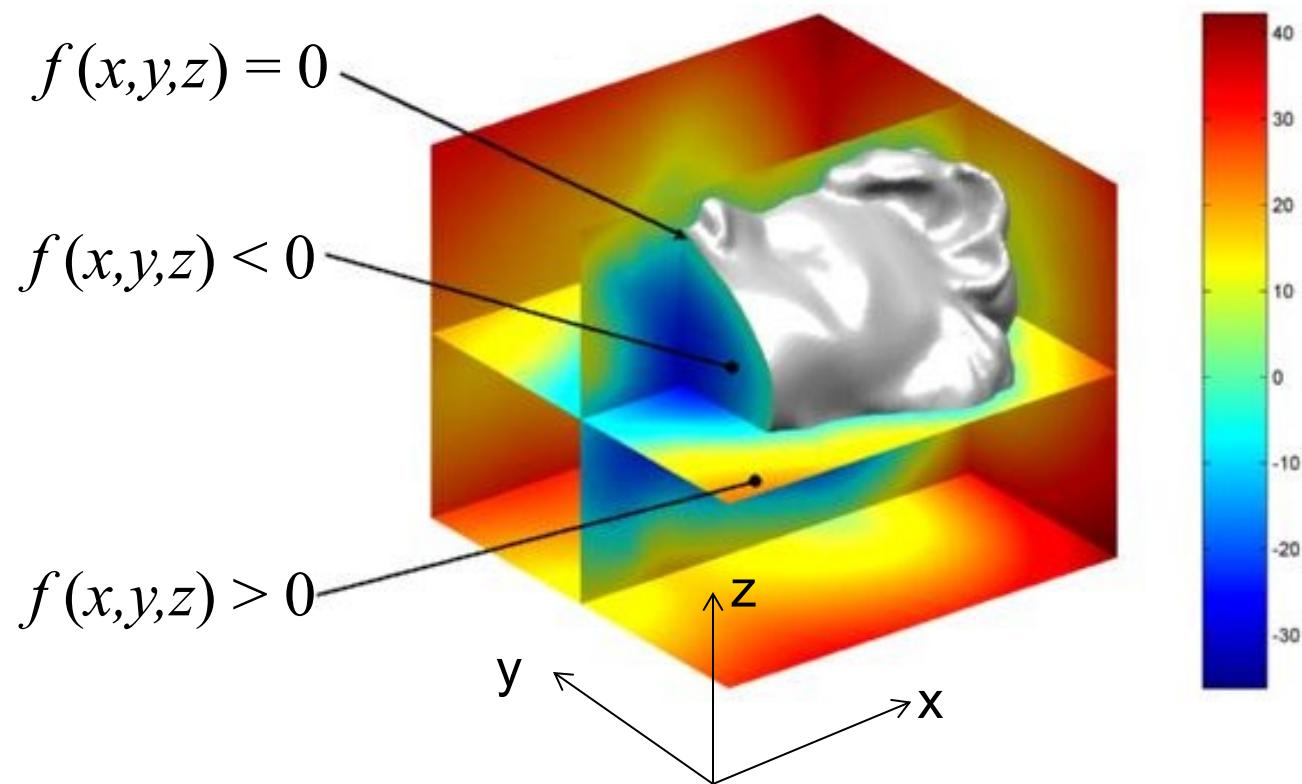
- Surface defined as set of points \mathbf{p} such that

$$f(\mathbf{p}) = 0, \quad \text{where} \quad f : \mathbf{R}^3 \rightarrow \mathbf{R}, \mathbf{p} = (x, y, z)$$

- Implicit representation defines a **solid**
 - Including the interior volume
 - Not just the surface

Implicit surfaces

- 3D visualization, $\mathbf{p}=(x,y,z)$



Implicit surfaces

http://en.wikipedia.org/wiki/Surface#Extrinsically_defined_surfaces_and_embeddings

Ray-surface intersection

- Given a ray with origin e , direction d

$$p(t) = e + td, \quad p, e, d \in \mathbf{R}^3, t \in \mathbf{R}$$

- Ray-surface intersection: solve for t such that

$$f(p(t)) = 0$$

- In general, non-linear equation of variable t
 - Use numerical root-finding algorithm
http://en.wikipedia.org/wiki/Root-finding_algorithm
 - For example, secant method
http://en.wikipedia.org/wiki/Secant_method

Infinite plane

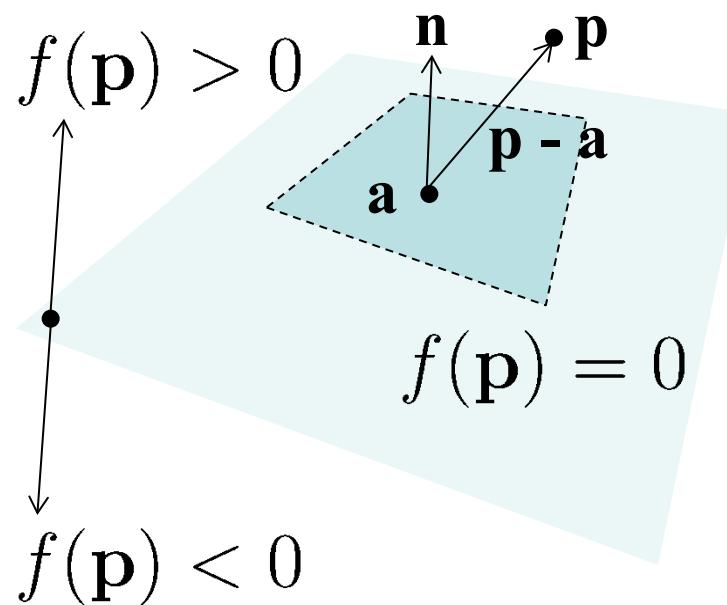
- Unit normal \mathbf{n} , arbitrary point on plane \mathbf{a}
- Plane defined as set of points \mathbf{p} where

$$f(\mathbf{p}) = (\mathbf{p} - \mathbf{a}) \cdot \mathbf{n} = 0$$

- $f(\mathbf{p})$ is distance to plane!
- Intersection with $\mathbf{p}(t)$

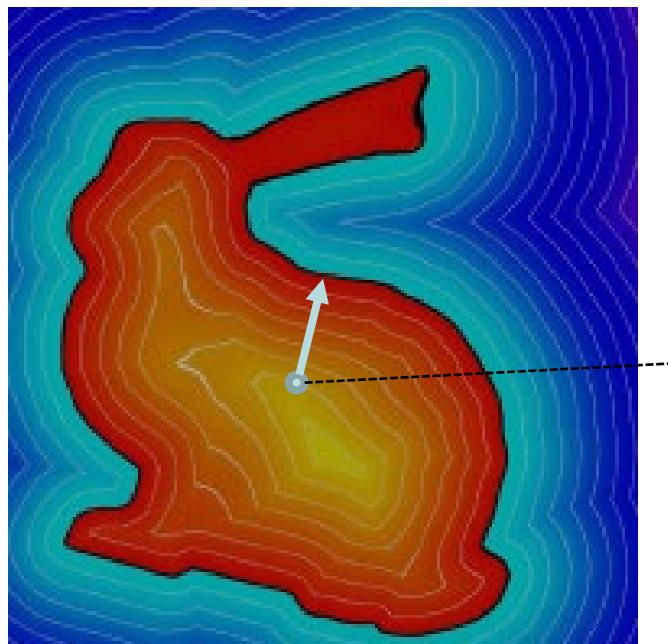
$$(\mathbf{e} + t\mathbf{d} - \mathbf{a}) \cdot \mathbf{n} = 0$$

Point on ray
 $t = \frac{(\mathbf{a} - \mathbf{e}) \cdot \mathbf{n}}{\mathbf{d} \cdot \mathbf{n}}$



Signed distance function (SDF)

- Implicit representation of surfaces where function values happen to represent signed distance to closest point on surface



Value of f = signed distance
to surface

Parametric surfaces

- Parametric surfaces are **boundary representations** (b-reps)

http://en.wikipedia.org/wiki/Boundary_representation

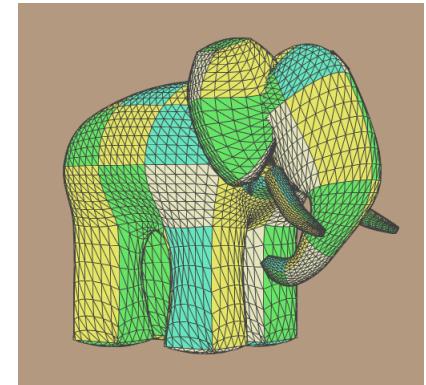
- „Solid object is given by describing its boundary between solid and non-solid“
- In general, surfaces can be described either implicitly or parametrically

Parametric surfaces

- Parameters u, v , typically in rectangular domain

$$\left. \begin{array}{l} x = f(u, v) \\ y = g(u, v) \\ z = h(u, v) \end{array} \right\} \text{Point } \mathbf{p} \text{ on surface } \mathbf{p} = (x, y, z)$$

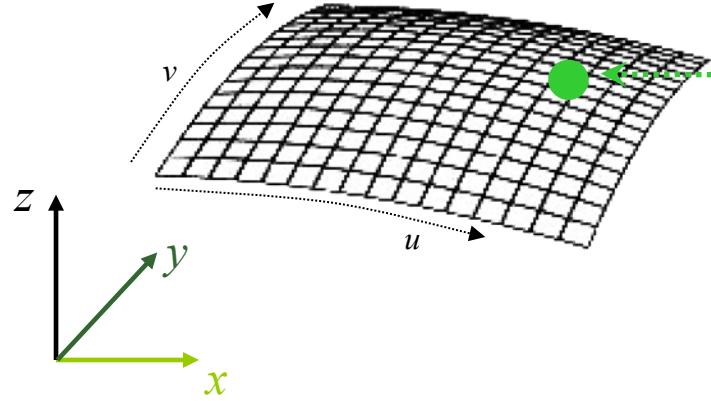
- Complicated surfaces require fitting together many parametric patches (charts)



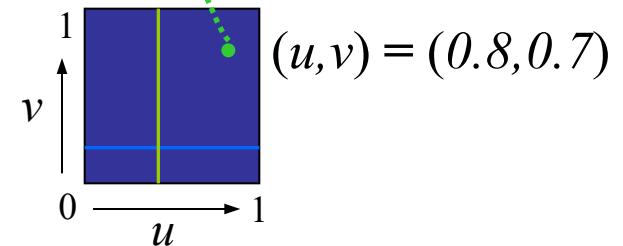
http://en.wikipedia.org/wiki/B%C3%A9zier_surface

Parametric patches (charts)

- Visualization



$$\begin{aligned}x &= f(0.8, 0.7) \\y &= g(0.8, 0.7) \\z &= h(0.8, 0.7)\end{aligned}$$



2D parameter domain,
quadrilateral patch

Parametric surfaces

- Ray-surface intersection

Point on ray = point on surface

$$e_x + t d_x = f(u, v)$$

$$e_y + t d_y = g(u, v)$$

$$e_z + t d_z = h(u, v)$$

Ray origin, direction:

$$\mathbf{e} = (e_x, e_y, e_z)$$

$$\mathbf{d} = (d_x, d_y, d_z)$$

- Solve 3 equations for 3 unknowns t, u, v
- Easy to solve for parametric planes (see ray-triangle intersection)
- Non-linear equation for higher order surfaces (e.g., NURBS)
 - Requires iterative solution

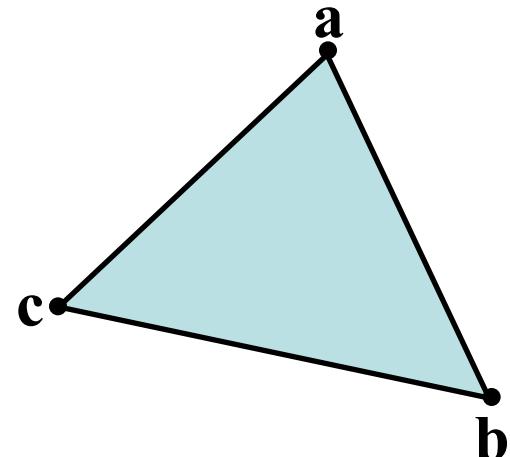
Ray-triangle intersection

- Parametric description of a triangle with vertices $\mathbf{a}, \mathbf{b}, \mathbf{c}$ (parameters α, β instead of u, v)

$$\mathbf{p} = \mathbf{a} + \beta(\mathbf{b} - \mathbf{a}) + \gamma(\mathbf{c} - \mathbf{a})$$

where $\beta > 0, \gamma > 0, \beta + \gamma < 1$

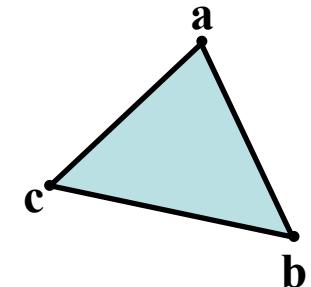
- Convention: vertices are ordered **counter-clockwise** as seen from the **outside**



Ray-triangle intersection

- Outward-pointing normal

$$\mathbf{n} = (\mathbf{b} - \mathbf{a}) \times (\mathbf{c} - \mathbf{a})$$



- Triangle in **barycentric coordinates** α, β, γ

$$\mathbf{p} = \alpha\mathbf{a} + \beta\mathbf{b} + \gamma\mathbf{c}$$

$$\text{where } \alpha + \beta + \gamma = 1,$$

$$0 < \alpha < 1, 0 < \beta < 1, 0 < \gamma < 1$$

- Given two of the coordinates, can easily compute third, e.g. $\alpha = 1 - \beta - \gamma$

Ray-triangle intersection

- Intersection condition

$$\mathbf{e} + t\mathbf{d} = \mathbf{a} + \beta(\mathbf{b} - \mathbf{a}) + \gamma(\mathbf{c} - \mathbf{a})$$

Point on ray Point on triangle

- One equation for each coordinate x,y,z

$$e_x + t d_x = a_x + \beta(b_x - a_x) + \gamma(c_x - a_x)$$

$$e_y + t d_y = a_y + \beta(b_y - a_y) + \gamma(c_y - a_y)$$

$$e_z + t d_z = a_z + \beta(b_z - a_z) + \gamma(c_z - a_z)$$

Ray-triangle intersection

- In matrix form

$$\begin{bmatrix} a_x - b_x & a_x - c_x & d_x \\ a_y - b_y & a_y - c_y & d_y \\ a_z - b_z & a_z - c_z & d_z \end{bmatrix} \begin{bmatrix} \beta \\ \gamma \\ t \end{bmatrix} = \begin{bmatrix} a_x - e_x \\ a_y - e_y \\ a_z - e_z \end{bmatrix}$$

- Solve for β, γ, t using Cramer's rule
(Shirley page 157 
- Ray intersects triangle if $0 < \beta + \gamma < 1$

Example: unit sphere

- Radius 1, centered at origin (0,0,0)
- Implicit
 - $f(x,y,z) = x^2 + y^2 + z^2 - 1$
- Parametric
 - Spherical coordinates (polar, azimuthal angles) θ, ϕ
 - $x = \sin(\theta)\cos(\phi), y = \sin(\theta)\sin(\phi), z = \cos(\theta)$

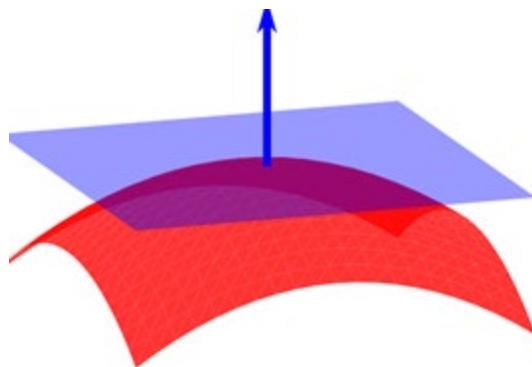
What is easier to ray trace, implicit or parametric sphere?

- A. Implicit equations
for sphere
- B. Parametric
equations for
sphere

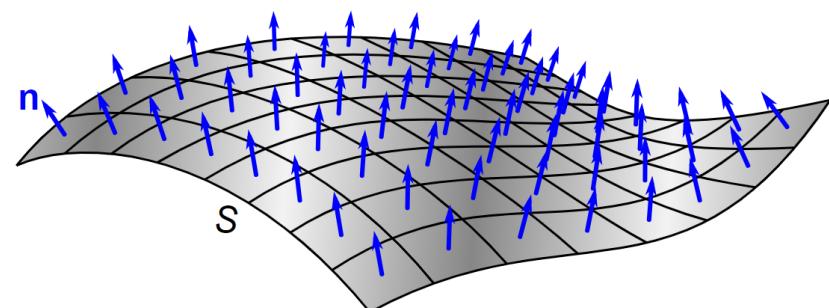
Surface normals

[https://en.wikipedia.org/wiki/Normal_\(geometry\)](https://en.wikipedia.org/wiki/Normal_(geometry))

- Unit vector perpendicular to tangent plane at each point on surface (i.e., surface normal is normal of tangent plane at each point)
- Required for shading calculations



Surface normal is
perpendicular to tangent
plane at each point



Normal field of curved
surface

Surface normals

- How to compute for
 - Parametric surfaces?
 - Implicit surfaces?

Surface normals

- Implicit surfaces: gradient of f

$$\mathbf{n} = \nabla f(\mathbf{p}) = \left(\frac{\partial f(\mathbf{p})}{\partial x}, \frac{\partial f(\mathbf{p})}{\partial y}, \frac{\partial f(\mathbf{p})}{\partial z} \right)$$

- Parametric surface: cross product of tangent vectors (partial derivatives)

$$\mathbf{n}(u, v) = \left(\frac{\partial x}{\partial u}, \frac{\partial y}{\partial u}, \frac{\partial z}{\partial u} \right) \times \left(\frac{\partial x}{\partial v}, \frac{\partial y}{\partial v}, \frac{\partial z}{\partial v} \right)$$

Transforming Normal Vectors

- Normal \mathbf{n} is perpendicular to tangent \mathbf{t}

$$\mathbf{n}^T \cdot \mathbf{t} = 0$$

- Under transformation \mathbf{Mt} , find transformation \mathbf{X} of normal \mathbf{n} that satisfies normal constraint

$$(\mathbf{Xn})^T \cdot \mathbf{Mt} = 0$$

$$\mathbf{n}^T \mathbf{X}^T \mathbf{Mt} = 0$$

$$\Rightarrow \mathbf{X}^T \mathbf{M} = \mathbf{I} \Rightarrow \mathbf{X} = \mathbf{M}^{-1^T}$$

- Note: $\mathbf{M}^{-1^T} = \mathbf{M}$ if and only if \mathbf{M} includes only rotation and translation (no shear)

Rendering assignments

- Assignments available on UMD ELMS/Canvas
<https://myelms.umd.edu>
- Extend an educational renderer based on Python and tensors
- First assignment
 - Getting started with base code
 - Deadline: Thursday Sept. 11
 - Turn in via file upload to ELMS/Canvas

Homework policies

- No public git forks of your homework!
 - Of course, a private git repo on github or elsewhere in the cloud is a good idea
- Adhere to the academic integrity policies of UMD/UMD CS

<http://osc.umd.edu/OSC/Default.aspx>

<http://www.cs.umd.edu/class/resources/academicIntegrity.html>

- Discussing homework with colleagues is fine, but no sharing of code
- No copying of code found online
- No AI generated code

Next

- Acceleration structures for ray tracing
- Note: no class session on 9/4, instead
 - Read through slides independently
 - Ask questions to AI chatbot (Canvas Virtual Study Assistant, or other one)
 - Submit assignment (most interesting question and answer from AI, in your opinion) on Canvas