

**CMSC740**  
**Advanced Computer Graphics**

Fall 2025  
Matthias Zwicker

# Background terminology

---

The course assumes you are familiar with the following terms. If you are not, you should read up on these via Wikipedia, chatbots, etc.

- (Artificial) neuron
  - Weights, bias, activation function
- Neural network architectures
  - Layers, hidden layers
  - Convolutional neural networks
- Normalization
  - Batch normalization
- Supervised learning
  - Loss functions, including L2 loss, L1 loss
  - Training data
- Stochastic gradient descent
  - Automatic differentiation/backpropagation
  - Adam

# Deep learning/AI in graphics

---

- What problems can we address?
- What is special about deep learning? Why is it useful compared to other techniques?
- Limitations of deep learning/AI techniques?

# Problems we'd like to address with deep learning

---

## In computer graphics/vision

- Image/video -> class label
- Image/video -> set of 2D regions containing objects
- Image/video -> text description
- Image/video -> 3D geometry
- Low quality image -> high quality image
- Text description -> 3D object
- Text description -> video

## Others

- Text in one language -> Translation to other language
- Text question -> answer
- Information about an environment (e.g. sensor data) -> best action of an agent to reach a goal
- Etc.

**Abstractly:** find mapping from complex, high-dimensional inputs to complex, high-dimensional outputs


# Deep neural networks

---

- Powerful type of non-linear functions
  - Specific function given by (potentially huge, billions) number of **parameters (weights, biases)**
  - Can **approximate any input-output relationship** by finding suitable parameters (universal approximation theorem) [https://en.wikipedia.org/wiki/Universal\\_approximation\\_theorem](https://en.wikipedia.org/wiki/Universal_approximation_theorem)
  - Work with **high-dimensional inputs, outputs**
  - Work with “brute force”, raw data representations (e.g., image/video as long vector of pixel values)
- Alternative types of non-linear functions?

# Deep neural networks


- Basic computational element: artificial neurons
  - Inputs  $x_i$ , output  $y_{\text{pred}}$
  - Defined by weights  $w_i$ , bias  $b$ , activation function  $f$
  - Parameters to be optimized: **weights, biases**



## Activation Functions


### Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




### tanh

$$\tanh(x)$$




### ReLU

$$\max(0, x)$$




### Leaky ReLU

$$\max(0.1x, x)$$



### Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$




### ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$

# Deep neural networks

- Organized in **layers**
  - Thousands to millions of neurons per layer
  - Dozens to hundreds of layers (interior layers called **hidden layers**)
  - Billions to trillions of connections (weights)
  - Connections between layers, within layers, or across layers (skipping layers inbetween)
- Architectures (how neurons are connected)
  - Multilayer perceptrons (MLPs)
  - Convolutional neural networks (CNNs)
  - Residual networks
  - Transformers, multi-head attention
  - Etc.



# Supervised learning

---

- Training data: set of corresponding **pairs of desired inputs and outputs** (e.g., image, text description, etc.)
- Objective: Find **mapping (function)** that fits/approximates training data; i.e., map each input to corresponding known output in training data
  - Minimize **scalar loss function** to quantify quality of fit/approximation
- Questions/challenges?
  - Collecting training data
  - Mathematical form of mapping, loss function
  - How to represent input, output
  - How to find mapping function that achieves objective (minimizes loss)

# Stochastic gradient descent

---

- Simple algorithm to **find function parameters (weights, biases)**, given training data and **loss function**
  - Loss function: quantify difference between training data (desired outputs) and network outputs as scalar number
- Iteratively
  - Compute gradient of loss function wrt. function parameters (weights, biases), using random (stochastic) subset (batch) of training data
  - Gradient: vector of same length as parameters, gives small change for each parameter that will reduce loss the most
  - Update parameters using gradient
- Advantages/disadvantages/alternatives?

# Backpropagation

---

- Algorithm to compute gradient of neural network computations using automatic differentiation
- **Automatic differentiation:** algorithm for automatically computing derivatives of arbitrary concatenation of elementary functions (with known derivatives) using chain rule
  - Implemented using software libraries that maintain data structures, intermediate results to calculate chain rule “in the background”, without user requiring to add much code

# (Mostly) open questions

---

- Generalization: how to characterize behavior of mapping function on inputs not seen in training data?
- Robustness: how to ensure that small changes in inputs or training data preserve accuracy of mapping function?
- Architecture design: given training data, how to characterize network architectures that will work well (accuracy on training data, generalization, robustness)?

# Recap

---

- Rendering equation
- Monte Carlo path tracing
- (Multiple) importance sampling
  - Bidirectional path tracing
- Advanced appearance models
  - BRDFs/BSDFs/BSSRDFs to model specific materials
- Participating media

# The good

---

- Simulate light transport in a physically accurate (in relation to human perception) manner
- Include complex material appearance models, complex geometry
- Generate photorealistic images (with sufficient modeling effort and render time)



<https://corona-renderer.com/gallery>

# The bad


---

- Noise/variance, slow convergence,  $O(1/\sqrt{n})$



55 sec, 32 spp, 12 threads



<http://cgg.unibe.ch/data/PG2013/>



100 sec, 128 spp, 12 threads

# The solution?

---




# The solution?

---



# Image-to-image translation




Various applications see <https://phillipi.github.io/pix2pix/>

Neural network architectures suitable  
for image-to-image translation?

# DnCNN: straightforward denoising with deep CNNs

- “Beyond a Gaussian Denoiser: Residual Learning of Deep CNN for Image Denoising”, Zhang et al., 2016
- Training data: pairs of noisy/clean images (add known noise to clean images)  
<https://github.com/cszn/DnCNN>
- 20 layers, 64 3x3 convolution filters in each layer (64 channels per layer)
- Batch normalization (additional computation to scale, shift outputs of each layer to zero mean, unit variance in each training batch)
- Predict residual (known noise) instead of pixel value
- L2 loss: for each pixel, square of difference between ground truth and network output; sum over all pixels, (no adversarial, L1, or feature loss)
- No downsampling/upsampling, no skip connections
- Train/optimize weights, biases on multiple noise levels




Conv: convolutional layer; BN: batch normalization; ReLU: rectified linear activation fct.

# DnCNN

- Ablation study (comparison among different design decisions)
- Epoch: number of passes through entire training data set during gradient descent


Accuracy



(a) SGD

Training progress

Stochastic gradient descent



(b) Adam

Adam (adaptive moment estimation) optimizer, improved version of SGD

# DnCNN




Fig. 6. Color image denoising results of one image from the DSD68 dataset with noise level 35.





Fig. 7. Color image denoising results of one image from the DSD68 dataset with noise level 45.

# Denoising, generic approach

- Noisy pixel values  $c_p$  at each pixel  $p$
- Denoising: replace each pixel with weighted average over neighbor pixels  $c_q$

$$\hat{c}_p = \sum_{q \in N(p)} w_{pq} c_q$$



- Weights should be  $0 < w < 1$ , add up to one, (square) neighborhood  $N_p$
- Assuming pixels  $c_q$  are independent random variables
  - Variance

$$V \left[ \sum_{q \in N(p)} w_{pq} c_q \right] = \sum_{q \in N(p)} (w_{pq}^2 V [c_q])$$

Variance reduction  
 $\sum_{q \in N(p)} w_{pq}^2 << 1$

- Expected value, if  $c_q$  are i.i.d (independent, identically distributed, [https://en.wikipedia.org/wiki/Independent\\_and\\_identically\\_distributed\\_random\\_variables](https://en.wikipedia.org/wiki/Independent_and_identically_distributed_random_variables))

$$E \left[ \sum_{q \in N(p)} w_{pq} c_q \right] = \left( \sum_{q \in N(p)} w_{pq} \right) E [c_p] = E [c_p], \quad \text{where i.i.d implies } E [c_p] = E [c_q]$$

# Denoising, classical approach

---

- Example: box filter,  $w_{pq} = 1/|N(p)|$ 
  - Variance reduction by  $1/|N(p)|$
  - Larger filter (larger size of neighborhood  $|N(p)|$ ) leads to more variance reduction
  - But: **bias** (pixels are not i.i.d in practice, i.e.  $E[c_p] \neq E[c_q]$ , hence box filter changes expected value of pixel, in general), causes blur
- Goal: techniques to determine filter weights in range  $0 < w < 1$ , such that bias is avoided
- Ideas?
  - Make  $w_{pq}$  high if expected pixel values  $E[c_p]$  and  $E[c_q]$  are similar, otherwise, make  $w_{pq}$  low


# Denoising using kernel prediction

---

- “Kernel-Predicting Convolutional Networks for Denoising Monte Carlo Renderings”, Bako et al., 2017
- Predict weights  $w_{pq}$  (“kernel”) in neighborhood of each pixel using neural network
- Other ideas to improve performance of denoising in the case of Monte Carlo rendered images?

# CNN to denoise Monte Carlo renderings

“Kernel-Predicting Convolutional Networks for Denoising Monte Carlo Renderings”, Bako et al., 2017



Any per-pixel data produced by renderer that may be useful for denoising (depth, normals, BRDF parameters)

Bako et al., SIGGRAPH 2017

[http://cvc.ucsbd.edu/graphics/Papers/SIGGRAPH2017\\_KPCN/](http://cvc.ucsbd.edu/graphics/Papers/SIGGRAPH2017_KPCN/)

# Specialties (vs usual denoising CNNs)

---

- Take advantage of additional information (besides RGB) produced by renderer
  - Additional per-pixel input channels (normal, depth, BRDF parameters) stacked together with RGB colors
- Diffuse/specular separation of input
  - Diffuse: albedo divide
  - Specular: color-luminance separation, log-transform

# Specialties (vs usual denoising CNNs)

---

## Denoising kernel prediction

- Network output at each pixel  $p$  is  $k \times k$  vector  $\mathbf{z}_p^L$
- Weight for pixel  $q$  in filter of pixel  $p$

$$w_{pq} = \frac{\exp([\mathbf{z}_p^L]_q)}{\sum_{q' \in \mathcal{N}(p)} \exp([\mathbf{z}_p^L]_{q'})},$$

Ensure weight normalization  
(sum to 1)

- Filtered output at pixel  $p$

$$\hat{\mathbf{c}}_p = g_{\text{weighted}}(\mathbf{x}_p; \theta) = \sum_{q \in \mathcal{N}(p)} \mathbf{c}_q w_{pq}.$$

- Advantage: can prove faster training convergence in simplified case; better generalization in practice

# Network, training & test data

---

## Network

- 8 hidden layers, 100 convolution kernels (channels) of  $5 \times 5$  in each layer
- L1 loss (for each pixel, absolute difference between network output and ground truth, sum over all pixels)


## Training data

- 600 frames from movie “Finding Dori”
  - Noisy: 32 spp
  - Reference: 1024 spp

## Test data

- 25 frames from other movies
  - Motion blur, depth of field, glossy reflections, and global illumination

# Results



Bako et al., SIGGRAPH 2017

[http://cvc.ucsb.edu/graphics/Papers/SIGGRAPH2017\\_KPCN/](http://cvc.ucsb.edu/graphics/Papers/SIGGRAPH2017_KPCN/)


# Denoising animations

---

- Frame-by-frame denoising leads to temporal flickering in animations
- Need to filter in 3D space-time
- How to implement this with CNNs?

# Improvement of Bako et al. 2017


- “Denoising with Kernel Prediction and Asymmetric Loss Functions”, Vogels et al., ACM Siggraph 2018
  - Residual blocks (avoid vanishing gradient problem, [https://en.wikipedia.org/wiki/Vanishing\\_gradient\\_problem](https://en.wikipedia.org/wiki/Vanishing_gradient_problem))
  - Modular training (source encoder, spatial-feature extractor, kernel predictor) allows specialization to different renderers with less training data



Input of res. block  
added to its output,  
forming “shortcut”

# Temporal denoising


- Input  $2M+1$  frames simultaneously
- Use source encoder and spatial-feature extractor modules pre-trained on single frames
- Warp frames to align with reference  $i$  using **motion vectors or optical flow**
- Predict single  $21 \times 21 \times (2M+1)$  kernel to denoise frame  $i$



<http://zurich.disneyresearch.com/~fabreric/publications/vogels-2018-kpal.pdf>


# Asymmetric loss

- Loss function  $\ell'_\lambda$ 
  - Ground truth  $r$ , network output  $d$ , noisy color  $c$ , Heaviside step function  $H$ , user parameter  $\lambda$ 
$$\ell'_\lambda(d, r, c) = \ell(d, r) \cdot (1 + (\lambda - 1)H((d - r)(r - c)))$$
  - “Penalize network outputs  $d$  more if they are on the other side of the reference than the noisy input  $c$ ”
- User can adjust  $\lambda$  to control bias/variance tradeoff



# Results

---



<http://zurich.disneyresearch.com/~fabricer/publications/vogels-2018-kpal.pdf>

# Interactive, GPU rendering

---


- Two similar products from Nvidia using recurrent convolutional auto-encoder
  - Auto-encoder: network that consists of encoder and decoder parts
- Denoising for ray tracing built into Nvidia Optix
  - Optimized for inference performance, milliseconds per frame
- Nvidia DLSS

<https://www.nvidia.com/en-us/geforce/technologies/dlss/>

# Recurrent neural network RNN

- Chaitanya et al, Interactive Reconstruction of Monte Carlo Image Sequences using a Recurrent Denoising Autoencoder, SIGGRAPH 2017

[http://research.nvidia.com/sites/default/files/publications/dnn\\_denoise\\_author.pdf](http://research.nvidia.com/sites/default/files/publications/dnn_denoise_author.pdf)



# Specialties

---

- Input to network
  - RGB colors
  - “G-buffer” with view-space shading normals (2D vectors), depth, material roughness
  - Divided by diffuse albedo (remove texture complexity from input)
- Loss
  - Spatial: L1, L1 on spatial gradients
  - Temporal: L1 on temporal gradients

# Training

- “Back propagation through time”


[https://en.wikipedia.org/wiki/Backpropagation\\_through\\_time](https://en.wikipedia.org/wiki/Backpropagation_through_time)

- Replicate feed forward parts to unroll recurrent connections

- “Adam” SGD solver

<https://arxiv.org/pdf/1412.6980.pdf>






- Gradients are noisy in SGD
- Compute moving average over SGD steps



Two replicated feed forward steps for frame  $i-1$ ,  $i$ , recurrent connection from  $i-1$  to  $i$

# Results (inputs with 1spp)

[http://research.nvidia.com/sites/default/files/publications/dnn\\_denoise\\_author.pdf](http://research.nvidia.com/sites/default/files/publications/dnn_denoise_author.pdf)

	Our reconstruction result	MC Input	AAF	EAW	SBF	Our	Reference
CORNELLBOX		 	 	 	 	 	 
SPOONZA		 	 	 	 	 	 
GLOSSYSPOONZA		 	 	 	 	 	 

# Results

---

- Video

<http://research.nvidia.com/publication/interactive-reconstruction-monte-carlo-image-sequences-using-recurrent-denoising>

# Computation times

---

- 54.9ms for 1280×720 pixels on NVIDIA (Pascal) Titan X
- OptiX-based path tracer from 70ms to 260ms in SanMiguel for 1 spp

<https://developer.nvidia.com/optix>

# Deep learning for rendering

---

- What else could we learn to improve rendering efficiency?
  - Keep more data from the renderer than 2D images (e.g., operate on 4D images (light fields), for denoising, interpolation)
  - Use image-to-image translation to add expensive light transport effects (global illumination, subsurface scattering, motion blur, ...)
  - Learn the entire rendering function  
scene description -> neural network -> image with global illumination

# Learn to importance sample

---

- “Learning to Importance Sample in Primary Sample Space”, Zheng, Zwicker, 2018  
<https://arxiv.org/abs/1808.07840>
- “Neural Importance Sampling”, Muller et al. 2018  
<https://arxiv.org/abs/1808.03856>