

CMSC740
Advanced Computer Graphics

Fall 2025
Matthias Zwicker

Deep learning in graphics

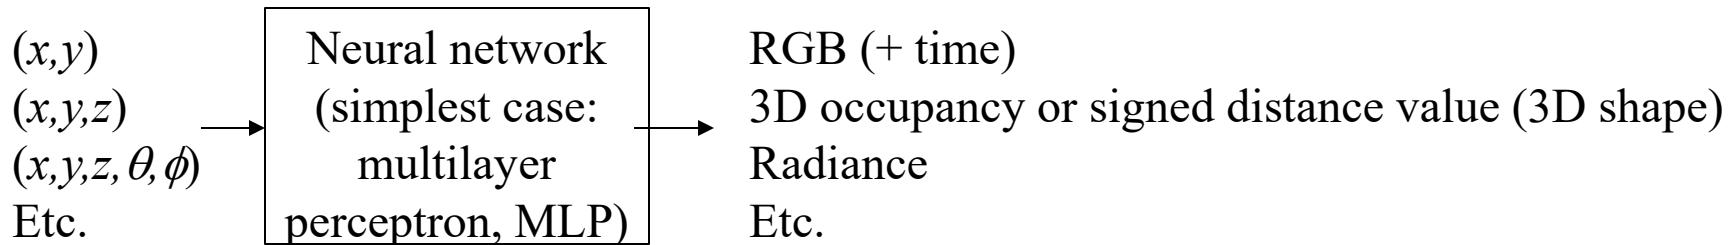
- So far
 - Denoising (map noisy image to clean image)

Neural representations in graphics/vision

- General idea: represent continuous quantities/functions in computer graphics with neural networks
- Examples
 - Geometry (network implements implicit surface representation using signed distance fields, or binary occupancy/inside-outside function)
 - Radiance fields (given ray, network returns radiance)
 - Images, videos (given pixel coordinates, network returns color)
 - Etc.

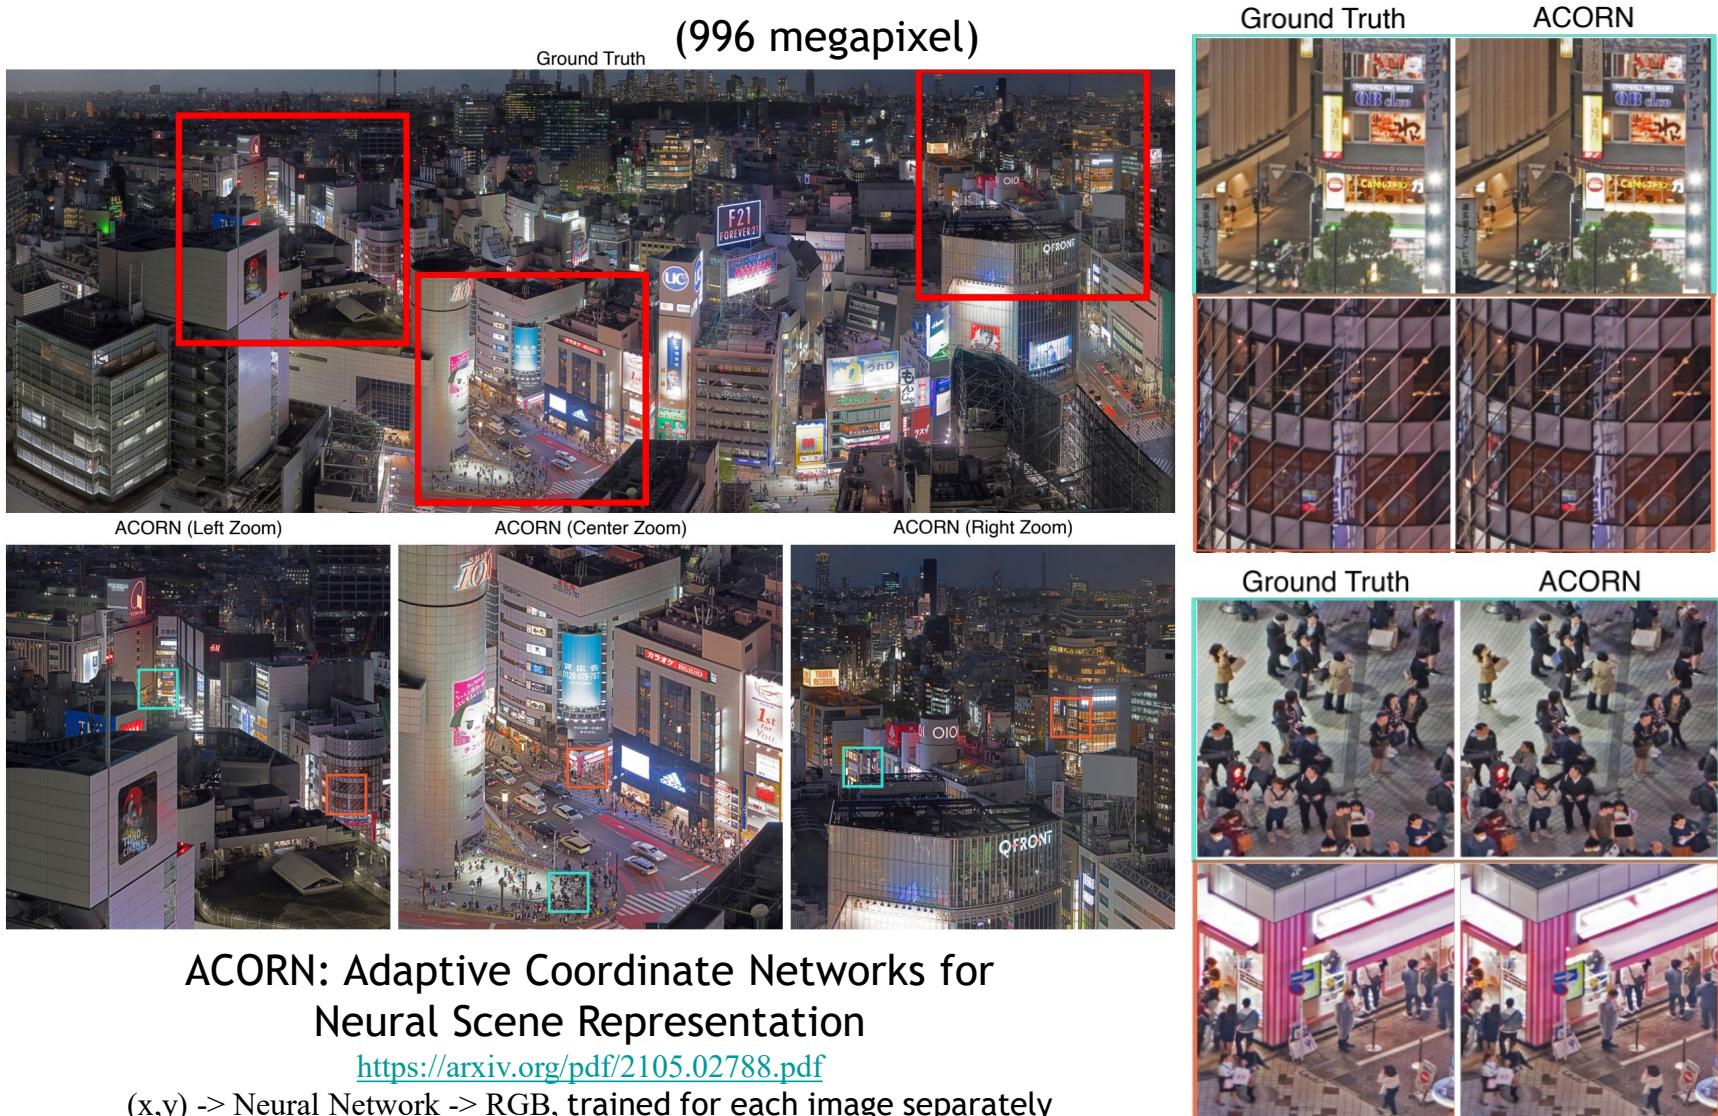
General approach

- Naïve: MLP as function representation for various objects



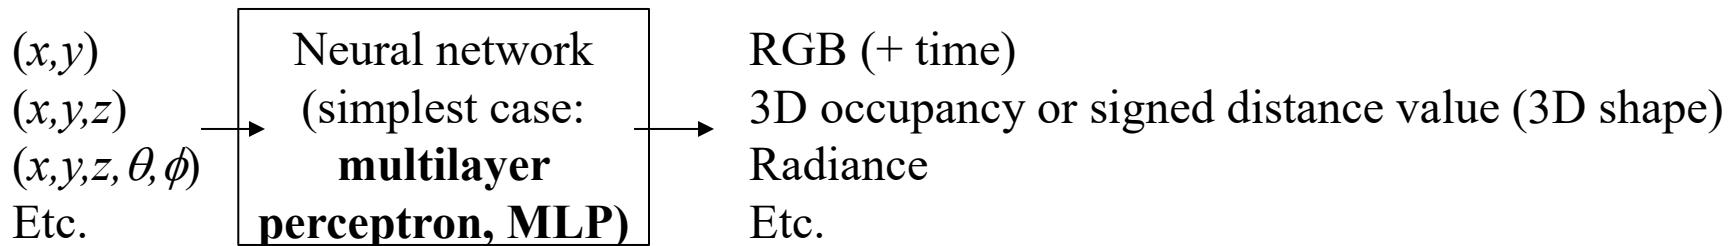
- Approach: train MLP for each “object” individually, using reconstruction loss (e.g., L2-loss distance between output and ground truth)
- “Objects”:
 - Radiance fields
 - 3D geometry
 - Images
 - Videos

Example: Image represented as NN

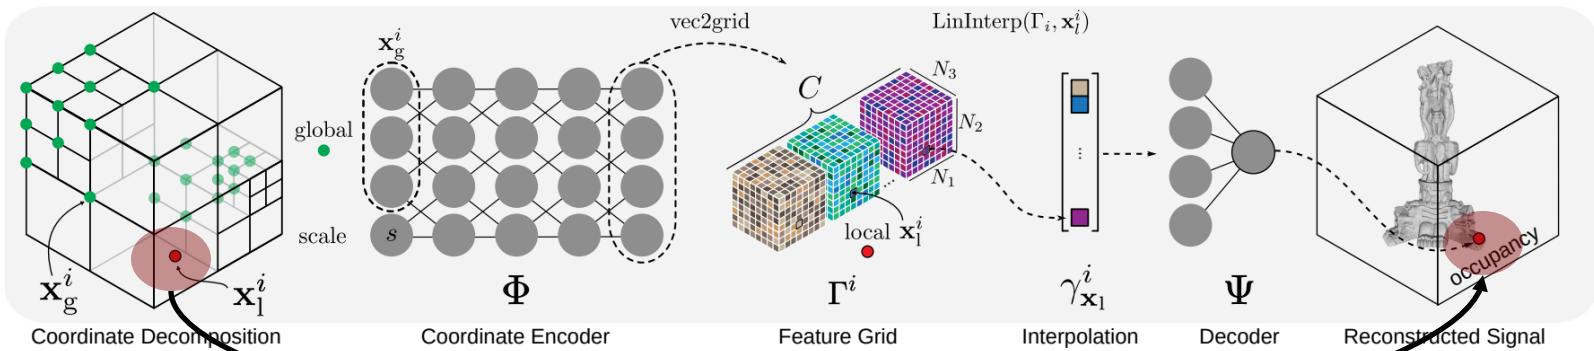


Approach

- Naïve



- Better: replace MLP with more sophisticated architectures (e.g., ACORN: Adaptive Coordinate Networks)



<https://arxiv.org/pdf/2105.02788.pdf>

Advantages, disadvantages

- Conventional techniques: Splines, wavelets, etc.
- Advantages of neural networks
 - Represent complicated functions efficiently, with low storage cost (adaptive, nonlinear); most useful if desired function is given as solution of some equation (PDE, integral equation, etc.)
 - Leverage deep learning infrastructure (GPUs, automatic differentiation, numerical optimization techniques for neural networks, Python libraries)
- Disadvantages of neural networks
 - Requires nonlinear optimization (instead of linear as for some other techniques)
 - Slower to evaluate than conventional techniques
 - Relation of network architectures to accuracy, convergence properties not understood as well as for conventional techniques

Application in graphics: neural representation to solve rendering equation

- First: background on traditional radiosity technique to solve rendering equation

[https://en.wikipedia.org/wiki/Radiosity_\(computer_graphics\)](https://en.wikipedia.org/wiki/Radiosity_(computer_graphics))

- Radiosity solves simplified rendering equation only for **diffuse** surfaces
 - Instead of radiance, use radiosity that doesn't depend on direction at each surface point (due to diffuse-only reflection)
- Then: using neural networks
 - Works for general BRDFs

Radiosity

(1984) ([https://en.wikipedia.org/wiki/Radiosity_\(computer_graphics\)](https://en.wikipedia.org/wiki/Radiosity_(computer_graphics)))

- Rendering equation restricted to diffuse surfaces (radiosity B instead of radiance L , no directional dependence; diffuse BRDF ρ , scene surfaces M)

$$B(\mathbf{x}') = E(\mathbf{x}') + \rho(x') \int_M B(\mathbf{x}) G(\mathbf{x} \leftrightarrow \mathbf{x}') d\mathbf{x}, \quad \begin{matrix} \text{3-point form,} \\ \text{area integration} \end{matrix}$$

- Key idea: discretization of desired, continuous radiosity $B(x)$ using piecewise constant functions B_i over mesh elements i

$$B_i = E_i + \rho_i \sum_{j=1}^n F_{ij} B_j$$

Sum over all mesh faces approximates
integral over scene surfaces M

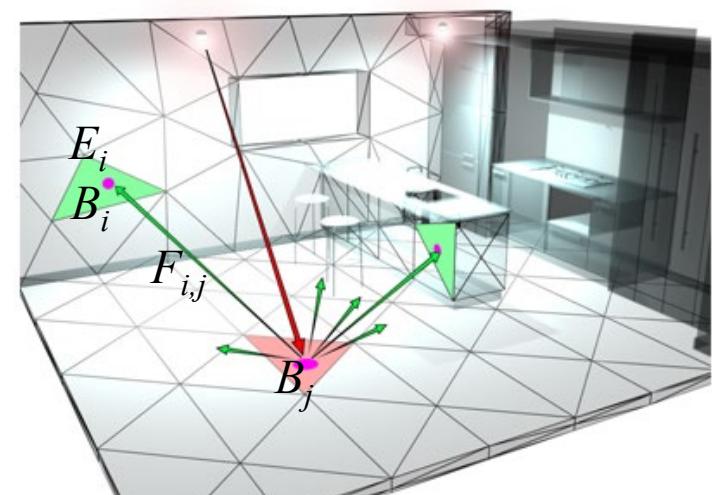
Constant radiosity
over mesh face i

Constant emission
over mesh face i

- Form factors $F_{i,j}$ between mesh faces i, j

$$F_{ij} = \frac{1}{A_i} \int_{A_i} \int_{A_j} G(\mathbf{x} \leftrightarrow \mathbf{x}') dA_j dA_i,$$

Geometry term Mesh faces



Radiosity solution

- Need to solve linear system for unknowns B_i

$$B_i = E_i + \rho_i \sum_{j=1}^n F_{ij} B_j$$

- Goal: minimize difference between left- and right-hand side to find “equilibrium”

$$\underset{\mathbf{B}}{\operatorname{argmin}} \sum_i \left\| B_i - \left(E_i + \rho_i \sum_{j=1}^n F_{ij} B_j \right) \right\|_2$$

- Solution ($n \times n$ matrix \mathbf{F} , column vectors \mathbf{B}, \mathbf{E})

$$\mathbf{B} = (\mathbf{I} - \rho \mathbf{F})^{-1} \mathbf{E}$$

Challenges

- High mesh resolution (i.e., small piecewise constant elements) required for high accuracy
 - Very large linear systems to solve
- Extending to arbitrary BRDFs, radiance fields requires 4D discretization
 - Instead of 2D mesh elements for diffuse surfaces
- Resurrect concept of radiosity using neural networks?
 - Use neural networks to represent radiance, instead of conventional approach such as piecewise constant functions over mesh

Neural radiosity

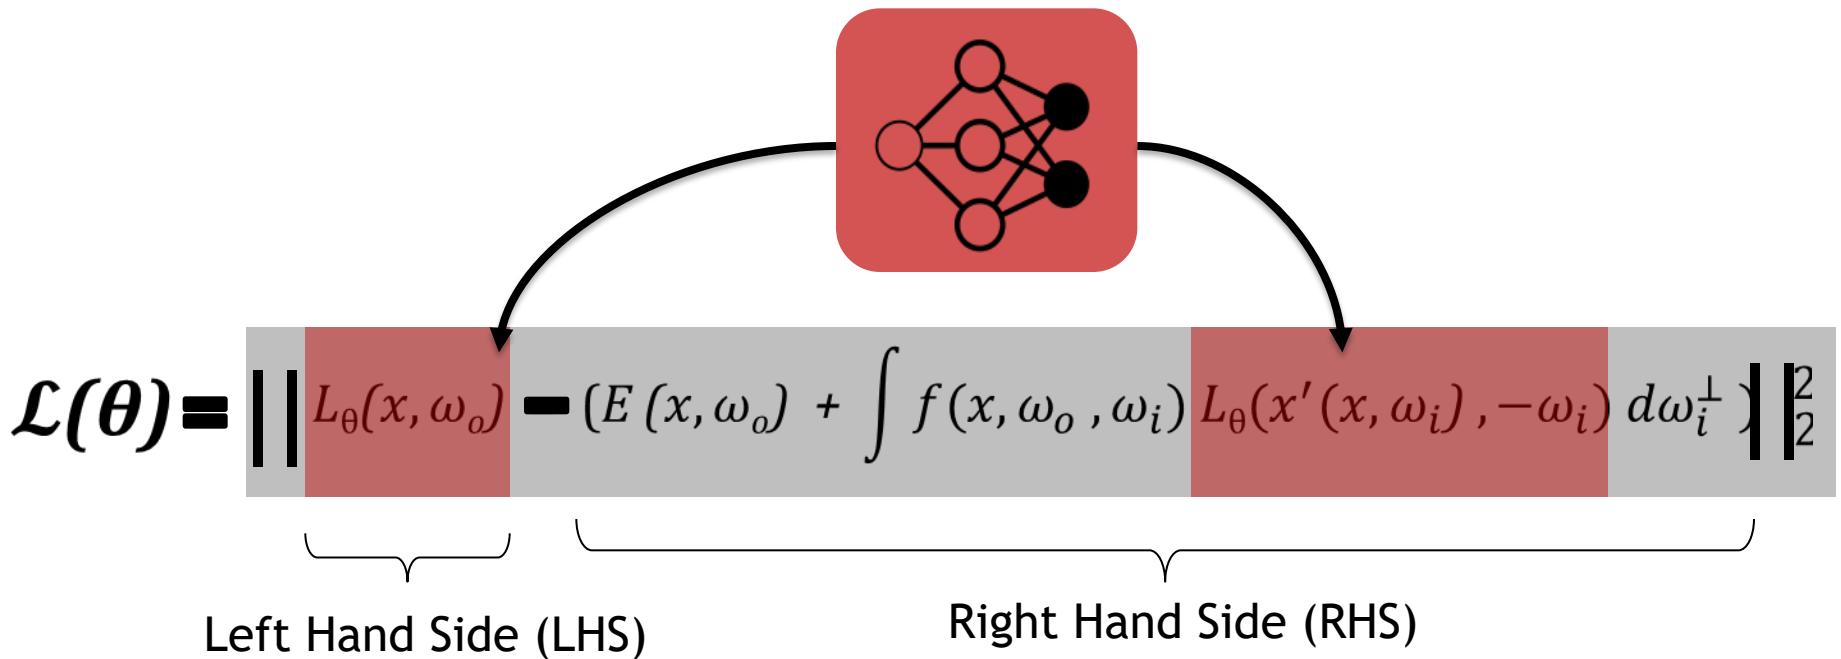
Neural network with parameters θ
to represent radiance field L_θ



$$L_\theta(x, \omega_o) \neq E(x, \omega_o) + \int f(x, \omega_o, \omega_i) L_\theta(x'(x, \omega_i), -\omega_i) d\omega_i^\perp$$



Training/optimization


$$\mathcal{L}(\theta) = \| L_\theta(x, \omega_o) - (E(x, \omega_o) + \int f(x, \omega_o, \omega_i) L_\theta(x'(x, \omega_i), -\omega_i) d\omega_i^\perp) \|_2^2$$

Left Hand Side (LHS) Right Hand Side (RHS)

Training loss L for network parameters θ
as norm of residual of rendering equation

Estimating norm of residual

- Residual: difference between LHS, RHS

$$r_\theta(x, \omega_o) = L_\theta(x, \omega_o) - (E(x, \omega_o) + \int f(x, \omega_o, \omega_i) L_\theta(x'(x, \omega_i), -\omega_i) d\omega_i^\perp)$$

- Monte Carlo sampling surface locations x_j , directions $\omega_{o,j}$

$$\mathcal{L}(\theta) = \|r_\theta(x, \omega_o)\|_2^2 = \frac{1}{N} \sum_{j=1}^N \frac{r_\theta(x_j, \omega_{o,j})}{p(x_j, \omega_{o,j})}$$

- Minimizing loss: stochastic gradient descent using batches of Monte Carlo samples

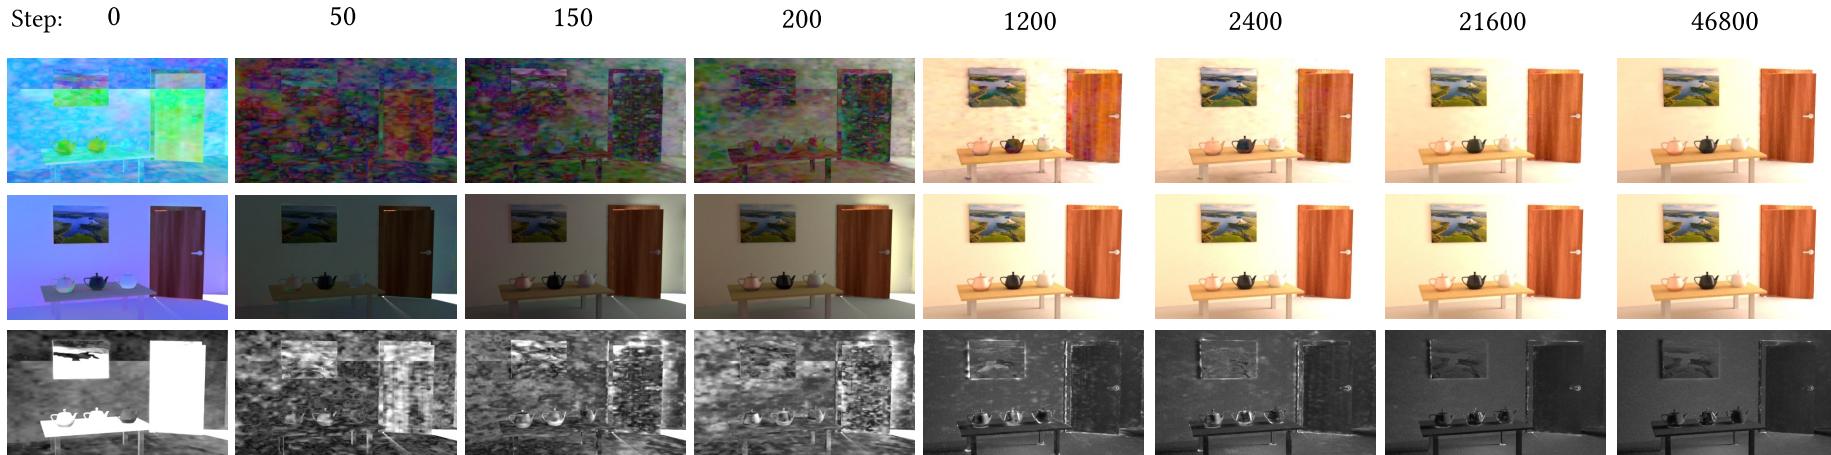
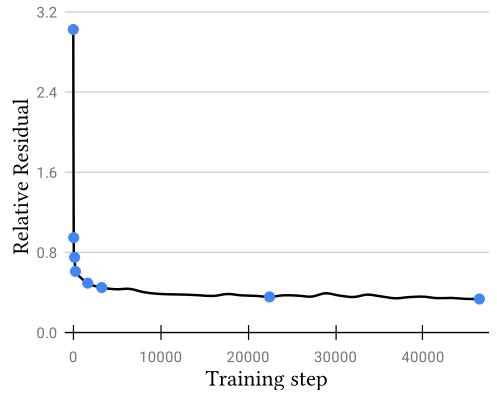
Network architecture

- Instead of naïve MLP, use sparse multi-resolution feature grid, details in paper

<https://arxiv.org/abs/2105.12319>

Training/optimization

Network needs to be trained for each scene!



As training progresses, left- and right-hand side become more similar, residual vanishes, radiance field converges to solution of rendering equation

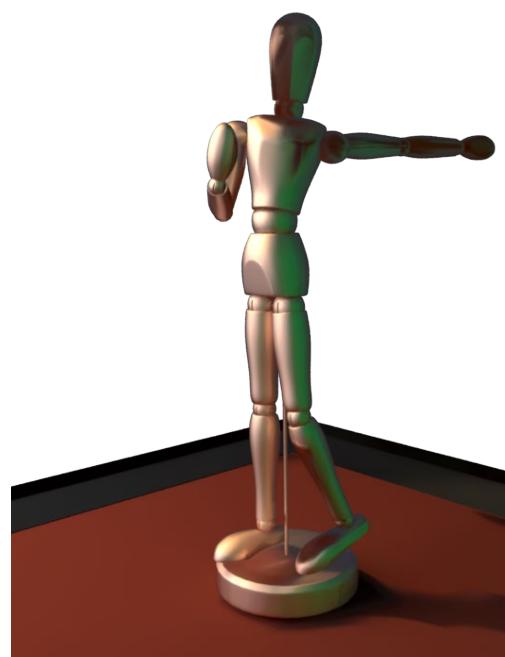
Dynamic scenes

In dynamic scenes (animations), can “re-use” and adapt network trained for initial scene to save training time

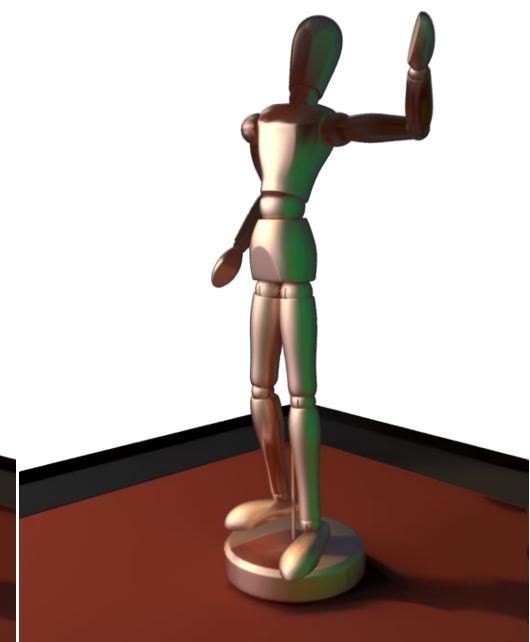


LHS - Initial scene
Rest Pose

Same
network
→



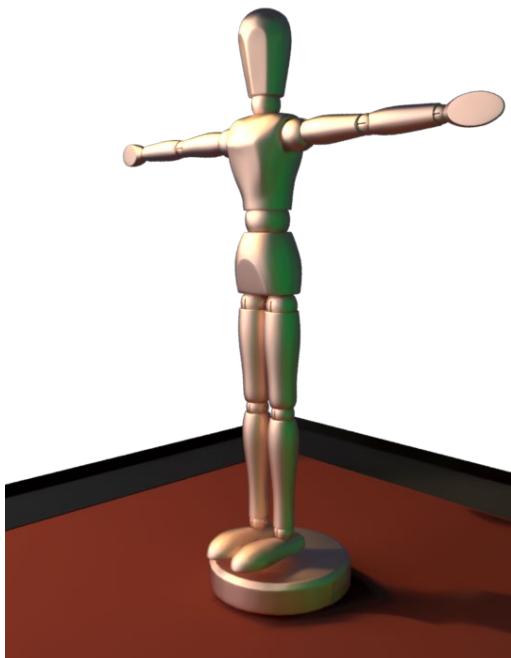
LHS - Stretch Pose



LHS - Speak Pose

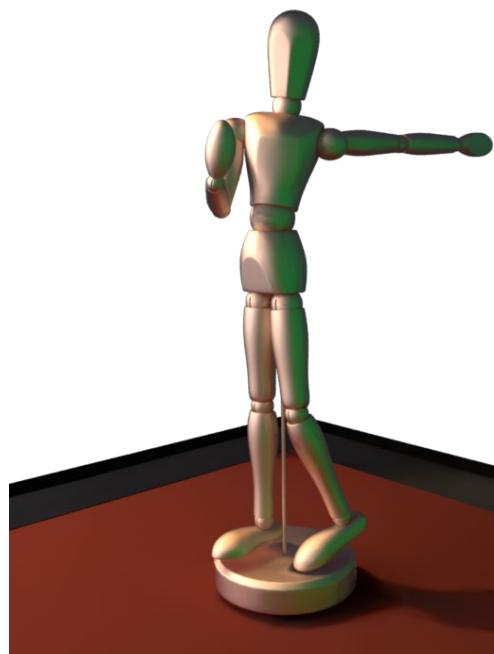
Without fine-tuning (retraining), image looks roughly ok but network doesn't represent correct solution after scene geometry changes

Dynamic scenes

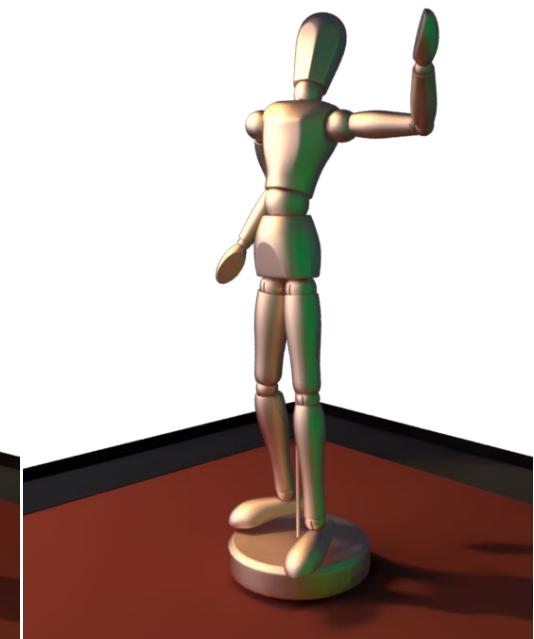


LHS - Initial scene
Rest Pose

After
fine
tuning
→



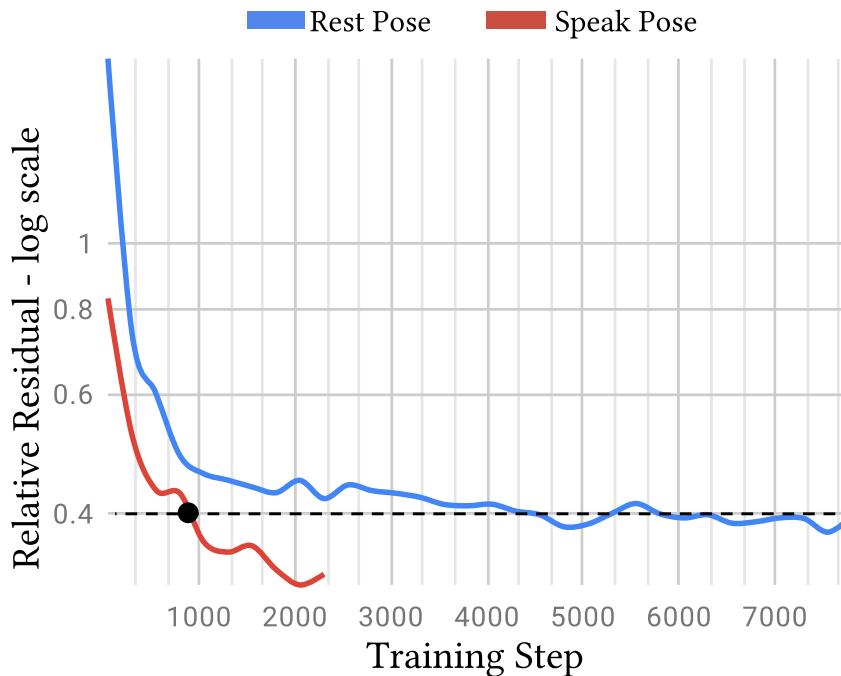
LHS - Stretch Pose



LHS - Speak Pose

Accurate radiance field after network was updated to represent new scene

Fine tuning



Fine-tuning (red line) converges faster than original training (blue line)

Related work: Real-time Neural Radiance Caching for Path Tracing

<https://arxiv.org/abs/2106.12372>

https://research.nvidia.com/publication/2021-06_real-time-neural-radiance-caching-path-tracing

Neural Radiosity Conclusions

- Neural networks can accurately represent complex radiance fields
 - Can use as representation for unknown functions of numerical optimization problems, such as rendering equation
- Neural radiosity requires per scene training, which can be slow
 - Training time can be amortized over multiple frames in dynamic scenes
- Real-time performance possible (real-time neural radiance caching)