


CMSC740

Advanced Computer Graphics

Matthias Zwicker
Fall 2025

Surface reconstruction


- Input: point cloud
- Output: watertight triangle mesh



Challenges

Input data quality

- Noise
- Holes



Input

How to assess
output quality?



Output

Quality measures

- Hausdorff distance: distance between two subsets of metric space
 - Surface: subset of all 3D points
 - “Greatest of all distances from point in one set to closest point in other set”
- Distance of point a to subset B


$$d(a, B) = \inf_{b \in B} d(a, b)$$

- Definition of Hausdorff distance

$$d_H(X, Y) = \max \left\{ \sup_{x \in X} d(x, Y), \sup_{y \in Y} d(Y, x) \right\}$$

- Efficient computation non-trivial

<https://dl.acm.org/doi/abs/10.3233/ICA-170544>



Quality measures


- Chamfer distance between discrete point sets

$$d_{CD}(X, Y) = \sum_{x \in X} \min_{y \in Y} \|x - y\|_2^2 + \sum_{y \in Y} \min_{x \in X} \|x - y\|_2^2$$

- In practice, random sampling of continuous surfaces to compute sum

Quality measures

- Preserve topology (number of holes, handles in closed surface)



Two examples today

1. Combinatorial

- Connect input data points to mesh

2. Function fitting

- Implicit function in 3D

- Many, many more variations from recent research
- «A Survey of Surface Reconstruction from Point Clouds”, Berger et al. 2016

<https://hal.inria.fr/hal-01348404v2/document>

Two examples today

1. Combinatorial

- Connect input data points to mesh

2. Function fitting

- Implicit function in 3D

Gradients

- Implicit surface defined by function
 $f(\mathbf{x}) : \mathbb{R}^3 \rightarrow \mathbb{R}$, $\mathbf{x} = (x, y, z)$, as set $\{\mathbf{x} \mid f(\mathbf{x})=C\}$


- Gradients

$$\nabla f(\mathbf{x}) = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \\ \frac{\partial f}{\partial z} \end{bmatrix}$$

- Gradients are normals of isosurfaces
 $\{\mathbf{x} \mid f(\mathbf{x})=const\}$, in particular $\{\mathbf{x} \mid f(\mathbf{x})=0\}$
for any value of *const*


Overview

- Input: points p_i , optional normals n_i




Overview

1. Find function f such that for all i : $f(\mathbf{p}_i) = 0$, **and** additional constraints, such as $\|\nabla f(\mathbf{p})\| = 1$ for all \mathbf{p} , or $\nabla f(\mathbf{p}_i) = \mathbf{n}_i$ for all i if normals \mathbf{n}_i available



Overview


2. Triangulate isosurface $\{\mathbf{x} \mid f(\mathbf{x})=0\}$



Poisson surface reconstruction

- “Screened Poisson Surface Reconstruction”, Kazhdan and Hoppe, ACM TOG 2013
- Standard approach often used in practice
- Implicit function is defined by 3D grid of function values, with trilinear interpolation
 - Implicit function represented as piecewise linear function https://en.wikipedia.org/wiki/Piecewise_linear_function
- Code, etc.

<http://www.cs.jhu.edu/~misha/Code/PoissonRecon/Version8.0/>
https://doc.cgal.org/latest/Poisson_surface_reconstruction_3/index.html




Distance function

- Implicit function with $\|\nabla f(\mathbf{p})\|=1$ for all \mathbf{p}
 - Signed distance function (SDF), if gradient is continuous
- Deep signed distance function: use neural network to represent $f(\mathbf{x})$
 - Instead of piecewise linear, or spline (piecewise polynomial) function

Deep SDF fitting

- “IGR: Implicit Geometric Regularization for Learning Shapes” (ICML 2020), <https://github.com/amosgropp/IGR>
- Reconstruct continuous SDF represented by neural network from point cloud
 - Optimization (“training”) for each point cloud, no data set required



Formulation

- Point cloud $\mathcal{X} = \{\mathbf{x}_i\}_{i \in I} \subset \mathbb{R}^3$
- Optional normals $\mathcal{N} = \{\mathbf{n}_i\}_{i \in I} \subset \mathbb{R}^3$
- Multilayer perceptron (MLP), parameters θ

$$f(\mathbf{x}; \theta) : \mathbb{R}^3 \times \mathbb{R}^m \rightarrow \mathbb{R}$$

- Loss Ensure f is distance function

$$\ell(\theta) = \ell_{\mathcal{X}}(\theta) + \lambda \mathbb{E}_{\mathbf{x}} \left(\|\nabla_{\mathbf{x}} f(\mathbf{x}; \theta)\| - 1 \right)^2$$


$$\ell_{\mathcal{X}}(\theta) = \frac{1}{|I|} \sum_{i \in I} \left(|f(\mathbf{x}_i; \theta)| + \tau \|\nabla_{\mathbf{x}} f(\mathbf{x}_i; \theta) - \mathbf{n}_i\| \right)$$

Constrain function
values at
input points

- If normals are available $\tau = 1$, else $\tau = 0$

Network architecture

- “Autodecoder”
- Input: Latent code L (learnable), 3D location
- Output: signed distance value (scalar)



FC: fully connected

Optimization




- Stochastic gradient descent with network parameters θ as variables
- Randomly sample 3D points x and data points x_i

Why does it work?

- Using a “soft penalty” to ensure unit length gradients usually doesn’t guarantee unit length gradients in solution
- Solution is not unique; why does stochastic gradient descent on network parameters find a “reasonable” solution?
- Paper shows some analysis why a planar point cloud leads to SDF that exactly represents plane (and not “wiggly” surface)

Results

- State of the art accuracy



Previous
SOTA

New
approach

Previous
SOTA

New
approach

<https://arxiv.org/pdf/2002.10099.pdf>

Results

- DGP: previous state of the art (also neural network based)
- Chamfer distance d_C (one sided d_C^\rightarrow)
- Hausdorff distance d_H (one sided d_H^\rightarrow)

	Method	Ground Truth		Scans	
		d_C	d_H	d_C^\rightarrow	d_H^\rightarrow
Anchor	DGP	0.33	8.82	0.08	2.79
	Ours	0.22	4.71	0.12	1.32
Daratech	DGP	0.2	3.14	0.04	1.89
	Ours	0.25	4.01	0.08	1.59
Dc	DGP	0.18	4.31	0.04	2.53
	Ours	0.17	2.22	0.09	2.61
Gargoyle	DGP	0.21	5.98	0.062	3.41
	Ours	0.16	3.52	0.064	0.81
Lord Quas	DGP	0.14	3.67	0.04	2.03
	Ours	0.12	1.17	0.07	0.98

<https://arxiv.org/pdf/2002.10099.pdf>

- Neural network based implicit function fitting (similar as IGR)
- Different loss function d
 - Neural network $f(c, q)$, trained over shape data set, learnable shape specific condition (latent code) c
 - Query location q_i , input point t_i closest to q_i ,


$$t'_i = q_i - f(c, q_i) \times \nabla f(c, q_i) / \|\nabla f(c, q_i)\|_2$$

$$d(\{t'_i\}, \{t_i\}) = \frac{1}{I} \sum_{i \in [1, I]} \|t'_i - t_i\|_2^2$$

Pull query q_i towards surface using direction and distance given by neural network f

Neurall-Pull

<https://github.com/mabaorui/NeuralPull>



Neural-pull

Table 1. Reconstruction comparison in terms of L2-CD ($\times 100$).

Dataset	DSDF	ATLAS	PSR	Points2Surf	IGR	Ours
ABC	8.41	4.69	2.49	1.80	0.51	0.48
FAMOUS	10.08	4.69	1.67	1.41	1.65	0.22
Mean	9.25	4.69	2.08	1.61	1.08	0.35

- CD: Chamfer distance
- DSDF, ATLAS,
- Points2Surf: neural network based
- PSR: Poisson surface reconstruction, piecewise polynomial implicit function

Triangulate isosurface



- Goal: construct triangle mesh that approximates implicit function
- Marching cubes approach

http://en.wikipedia.org/wiki/Marching_cubes


- Sample implicit function on 3D grid
- Construct triangulation of each grid cell separately

Triangulate isosurface

- Given values of f at grid vertices
- Triangulate zero-isosurface $\{\mathbf{x} \mid f(\mathbf{x})=0\}$
- Idea: visit all cells, triangulate each cell with sign changes at its vertices separately




Triangulated
cell



All cells
triangulated

Isosurface extraction in 3D


- Basic algorithm in 3D called **marching cubes**
http://en.wikipedia.org/wiki/Marching_cubes
- $2^8 = 256$ cases of possible configurations for sign changes
- Can reduce to small number of cases using symmetries (rotation, reflection)



15 basic cases for sign change configurations in marching cubes

Marching cubes

- Encode each cell as 8-bit number given by signs at vertices
- Look-up triangulation in precomputed look-up table
- Position vertices along cube edges by linear interpolation



Position
determined by
linear interpolation
of values f_0, f_1 at
neighboring cube
vertices

Two examples today

1. Combinatorial


- Connect input data points to mesh

2. Function fitting


- Implicit function in 3D

Combinatorial approaches

- Connect input points to form mesh
- Discrete number of possible connections, combinatorial problem
- Strategies?



Input




Output

Goals

- Guarantee watertight mesh, without self intersections, gaps, cracks
- Guarantee same topology (genus) as (unknown) original shape
 - Same number of holes, handles

[http://en.wikipedia.org/wiki/Genus_\(mathematics\)](http://en.wikipedia.org/wiki/Genus_(mathematics))



Topics

Background

- Voronoi diagram
- Delaunay triangulation
- Medial axis/surface
- Crust

Reconstruction algorithms

- Curve reconstruction using crust
- Extension to surfaces in 3D


Voronoi diagram

http://en.wikipedia.org/wiki/Voronoi_diagram

- Partition of space into regions given by set of points
 - Partition: non overlapping, complete coverage
- **Voronoi cell** of each input point is area closest to that input point

Voronoi vertex
(equidistant to
three or more
input points)

Voronoi edge
(equidistant to
two input points)



Voronoi cell
(convex polygon)

Input point

Voronoi diagram


Delaunay triangulation

http://en.wikipedia.org/wiki/Delaunay_triangulation

- Dual graph of Voronoi diagram

http://en.wikipedia.org/wiki/Dual_graph


- Delaunay vertex (input point) corresponding to each Voronoi cell
- Delaunay edge joining two neighboring Voronoi cells



Delaunay triangulation

Properties

- No Delaunay vertex lies inside any circumsphere of any Delaunay triangle
- Maximizes minimum angle in all triangles




Algorithms for Delaunay triangulation


- Non-trivial, in particular in higher dimensions
- Separate research area, computational geometry
https://en.wikipedia.org/wiki/Computational_geometry
- Standard implementations available
<http://www.qhull.org/>

Medial axis/surface

- Set of points with more than one closest point on curve/surface
- Can interpret as extension of Voronoi diagram to continuous curves/surfaces
- For curves in 2D, Voronoi vertices (vertices of Voronoi diagram) of dense set of sample points along curve approximate medial axis
 - Unfortunately, doesn't hold in 3D



Medial axis



Medial surface



Crust

Definition of Crust

- A graph on the sample points
- Each pair of sample points potentially forms an edge
- An edge belongs to the crust if it has circumcircle empty of
 - other sample points
 - **and** all Voronoi vertices given by the sample points

Curve reconstruction using „crust“






1. Compute Voronoi diagram of sample points
2. Compute Delaunay triangulation of sample points **and** Voronoi vertices
3. Crust is all edges of this Delaunay triangulation that contain only sample points („Voronoi“ filtering)



Delaunay triangulation of sample points and Voronoi vertices

Crust algorithm

- When does it work, when doesn't it?



- Intuitively, need „dense enough“ samples


Sampling criterion

Feature size

- Distance from point on (continuous) surface to nearest point on medial axis

r-sample

- A set of points is r -sample of surface, if distance from any point on surface to nearest sample is at most r times feature size
- „If the medial axis is close to the surface, we need to sample more densely to satisfy the r -sampling criterion“
- Medial axis is close to surface if curvature is high, or in thin structures
- R-sampling criterion takes into account both curvature and proximity of surface parts („object thickness“)



Visualization of feature size

Sampling theorem

Theorem (Amenta et al.) <http://www.cs.ucdavis.edu/~amenta/pubs/crust.pdf>


Crust of an r -sample from a smooth curve connects only adjacent points on the curve if $r \leq 0.25$

„Curve reconstruction works correctly if distance from any point on original, continuous surface to nearest sample is at most $r=0.25$ times feature size“

Extension to 3D


<http://www.cs.ucdavis.edu/~amenta/pubs/crust.pdf>

- In 3D, Voronoi vertices do not necessarily lie close to medial axis (which they do in 2D)
 - Even if surface is densely sampled
- Cannot use the exact same scheme as for curves




Observations

- Voronoi cell of blue sample is long and thin
- Roughly perpendicular to surface, extending out to medial surface
- Some Voronoi vertices do lie close to medial axis



Crust in 3D

- Instead of using all Voronoi vertices in Voronoi filtering step, for each sample use only the two Voronoi vertices farthest from the sample, instead of all Voronoi vertices
- One on each side of the surface
- Called „poles“
- First pole: farthest vertex
- Second pole: farthest vertex that forms negative dot product with first pole




Crust in 3D

<http://www.cs.ucdavis.edu/~amenta/pubs/crust.pdf>

1. Compute the Voronoi diagram of the sample points S
2. For each sample point s do:
 - (a) If s does not lie on the convex hull, let p^+ be the farthest Voronoi vertex of V_s from s . Let n^+ be the vector sp^+ .
 - (b) If s lies on the convex hull, let n^+ be the average of the outer normals of the adjacent triangles.
 - (c) Let p^- be the Voronoi vertex of V_s with negative projection on n^+ that is farthest from s .
3. Let P be the set of all poles p^+ and p^- . Compute the Delaunay triangulation of $S \cup P$.
4. Keep only those triangles for which all three vertices are sample points in S .

Voronoi filtering



s : sample point

V_s : voronoi cell of s

p^+, p^- : poles

n^+, n^- : directions from s to poles

Crust in 3D

<http://www.cs.ucdavis.edu/~amenta/pubs/crust.pdf>

Theorem (Amenta et al.)

Given an r -sample of a surface with $r < 0.06$.

Then:

- The crust of the sample **contains** a set of triangles forming a mesh topologically equivalent to the surface
- Every point on the crust lies within distance $5r^*d(p)$ of some point p on the surface, where $d(p)$ is the distance of p to the medial axis





Proof: see paper

Normal estimation and filtering

- Crust contains many small sliver triangles that are
 - Perpendicular to surface
 - Form T-junctions
- Additional filtering (removal of triangles) required
 - Based on fact that vectors to poles are nearly perpendicular to surface
- Throw away triangles whose normal differs too much from estimated normals
 - Can prove that the remaining triangles still form a subset forming a mesh topologically equivalent to the surface

Examples

<http://www.cs.ucdavis.edu/~amenta/pubs/crust.pdf>



Summary

- Combinatorial approach using Voronoi diagram, Delaunay triangulation, Crust
- Advantages
 - Guarantees correct topology under certain sampling criterion (**feature size, r-sample**)
- Disadvantages
 - Interpolating input points undesirable if input points are noisy