

CMSC740

Advanced Computer Graphics

Matthias Zwicker
Fall 2025

Scene reconstruction from images

- Inverse rendering problem: given images of static scene from different viewpoints, camera parameters; reconstruct scene parameters, such that rendered images match input
- Applications
 - Extract 3D representation from scene parameters
 - Novel view synthesis from arbitrary viewpoints (simply plug new camera parameters into f)
 - Re-lighting (required for most practical applications)
 - Material editing

Scene reconstruction from images

- Assume: have **rendering function** f

$$\text{image} = f(\text{scene}, \text{camera})$$

- Scene parameters (scene): could be array of triangle vertices, light sources, BRDF parameters, neural network-based geometry representation, etc.

Approach

- Optimize scene parameters θ to minimize rendering loss l over all input images I

$$\arg \min_{\text{scene}} \ell(\text{scene})$$

$$\ell(\text{scene}) = \sum_i \|\text{image}_i - f(\text{scene}, \text{camera}_i)\|$$

i	Input image	Rendered image
-----	----------------	-------------------

- Optimization using gradient descent

Challenges

- Suitable rendering function f
 - Need to compute gradients wrt. scene parameters

$$\frac{\partial \ell(\text{scene})}{\partial \text{scene}} = \frac{\partial \sum_i \|\text{image}_i - f(\text{scene}, \text{camera}_i)\|}{\partial \text{scene}}$$

- Automatic differentiation
- Suitable scene parameterization
 - Rendering function and scene representation need to be powerful enough to reproduce input images

NeRF

- “NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis”, Mildenhall et al, 2020 <https://arxiv.org/abs/2003.08934>
- Key ideas
 - **Volumetric** rendering function using volumetric emission and absorption
 - Volume density and emission (radiance) represented by **neural networks**
- Notes
 - Rendering function **not physically-based**, no model of light scattering on surfaces
 - Assumes mapping of input pixels to 3D rays is known (known camera location, orientation, field of view)
 - Doesn’t directly recover 3D surfaces due to volumetric rendering model
 - Enables **novel view synthesis** (rendering scene from different camera viewpoints), but not re-lighting; NeRF was first highly successful method for novel view synthesis that is based on inverse rendering

Volumetric rendering function

- Color $C(\mathbf{r})$ of ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$, origin \mathbf{o} , direction \mathbf{d}

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt$$

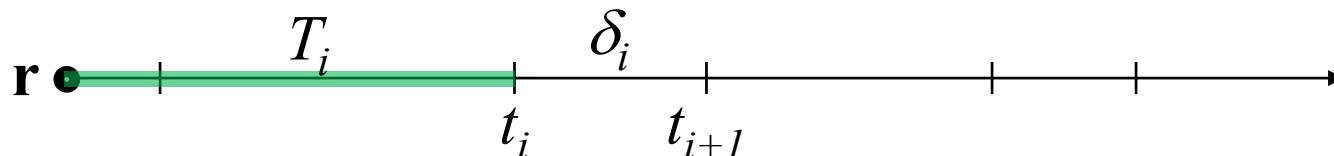
- Transmittance $T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$
- Density (scattering coefficient) $\sigma(\mathbf{r}(t))$
- Radiance (volumetric emission) $\mathbf{c}(\mathbf{r}(t), \mathbf{d})$
- Neural network to represent both σ and \mathbf{c}

Volumetric rendering function

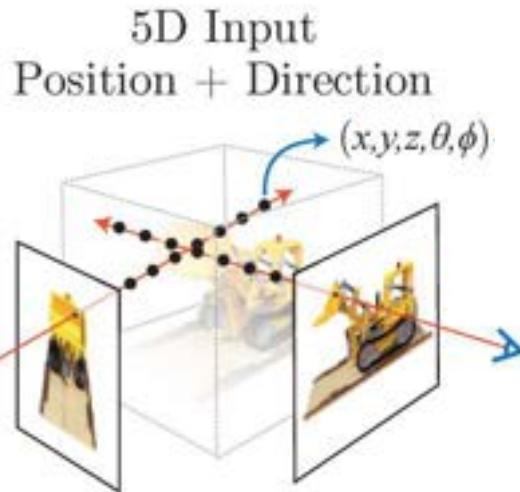
- Approximate integral using numerical quadrature (similar to ray marching)
 - N sample points along ray at locations t_i
 - Step size $\delta_i = t_{i+1} - t_i$

$$\hat{C}(\mathbf{r}) = \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i, \text{ where } T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$

Probability to reach eye Emission in segment i
 T_i $(1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i$
Probability for emission in i -th segment

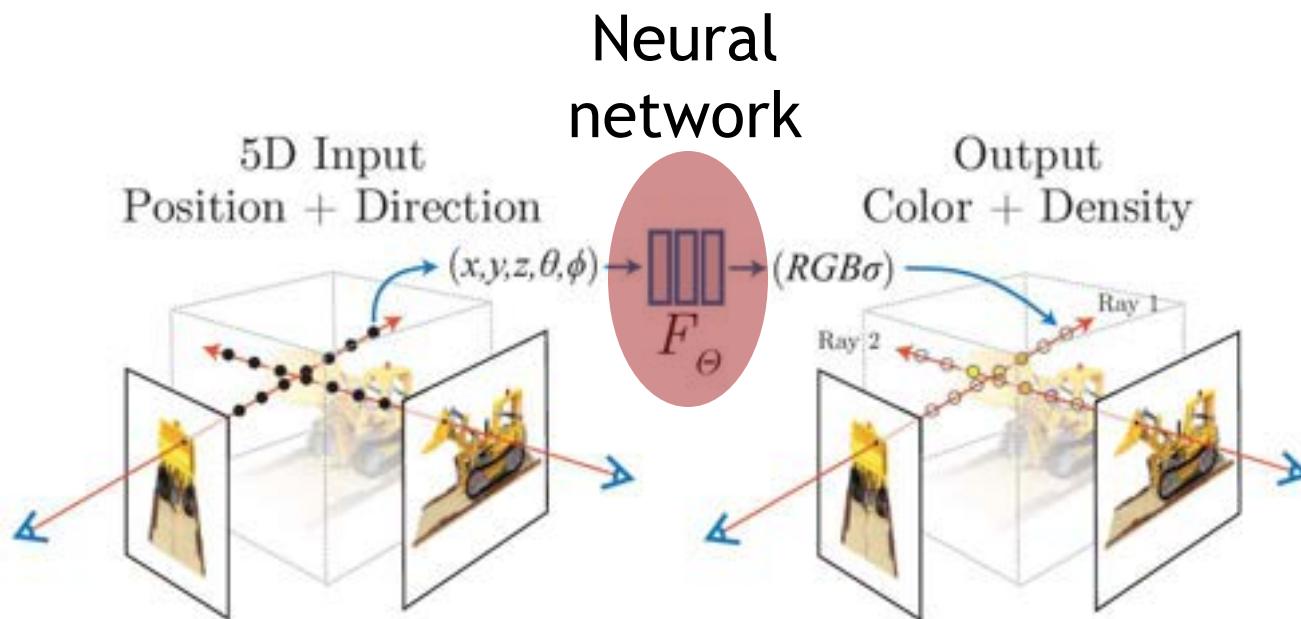


Volumetric rendering function



Ray marching, rays
defined by input images
Samples on rays with
positions, directions
 $(\mathbf{r}(t), \mathbf{d})$

Volumetric rendering function



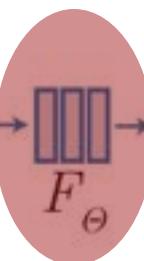
Ray marching, rays
defined by input images
Samples on rays with
positions, directions
 $(\mathbf{r}(t), \mathbf{d})$

Colors
 $\mathbf{c}(\mathbf{r}(t), \mathbf{d})$
Densities
 $\sigma(\mathbf{r}(t))$
for each
sample

Volumetric rendering function

Neural
network

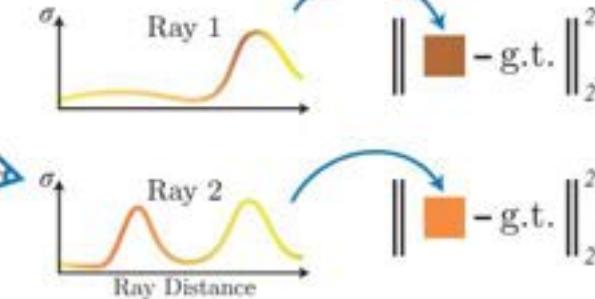
5D Input
Position + Direction



Output
Color + Density

Volume
Rendering
(Slide 7)

Rendering
Loss



Ray marching, rays
defined by input images
Samples on rays with
positions, directions
 $(\mathbf{r}(t), \mathbf{d})$

Colors
 $\mathbf{c}(\mathbf{r}(t), \mathbf{d})$
Densities
 $\sigma(\mathbf{r}(t))$
for each
sample

Optimize
colors,
densities to
minimize loss

Optimization

- Scene representation using neural network

$$F_{\Theta} : (\mathbf{x}, \mathbf{d}) \rightarrow (\mathbf{c}, \sigma)$$

At each 3d point \mathbf{x} , direction \mathbf{d} , network outputs radiance \mathbf{c} , density σ

- Loss

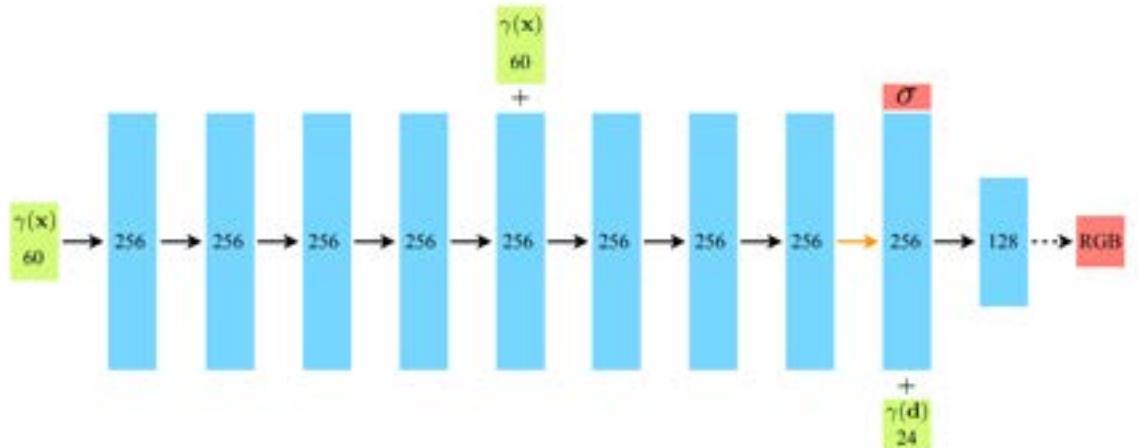
$$\mathcal{L} = \sum_{\mathbf{r} \in \mathcal{R}} \left\| \hat{C}(\mathbf{r}) - C(\mathbf{r}) \right\|_2^2$$

Sum over all Rendered Observed
rays/pixels color color in input

- Ground truth $C(\mathbf{r})$, rendered estimate $\hat{C}(\mathbf{r})$
- Gradients of pixels $C(\mathbf{r})$ colors wrt. network weights/biases Θ easy to compute using backpropagation (automatic differentiation)

Network architecture

- MLP (fully connected layers), original NeRF approach



- Positional encoding: instead of position p , network input is

$$\gamma(p) = (\sin(2^0 \pi p), \cos(2^0 \pi p), \dots, \sin(2^{L-1} \pi p), \cos(2^{L-1} \pi p))$$

- Applied to normalized coordinates p separately
- Position \mathbf{x} : $L=10$, direction \mathbf{d} : $L=4$
- Intuition: helps network to better learn high-frequency functions

Hierarchical sampling

- Train two networks simultaneously
- Coarse network: $N_c=64$ samples per ray
- Fine network: $N_f=128$ samples per ray
- Use coarse samples to define piecewise constant PDF along ray to importance sample fine samples
 - Probabilities for segments given by coarse samples

$$\hat{w}_i = w_i / \sum_{j=1}^{N_c} w_j$$

$$w_i = T_i(1 - \exp(-\sigma_i \delta_i))$$

“probability for radiance being emitted in i-th segment, and transmitted along ray all the way to eye”

Results

<https://www.matthewtancik.com/nerf>

- View synthesis
- Depth image reconstruction
- 3D surface reconstruction using marching cubes

Limitations of NeRF

- Doesn't attempt to represent surface directly (only volumetric density σ)
- Treats pixels as infinitesimal rays, doesn't take into account pixel areas
- Doesn't take into account imaging artifacts such as blur, over-/under-exposure
- Only works for static scenes
- Requires known camera parameters
- Training and rendering slow
- Requires many input images for high quality reconstruction
- Doesn't take into account potential appearance variation in input images (different illumination, time, time of year, etc.)
- Doesn't recover BRDF parameters and illumination
- Only uses 3D locations to predict scene (density, radiance); does not use correspondence information between 3D locations and images

Surface reconstruction

- How to reformulate problem to enable reconstruction of well-defined surfaces?

Surface reconstruction

- “NeuS: Learning Neural Implicit Surfaces by Volume Rendering for Multi-view Reconstruction” <https://arxiv.org/pdf/2106.10689.pdf>
- Optimize neural SDF (signed distance function) instead of density to represent surface
- But still use volumetric rendering similar to NeRF
- Challenge: relationship between SDF and density, which is required for volumetric rendering?

NeuS

- Neural SDF $f(\mathbf{p})$, 3D point \mathbf{p}
- Neural radiance field $c(\mathbf{p}, \mathbf{v})$, direction \mathbf{v}
- Rendering function (almost, but not exactly same as NeRF)

$$C(\mathbf{o}, \mathbf{v}) = \int_0^{+\infty} w(t) c(\mathbf{p}(t), \mathbf{v}) dt$$

$$w(t) = T(t)\rho(t), \text{ where } T(t) = \exp\left(-\int_0^t \rho(u) du\right)$$

Weight function

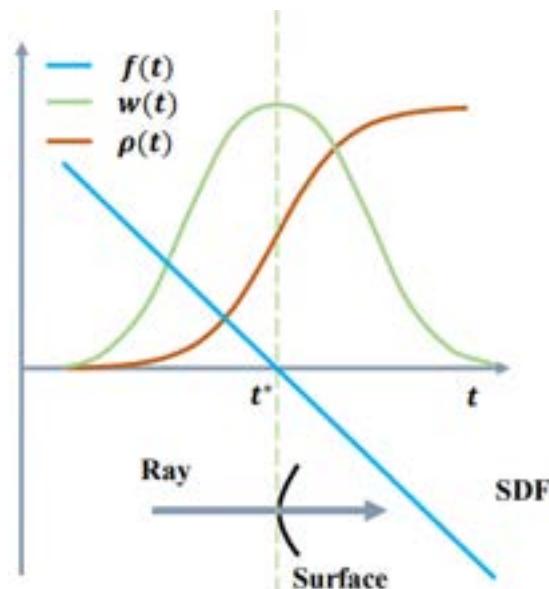
- Opaque density function ρ derived from SDF f

$$\rho(t) = \max\left(\frac{-\frac{d\Phi_s}{dt}(f(\mathbf{p}(t)))}{\Phi_s(f(\mathbf{p}(t)))}, 0\right) \quad \Phi_s(x) = (1 + e^{-sx})^{-1}$$

NeuS weight function w

Designed to be

- Unbiased: $w(t)$ has local maximum for value t when $f(\mathbf{p}(t)) = 0$
 - “Color contribution from point on surface is strongest”
- Occlusion aware: given two points $t_0 < t_1$, where $f(t_0) = f(t_1)$, then $w(t_0) > w(t_1)$



$w(t)$ is unbiased, maximum where $f=0$

Training

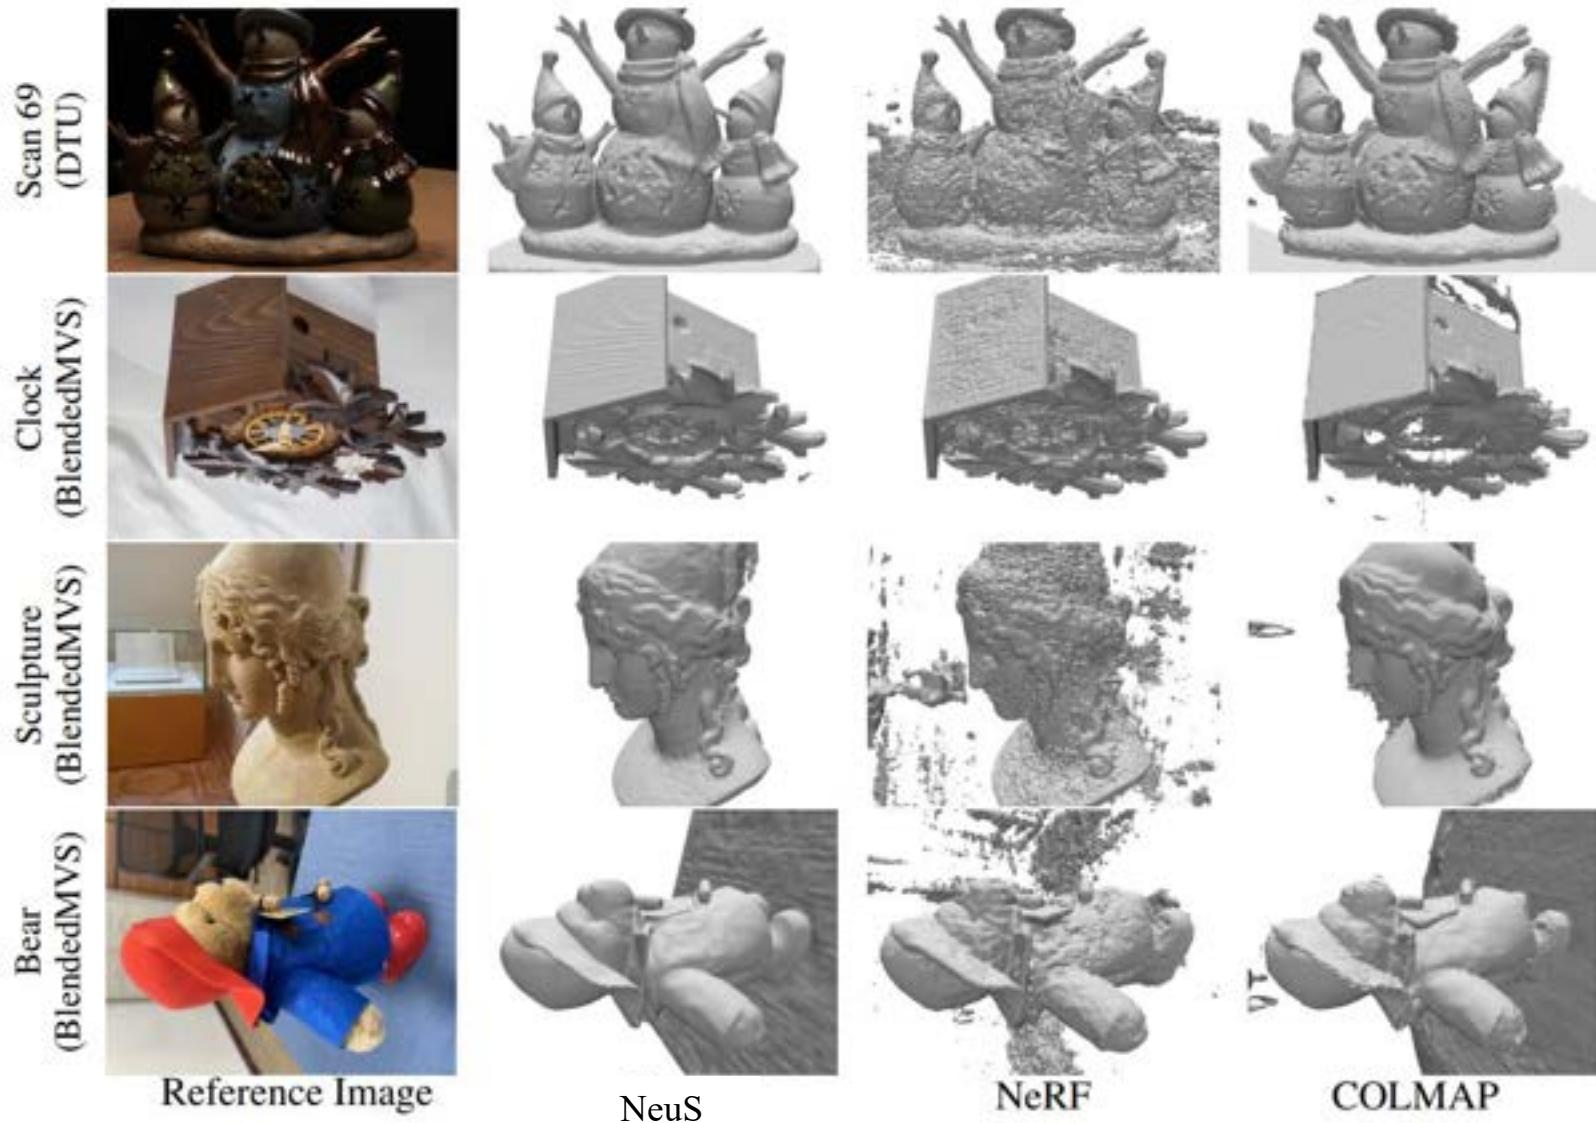
- Include Eikonal term to regularize SDF

$$\mathcal{L}_{reg} = \frac{1}{nm} \sum_{k,i} (\|\nabla f(\hat{\mathbf{p}}_{k,i})\|_2 - 1)^2$$

- Hierarchical sampling, similar to NeRF

Results

<https://lingjie0206.github.io/papers/NeuS/>



Results

- Chamfer distance to ground truth 3D scans

ScanID	w/ mask			w/o mask			
	IDR	NeRF	Ours	COLMAP	NeRF	UNISURF	Ours
scan24	1.63	1.83	0.83	0.81	1.90	1.32	1.00
scan37	1.87	2.39	0.98	2.05	1.60	1.36	1.37
scan40	0.63	1.79	0.56	0.73	1.85	1.72	0.93
scan55	0.48	0.66	0.37	1.22	0.58	0.44	0.43
scan63	1.04	1.79	1.13	1.79	2.28	1.35	1.10
scan65	0.79	1.44	0.59	1.58	1.27	0.79	0.65
scan69	0.77	1.50	0.60	1.02	1.47	0.80	0.57
scan83	1.33	1.20	1.45	3.05	1.67	1.49	1.48
scan97	1.16	1.96	0.95	1.40	2.05	1.37	1.09
scan105	0.76	1.27	0.78	2.05	1.07	0.89	0.83
scan106	0.67	1.44	0.52	1.00	0.88	0.59	0.52
scan110	0.90	2.61	1.43	1.32	2.53	1.47	1.20
scan114	0.42	1.04	0.36	0.49	1.06	0.46	0.35
scan118	0.51	1.13	0.45	0.78	1.15	0.59	0.49
scan122	0.53	0.99	0.45	1.17	0.96	0.62	0.54
mean	0.90	1.54	0.77	1.36	1.49	1.02	0.84

DTU dataset

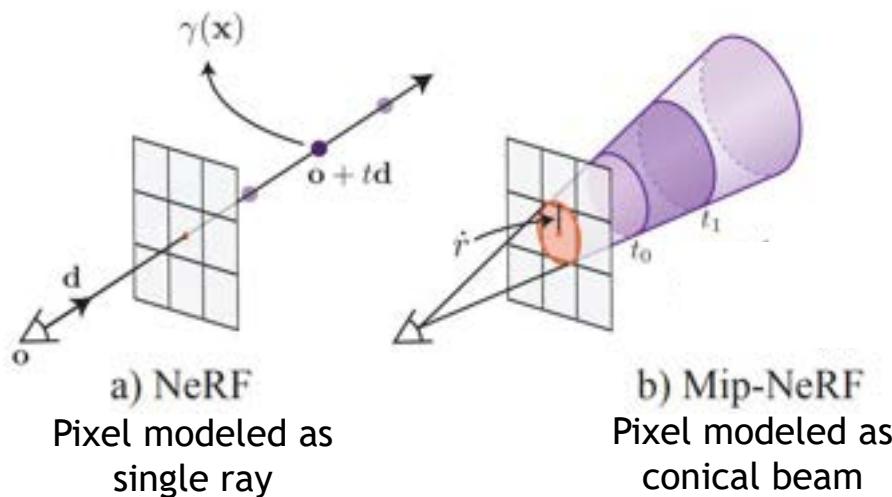
https://roboimagedata.compute.dtu.dk/?page_id=36

Pixel modeling

- How to increase reconstruction accuracy by more accurately modeling image formation in pixels of real cameras?

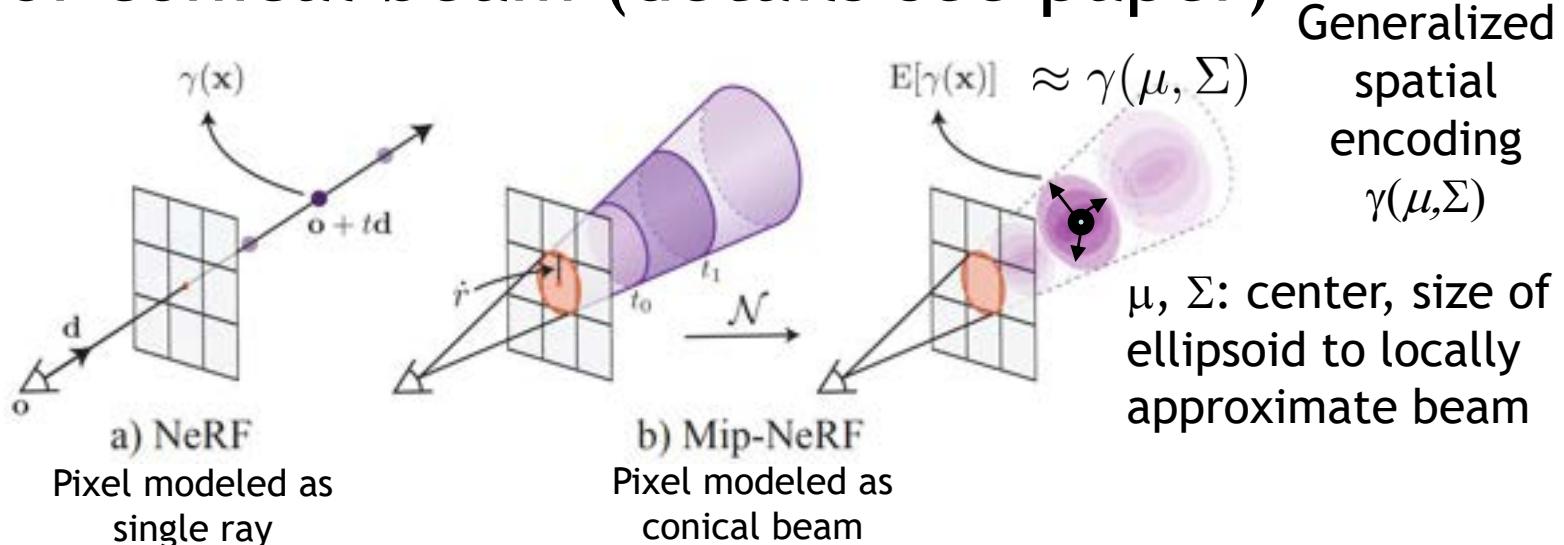
Mip-NeRF

- “Mip-NeRF: A Multiscale Representation for Anti-Aliasing Neural Radiance Fields”
<https://github.com/google/mipnerf>
- Key idea: radiance should take into account pixel size (anti-aliasing)

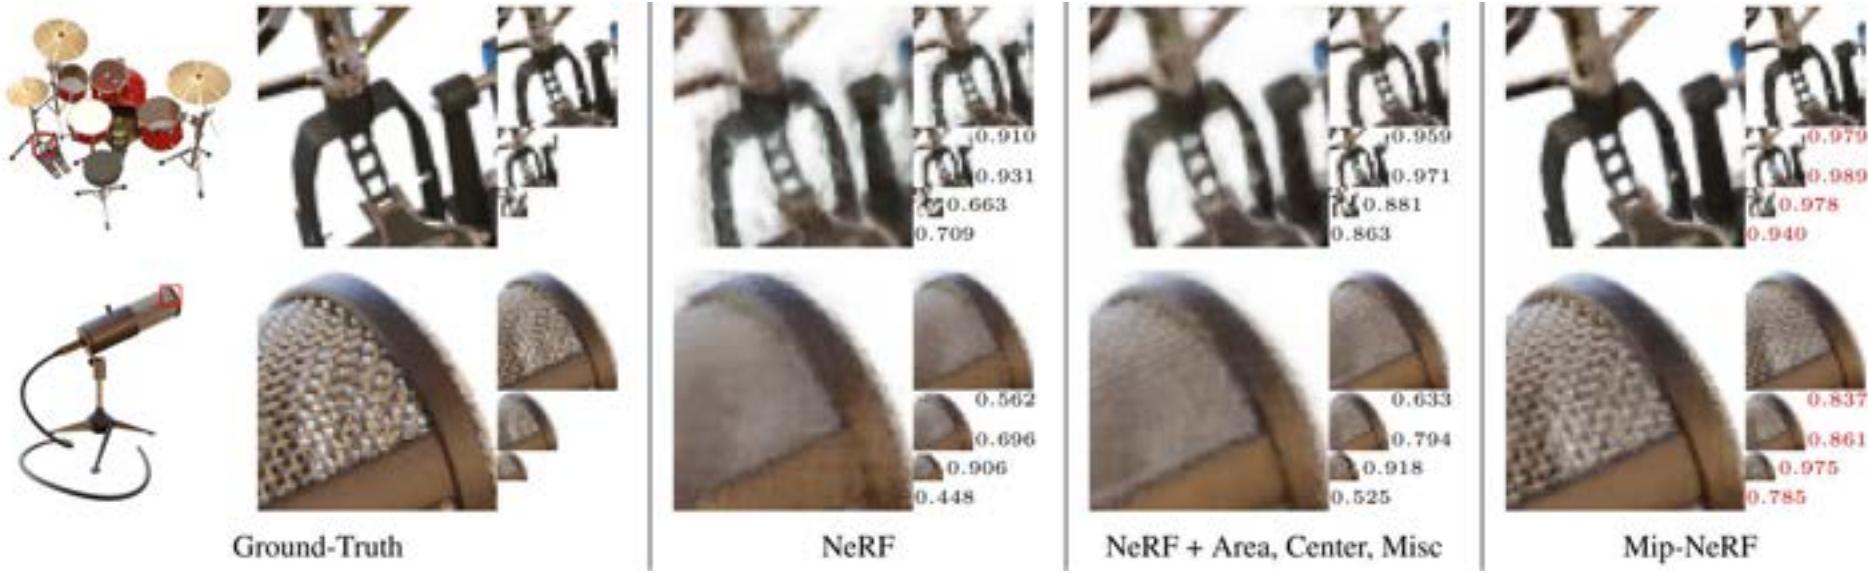


Mip-NeRF

- “Mip-NeRF: A Multiscale Representation for Anti-Aliasing Neural Radiance Fields”
<https://github.com/google/mipnerf>
- Key idea: radiance should take into account pixel size (anti-aliasing)
- Extend spatial encoding $\gamma(\mathbf{x})$ to capture local size of conical beam (details see paper)



Results



	PSNR ↑				SSIM ↑				LPIPS ↓					
	Full Res.	1/2 Res.	1/4 Res.	1/s Res.	Full Res.	1/2 Res.	1/4 Res.	1/s Res.	Full Res.	1/2 Res.	1/4 Res.	1/s Res.	Avg. ↓	Time (hours) # Params
NeRF (Jax Impl.) [11, 30]	31.196	30.647	26.252	22.533	0.9498	0.9560	0.9299	0.8709	0.0546	0.0342	0.0428	0.0750	0.0288	3.05 ± 0.04 1,191K
NeRF + Area Loss	27.224	29.578	29.445	25.039	0.9113	0.9394	0.9524	0.9176	0.1041	0.0677	0.0406	0.0469	0.0305	3.03 ± 0.03 1,191K
NeRF + Area, Centered Pixels	29.893	32.118	33.399	29.463	0.9376	0.9590	0.9728	0.9620	0.0747	0.0405	0.0245	0.0398	0.0191	3.02 ± 0.05 1,191K
NeRF + Area, Center, Misc.	29.900	32.127	33.404	29.470	0.9378	0.9592	0.9730	0.9622	0.0743	0.0402	0.0243	0.0394	0.0190	2.94 ± 0.02 1,191K
Mip-NeRF	32.629	34.336	35.471	35.602	0.9579	0.9703	0.9786	0.9833	0.0469	0.0260	0.0168	0.0120	0.0114	2.84 ± 0.01 612K
Mip-NeRF w/o Misc.	32.610	34.333	35.497	35.638	0.9577	0.9703	0.9787	0.9834	0.0470	0.0259	0.0167	0.0120	0.0114	2.82 ± 0.03 612K
Mip-NeRF w/o Single MLP	32.401	34.131	35.462	35.967	0.9566	0.9693	0.9780	0.9834	0.0479	0.0268	0.0169	0.0116	0.0115	3.40 ± 0.01 1,191K
Mip-NeRF w/o Area Loss	33.059	34.280	33.866	30.714	0.9605	0.9704	0.9747	0.9679	0.0427	0.0256	0.0213	0.0308	0.0139	2.82 ± 0.01 612K
Mip-NeRF w/o IPE	29.876	32.160	33.679	29.647	0.9384	0.9602	0.9742	0.9633	0.0742	0.0393	0.0226	0.0378	0.0186	2.79 ± 0.01 612K

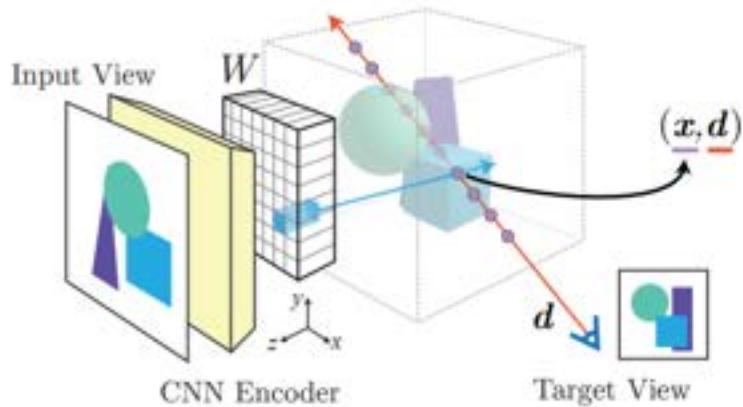
Leveraging correspondence for sparse input data

- Observation
 - To perform 3D reconstruction using correspondence and triangulation, we only need correspondence between **two views**
 - In contrast, inverse rendering typically requires observation of same scene point from **many viewpoints**
- How to leverage correspondences between 3D locations and 2D image regions to improve performance of inverse rendering under sparse input views?

Sparse input views

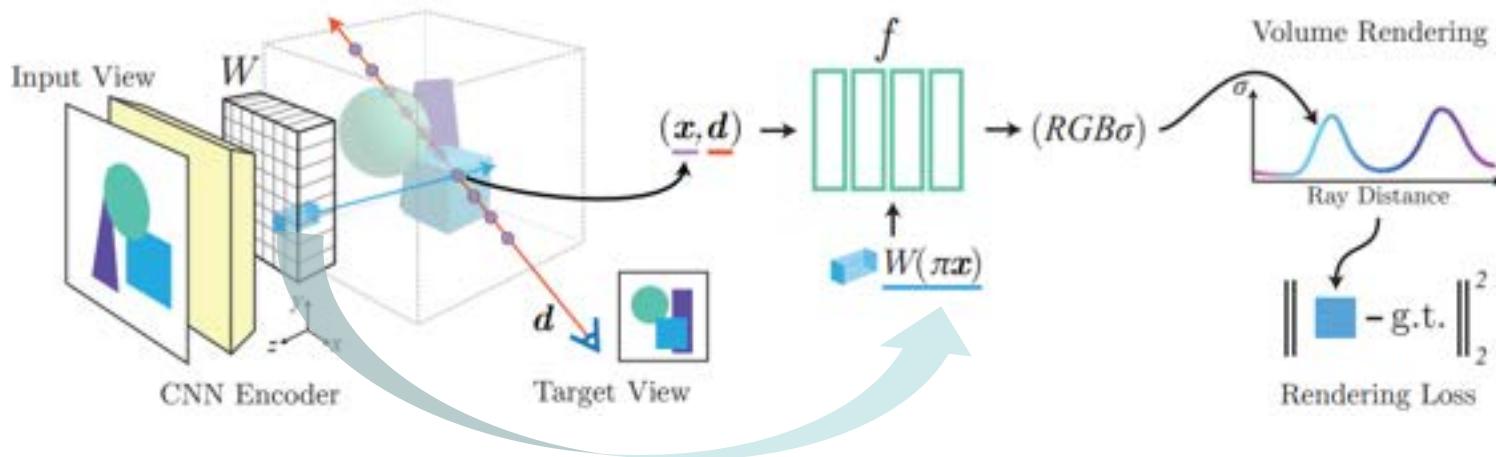
- “pixelNeRF: Neural Radiance Fields from One or Few Images” <https://arxiv.org/pdf/2012.02190.pdf>
- Key idea:
 - Project 3D sample points on rays to retrieve correspondence information between 3D sample points and images, encoded as neural network features
 - Use local image features to help NeRF network predict density, radiance

Architecture



- 2D feature maps W using CNN encoder

Architecture



- 2D feature maps \mathbf{W} using CNN encoder
- NeRF f
$$f(\gamma(\mathbf{x}), \mathbf{d}; \mathbf{W}(\pi(\mathbf{x}))) = (\sigma, \mathbf{c})$$
 - Projection of 3D sample location \mathbf{x} onto image plane $\pi(\mathbf{x})$
 - Positional encoding γ
 - Direction \mathbf{d} , density σ , radiance \mathbf{c}

Multiple input views

- Intermediate feature for each view i using network f_1

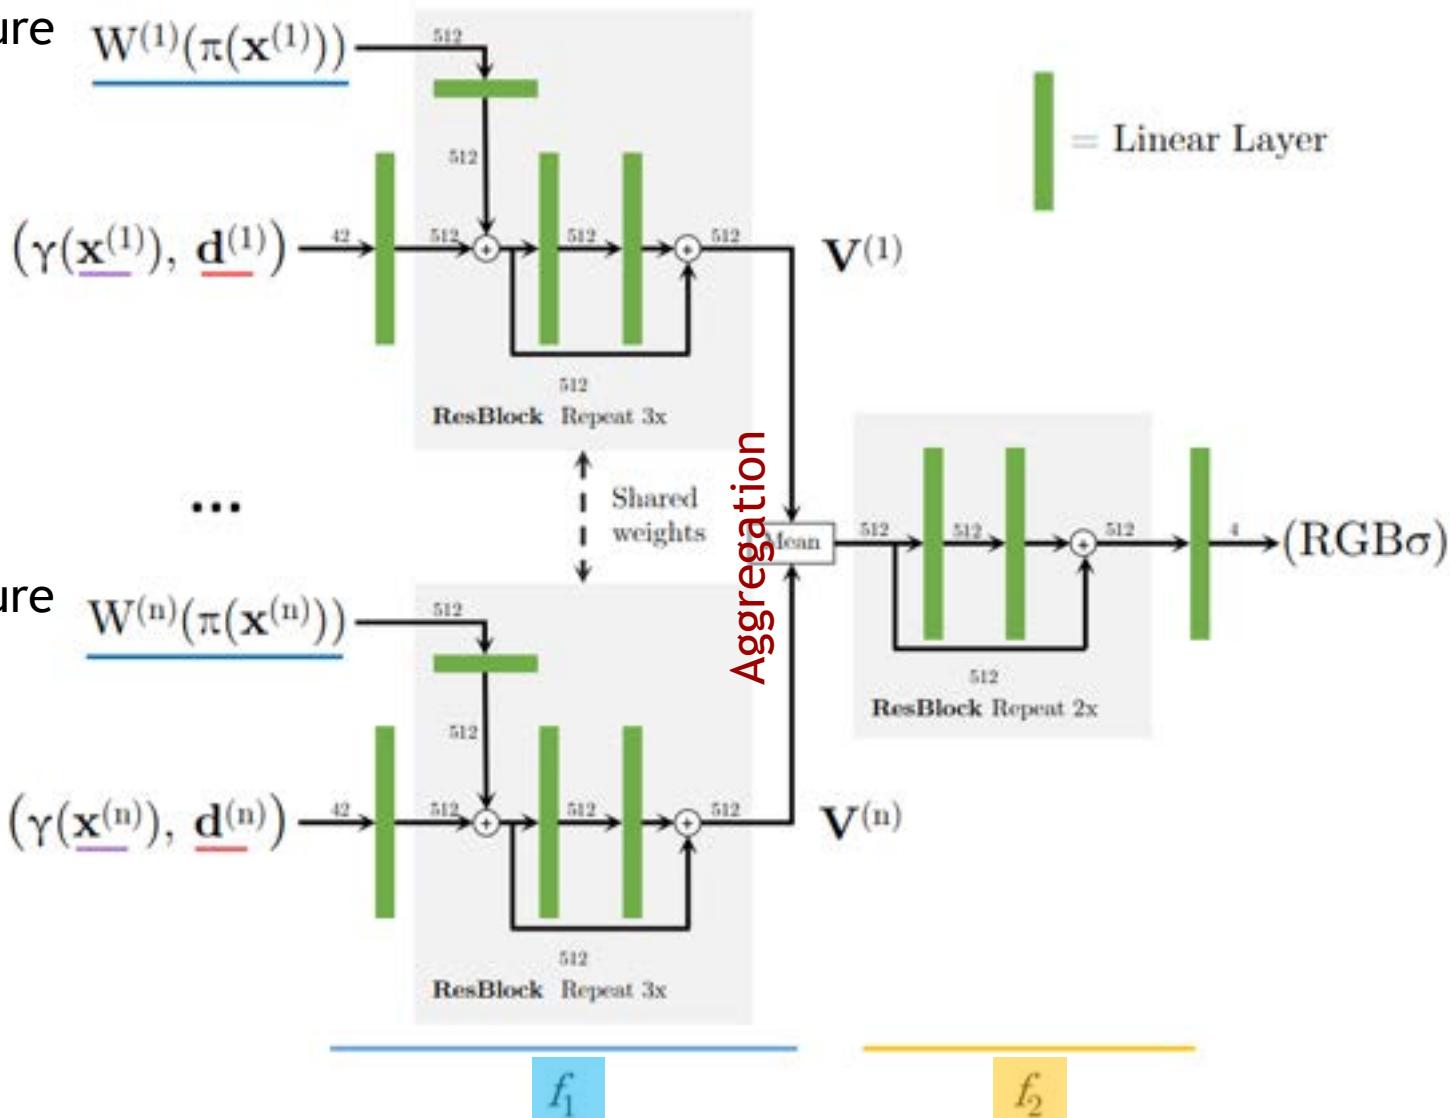
$$\mathbf{V}^{(i)} = f_1 \left(\gamma(\mathbf{x}^{(i)}), \mathbf{d}^{(i)}; \mathbf{W}^{(i)}(\pi(\mathbf{x}^{(i)})) \right)$$

- Aggregation (averaging ψ) over multiple views and color/density prediction using network f_2

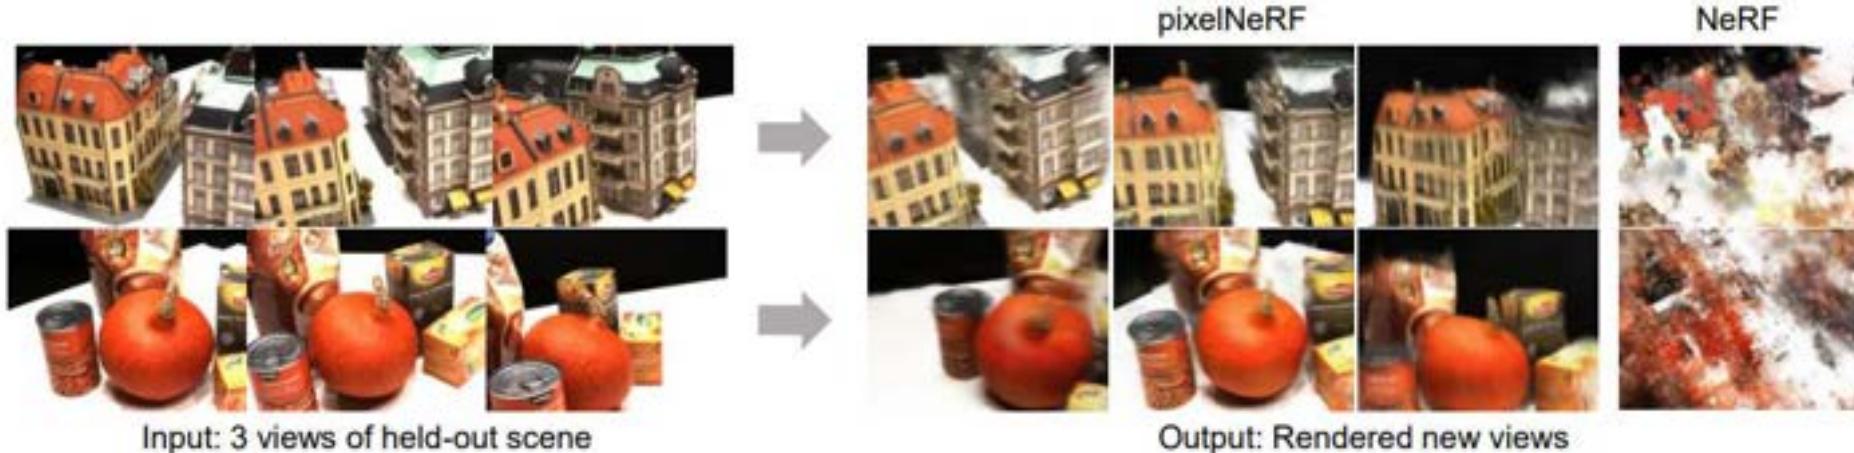
$$(\sigma, \mathbf{c}) = f_2 \left(\psi \left(\mathbf{V}^{(1)}, \dots, \mathbf{V}^{(n)} \right) \right)$$

Multiple input views

Pixel feature in view 1



Results



- Trained on data set, applied to test scene without scene-specific training

