

CMSC740

Advanced Computer Graphics

Matthias Zwicker
Fall 2025

Rendering avatars

- Goal
 - Construct avatar model (virtual human) from videos (using as little data as possible) of real person
 - Render avatars from any camera viewpoint, in arbitrary body pose
- Applications
 - Games, movies, AR/VR/XR
- Approach: deformable NeRF, custom deformation model to approximate humans

Approach

- Goal: Construct rendering function f for avatars
 $\text{image} = f(\text{camera parameters}, \text{body pose}, \text{body shape parameters}, \text{body appearance parameters})$
- Training: **inverse rendering**, optimize input parameters (camera parameters, body pose, body shape parameters, body appearance parameters) to match ground truth image data (videos of moving person, from one or multiple viewpoints)
- Rendering: given trained shape, body appearance parameters, evaluate $f(\text{render})$ for **arbitrary new camera parameters, body poses**

Approximate geometry

- Goal: take advantage of known, rough shape of humans
- Construct approximate 3D geometry using body pose, body shape parameters
- Rendering function using approximate geometry

$\text{image} = f(\text{camera parameters}, \mathbf{\textit{geometry}}, \text{body appearance parameters})$

$\mathbf{\textit{geometry}} = \mathbf{\textcolor{red}{g}}(\text{body pose, body shape parameters})$


- Two components
 - SMPL model for geometry $\mathbf{\textcolor{red}{g}}$
 - Extended NeRF for appearance f

Recent work

- HVTR: Hybrid Volumetric-Textural Rendering for Human Avatars, 3DV 2022
<https://www.cs.umd.edu/~taohu/hvtr/>
- Neural actor: neural free-view synthesis of human actors with pose control, ACM TOG 2021
<https://vcai.mpi-inf.mpg.de/projects/NeuralActor/>
- And many more...

Approximate geometry: SMPL


- SMPL: A Skinned Multi-Person Linear Model, ACM TOG 2015 <https://smpl.is.tue.mpg.de/>
- Parametric model for pose- and body shape-dependent geometry, represented as mesh
geometry = SMPL(body pose, body shape parameters)



SMPL geometries (meshes) for different pose, shape parameters

SMPL


- Based on idea of skeletal animation http://en.wikipedia.org/wiki/Skeletal_animation
- Equip (“rig”) manually created template mesh with **skeleton**
 - Joints connected to each other via rigid parts (“bones”),
https://en.wikipedia.org/wiki/Kinematic_chain
- Drive mesh deformation using skeleton
 - Mesh is attached to skeleton like skin
 - Skeleton pose given by set of rigid transformations, one per bone
 - Simple mathematical formulation
- Main method to produce CG character animation, also used in computer vision for character pose estimation



Meshes rigged with skeletons <http://www.geometry.caltech.edu/pubs/SZTDBG07.pdf>

Rigged template mesh

- Rigging: given template mesh in reference pose, construct skeleton and association of mesh with skeleton
- Rig consists of two parts
 - Skeleton
 - Skin weights: influence of each bone on each surface point



SMPL template mesh with joint locations, skin weights for each bone


Rigging

- Often manually
- Automatic methods exist
“Automatic rigging and animation of 3D characters”, Baran et al., <https://dl.acm.org/citation.cfm?id=1276467>
- Commercial software often specialized for human characters

<http://www.mixamo.com/>

Rigging

- Skeleton in reference pose




Skeleton constructed in Autodesk Maya

<https://www.cs.washington.edu/education/courses/458/07au/projects/project6/riggingTutorial/skeleton.htm>

Rigging

- Skin weights: visualizes for each bone how much it influences each point on mesh



Weight painting in Autodesk Maya

<http://jamesrburr.wordpress.com/2012/01/15/rigging-the-goblin-character/>

Skin deformation model


- Template mesh (surface): vertices in reference pose $\bar{\mathbf{x}}_i$
- Deformed vertices \mathbf{x}_i
- Bone weights w_{bi} of bone b for vertex i (matrix of size #bones x #vertices)
- Transformation matrix \mathbf{T}_b of bone b relative to reference pose

$$\mathbf{x}_i = \left(\sum_{b \in \text{bones}} w_{bi} \mathbf{T}_b \right) \bar{\mathbf{x}}_i$$

Locally weighted
rigid transformations
("skeletal animation",

https://en.wikipedia.org/wiki/Skeletal_animation

"linear blend skinning")



Template mesh/reference pose and
transformation of a bone

Forward kinematics

- Bone transformation matrices \mathbf{T}_b can be computed based on skeletal joint angles
https://en.wikipedia.org/wiki/Forward_kinematics
- Fix “root” joint, compute transformations relative to root step-by-step along kinematic chain using joint angles
 - Bone transformation is multiplication of transformation matrices for each steps along kinematic chain from root to bone
- Note: better to use dual quaternions rather than transformation matrices \mathbf{T}_b
 - Dual quaternions: representation of rigid motions such that weighted averaging “makes sense”, has desirable properties <https://users.cs.utah.edu/~ladislav/dq/index.html>

$$\mathbf{x}_i = \left(\sum_{b \in \text{bones}} w_{bi} \mathbf{T}_b \right) \bar{\mathbf{x}}_i$$

Replace with dual quaternion blending

So far


- Skin deformation model (linear blend skinning) for generic template body shape (mesh) provides pose-dependent geometry function

geometry = LinearBlendSkinning(body pose parameters, template mesh, bone weights)

where body pose parameters is set of joint angles

- Limitations
 - Skin deformation model cannot match detailed deformed body shapes
 - Single template mesh, no identity-specific shape

- Extensions of linear blend skinning



Template with joints,
blend weights

Identity-specific
shape details

Pose-specific
shape details

Final geometry using
linear blend skinning

Shape representation

- Template mesh unrolled into vector of xyz vertex positions $\bar{\mathbf{T}} \in \mathbb{R}^{3N}$
- Number of vertices $N = 6890$
- Skinning blend weights given by matrix $\mathcal{W} \in \mathbb{R}^{N \times K}$
 - Number of joints $K = 23$
- Identify specific and pose-dependent details given as \mathbb{R}^{3N} vectors storing 3D offset for each template vertex, called **blend shape** vectors
 - Identity specific blend shape vector $B_S \in \mathbb{R}^{3N}$
 - Pose-dependent blend shape vector $B_P \in \mathbb{R}^{3N}$

SMPL

- Extensions of linear blend skinning



Template with K joints,
blend weights W


$$\bar{\mathbf{T}} \in \mathbb{R}^{3N}$$

$$\mathcal{W} \in \mathbb{R}^{N \times K}$$

$$N = 6890, K = 23$$

SMPL

- Extensions of linear blend skinning



Template with K joints,
blend weights W

$$\bar{\mathbf{T}} \in \mathbb{R}^{3N}$$

$$\mathcal{W} \in \mathbb{R}^{N \times K}$$

$$N = 6890, K = 23$$

Identity-specific
shape details

$$\bar{\mathbf{T}} + B_S$$

SMPL

- Extensions of linear blend skinning



Template with K joints,
blend weights W

$$\bar{\mathbf{T}} \in \mathbb{R}^{3N}$$

$$\mathcal{W} \in \mathbb{R}^{N \times K}$$

$N = 6890, K=23$

Identity-specific
shape details

$$\bar{\mathbf{T}} + B_S$$


Pose-specific
shape details

$$\bar{\mathbf{T}} + B_S + B_P(\vec{\theta})$$

Vector of 3×23 $\vec{\theta}$
joint angles

SMPL

- Extensions of linear blend skinning



Template with K joints,
blend weights W

$$\bar{\mathbf{T}} \in \mathbb{R}^{3N}$$

$$\mathcal{W} \in \mathbb{R}^{N \times K}$$

$N = 6890, K=23$

Identity-specific
shape details

$$\bar{\mathbf{T}} + B_S$$

Pose-specific
shape details

$$\bar{\mathbf{T}} + B_S + B_P(\vec{\theta})$$

Vector of 3x23
joint angles $\vec{\theta}$

SMPL Training

- Goal: Given multi-view images of person, fit SMPL parameters to match input views (or even just single view)
- Problem: too many parameters, ill defined if only few input views
- Approach: use large data set of 3D scans to **pre-train** entire model first, then apply to new data
 - Split model parameters into **two sets**
 1. First set to be pre-trained **only on 3D data set of many identities and poses**, then fixed
$$\mathcal{W} \in \mathbb{R}^{N \times K} \quad \bar{\mathbf{T}}, B_P(\vec{\theta}) \in \mathbb{R}^{3N}$$
 2. Identity specific shape details B_S to be fit to training and new data

Identity-specific blend shapes B_s



- Split into pre-trained parameters and parameters that will be fit to new data
 - Linear model

$$B_S(\vec{\beta}; \mathcal{S}) = \sum_{n=1}^{|\vec{\beta}|} \beta_n \mathbf{S}_n \quad \mathbf{S}_n \in \mathbf{R}^{3N}$$

- Pre-trained “basis vectors” \mathbf{S}_n
- Shape coefficients (scalar weights) β_n , will be fit to new data
- Number of shape coefficients $|\vec{\beta}|$

Number of shape coefficients

$|\vec{\beta}|$



$|\vec{\beta}|$

DQBS: dual quaternion blend skinning

LBS: linear blend skinning

Pose-dependent blend shapes B_P

- Linear function $B_P(\vec{\theta}; \mathcal{P}) = \sum_{n=1}^{9K} (R_n(\vec{\theta}) - R_n(\vec{\theta}^*)) \mathbf{P}_n$
- Pre-trained vectors $\mathbf{P}_n \in \mathbf{R}^{3N}$ (#vertices N)
- Function R maps 23 joint angles $\vec{\theta}$ to vector consisting of 207 elements of all joint rotation matrices ($K=23$ joint matrices $\times 9$ elements)
 - R_n is n -th element in R
 - R computed using Rodrigues' formula
https://en.wikipedia.org/wiki/Rodrigues%27_rotation_formula
 - Joint angles of rest pose $\vec{\theta}^*$

Pre-training

- Using data set of thousands of 3D scans, registered to template mesh topology
 - Same mesh structure as template for all scans, 1-1 vertex correspondence
- Objective: optimize SMPL parameters (template mesh vertices T , blend weights W , identity blend shapes S_n , shape coefficients β_n , pose dependent blend shapes P_n) to minimize vertex registration error
 - Loss for each 3D scan: each vertex of SMPL model needs to match corresponding vertex in 3D scan
- Details of optimization see paper <https://smpl.is.tue.mpg.de/>

Note

- Description here omitted pose dependent joint locations (details see paper)

Applications

- Character animation for graphics (plugins available for many standard graphics tools, Maya, Blender, Unreal engine)

<https://smpl.is.tue.mpg.de/>

- Motion capture by fitting SMPL model to multiple camera views
- Fitting SMPL model to single view video using keypoints

<https://github.com/zju3dv/EasyMocap>

Neural Actor

- Neural Actor: Neural Free-view Synthesis of Human Actors with Pose Control, ACM TOG 2021, <https://vcai.mpi-inf.mpg.de/projects/NeuralActor/>
- Goal: Learn rendering function for virtual human from videos of person
 - Using approximate geometry

image = f (camera parameters, *geometry*, body appearance parameters)

geometry = g (body pose, body shape parameters)

- **SMPL** model for geometry g
- Extended NeRF for rendering f

Neural actor


Learn person-specific rendering function to match ground truth views

$$f(\text{camera parameters}, \textit{geometry}, \text{body appearance parameters})$$
$$\textit{geometry} = g(\text{body pose}, \text{body shape parameters})$$

Train f using inverse rendering
Geometry given by SMPL



Neural actor




SMPL
 $geometry = g(\text{body pose, body shape parameters})$
body pose parameters inferred from ground truth
images using separate techniques, e.g.

<https://arxiv.org/pdf/2204.12484.pdf>




Neural actor

- “Inverse deformation” to canonical space



Map sample points to canonical space (fixed pose using reference joint angles) by inverting linear blend skinning transformations

NeRF in canonical space




Same idea as “Nerfies” (deformable neural radiance fields), but with specialized deformation model for human bodies


Neural actor

- “Inverse deformation” to canonical space

NeRF in canonical space



Neural actor



Texture space features to learn detail deformations,
improve NeRF rendering

Training, results, discussion

- Person-specific multi-view video data sets, 11-12 cameras, approx. 30,000 training frames
- Results <https://vcai.mpi-inf.mpg.de/projects/NeuralActor/>
- Limitations?
- Related work
 - Fast, high quality rendering
<https://taohuumd.github.io/projects/hvtrpp/>
 - Including dynamic motion
<https://github.com/TaoHuUMD/SurMo?tab=readme-ov-file>