

CMSC740

Advanced Computer Graphics

Matthias Zwicker
Fall 2025

Limitations of NeRF

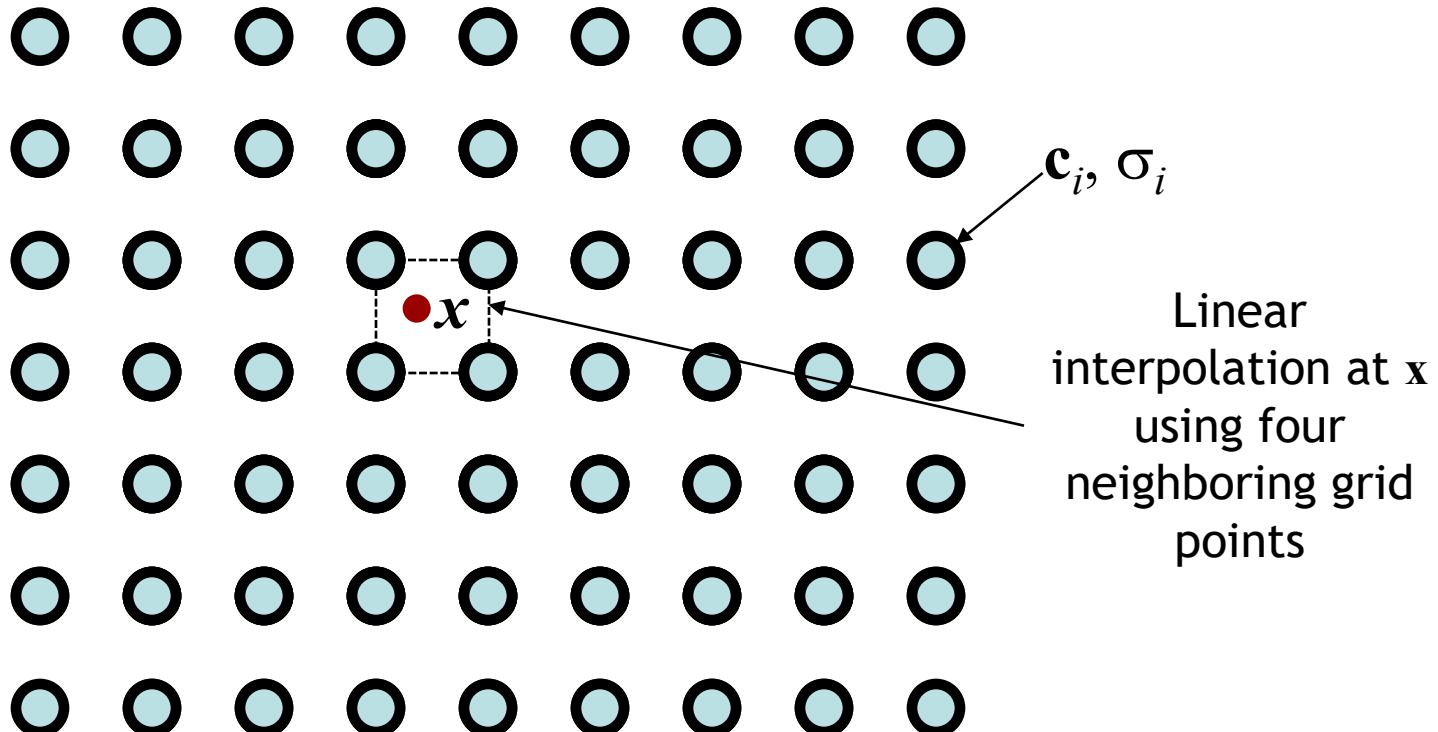
- Doesn't attempt to represent surface directly (only density σ)
 - SDF surface: NeuS, <https://lingjie0206.github.io/papers/NeuS/>
- Treats pixels as infinitesimal rays, doesn't take into account pixel areas
 - Extended spatial encoding to capture beam volume: MipNeRF, <https://jonbarron.info/mipnerf/>
- Only works for static scenes
 - Include temporal space deformation: Nerfies, <https://nerfies.github.io/>
- Doesn't take into account potential appearance variation in input images (different illumination, time, time of year, etc.)
 - Latent appearance vectors: NeRF in the wild, <https://nerf-w.github.io/>
- Doesn't recover BRDF parameters and illumination
 - Approximate model to better capture glossy reflection: RefNeRF, <https://dorverbin.github.io/refnerf/>
- Requires known camera parameters
- **Training and rendering slow**
- Requires many input images

Slow training and rendering

- Training/evaluating neural network for radiance and density is slow
 - Hundreds of network evaluations per ray, millions of rays per image
- Alternative representations for radiance, density that are faster to evaluate?
 1. Grid-based with interpolation function (e.g. linear)
 2. Radial (or elliptical) basis functions

Grid-based representation

- Parameters (to be optimized): radiance, density values at grid points c_i, σ_i
- Evaluation at arbitrary point x :
 $c(x) = \text{linear interpolation of } \{c_i \mid i \text{ in neighborhood of } x\}$

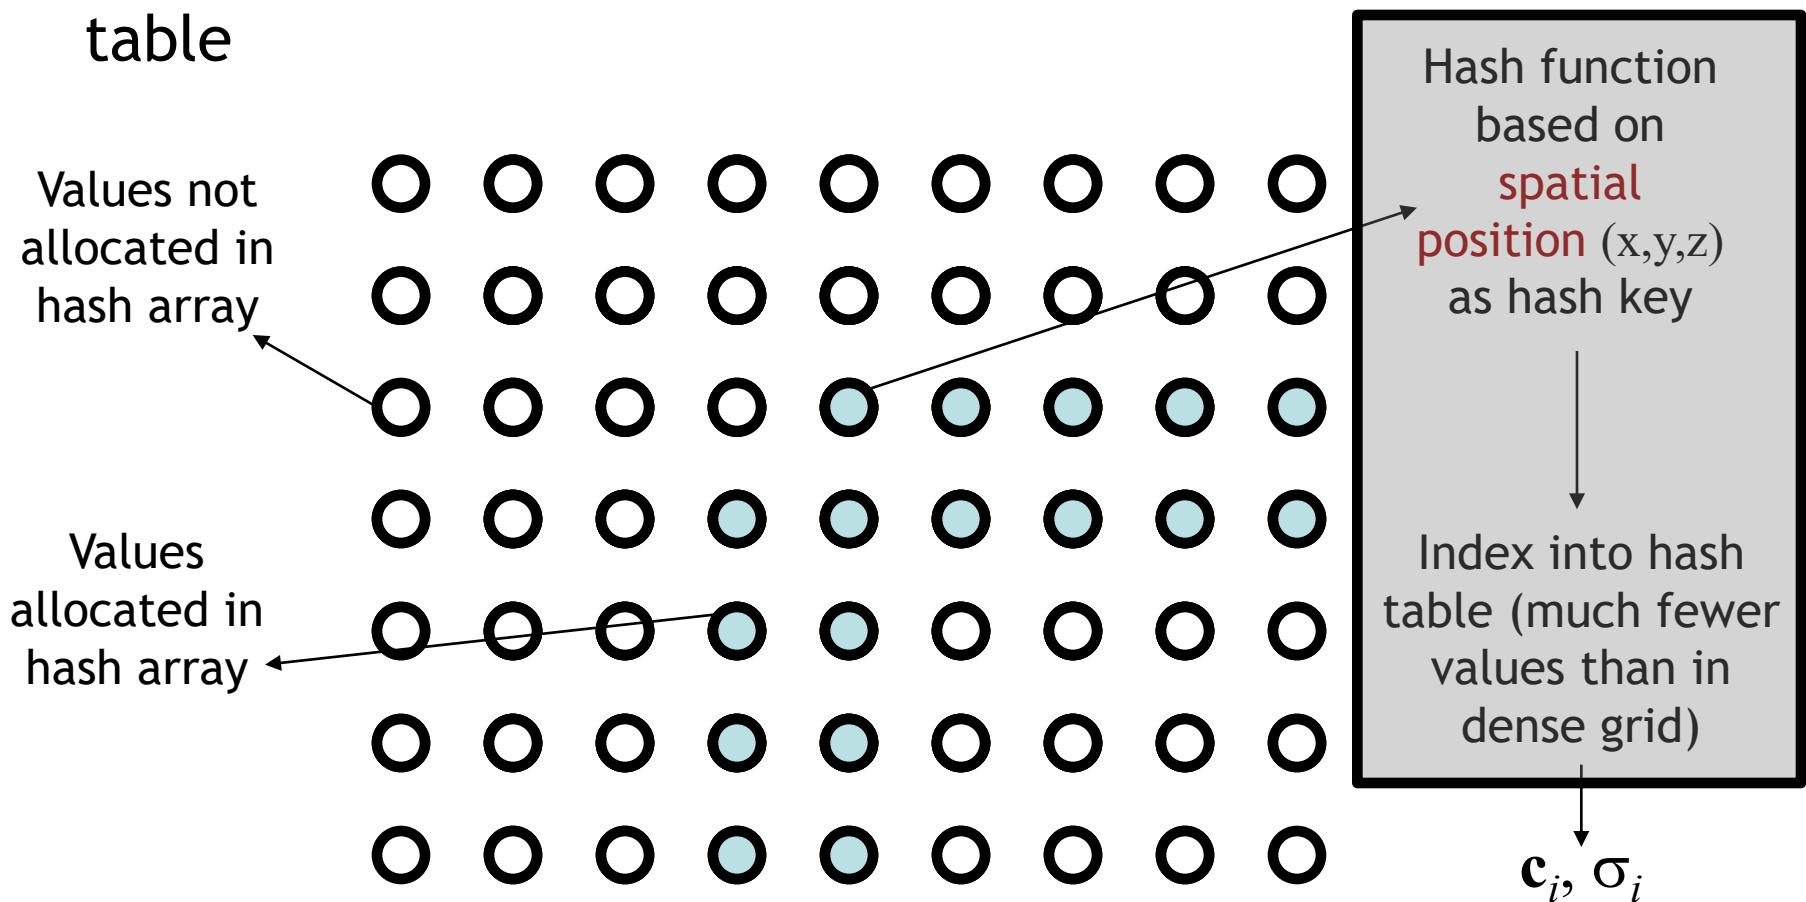


Challenges

- Dense grid requires a lot of memory
- How to pick appropriate grid resolution
- How to represent directional dependence
of radiance $c(\mathbf{x}, \mathbf{d})$, \mathbf{d} ray direction

Sparse grids using spatial hashing

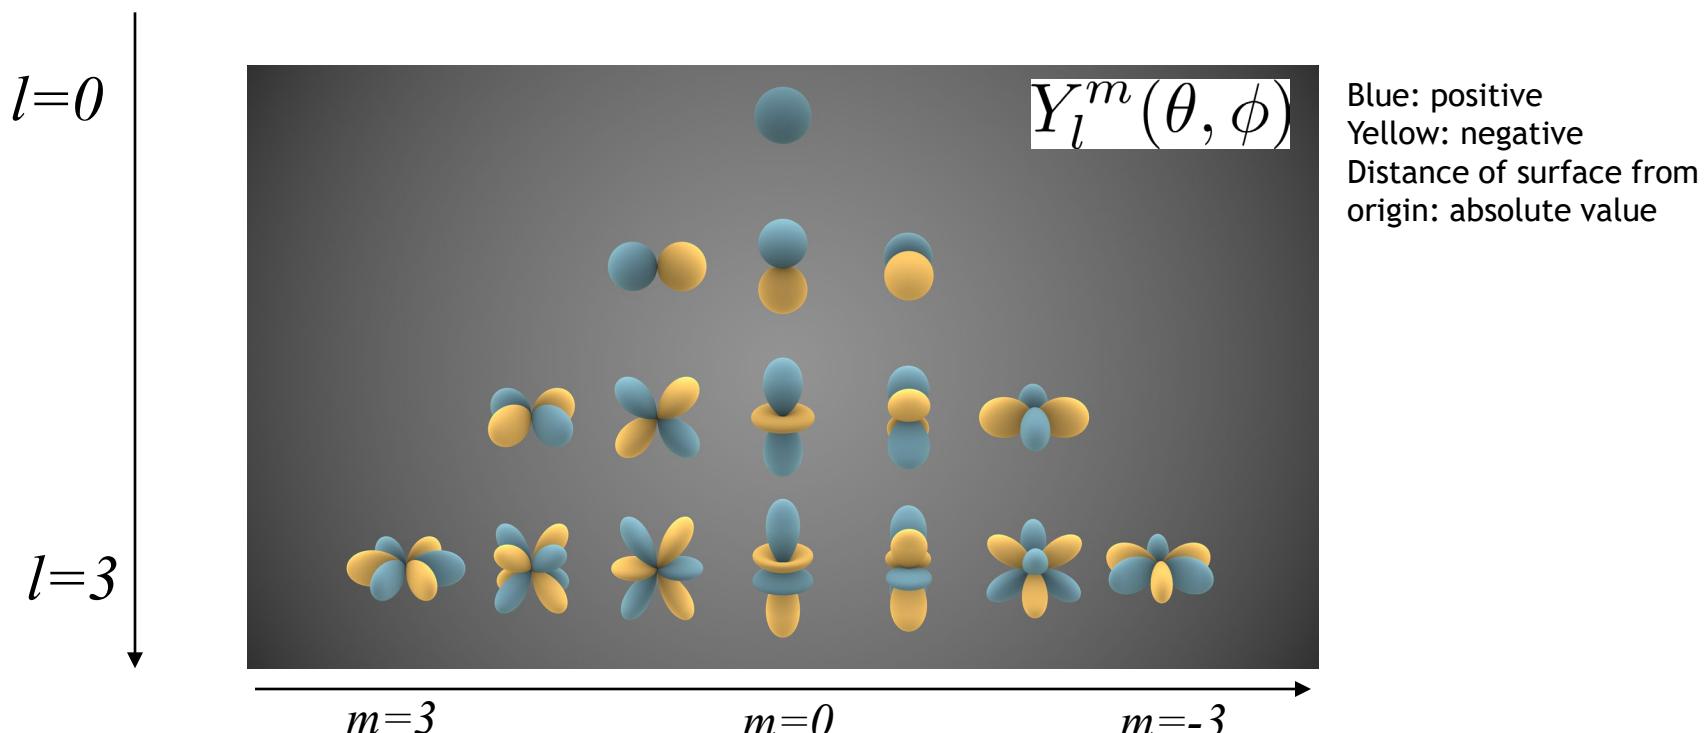
- Allocate storage for grid vertices only where needed (non-zero values)
- Spatial hashing: indirect lookup of grid value through hash table



Spatial hashing popular for many applications in computer graphics

Directional dependence: spherical harmonics

- Spherical harmonics: generalization of Fourier basis functions (sines, cosines of increasing frequencies) to sphere of directions $Y_l^m(\theta, \phi)$
 - “Frequency” l , “phase” m , spherical angles θ, ϕ



Directional dependence: spherical harmonics

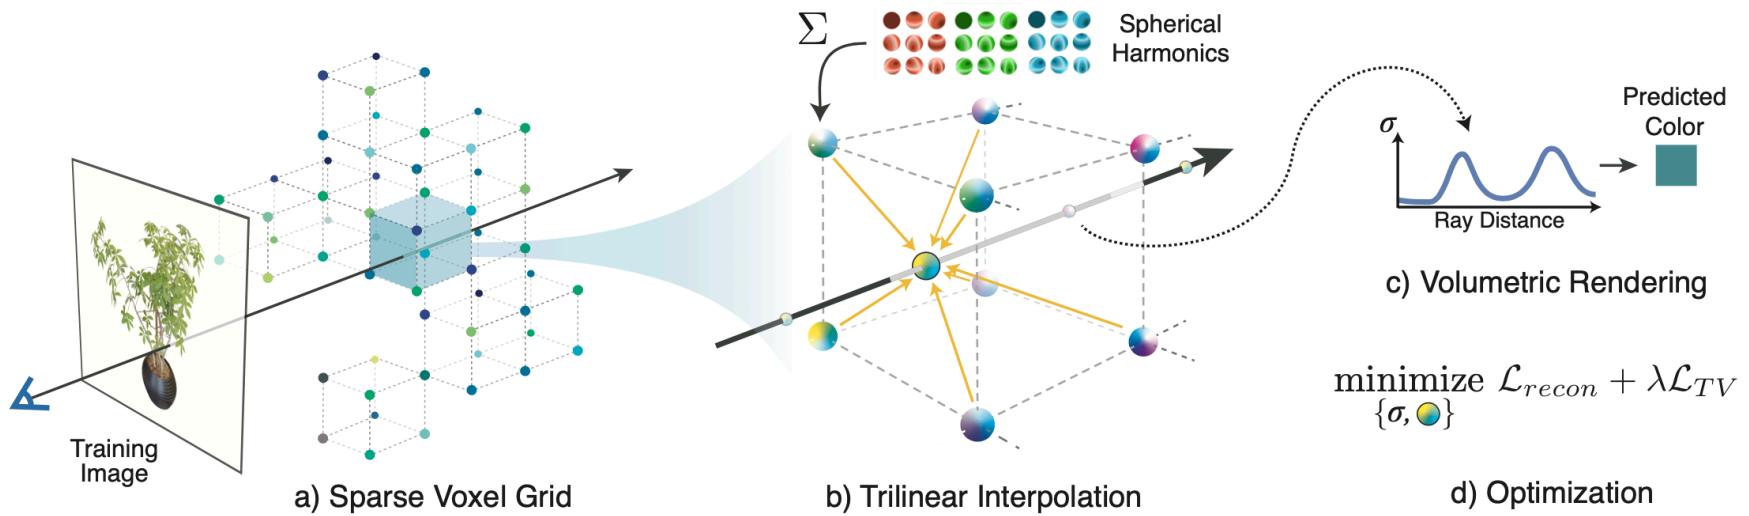
- Represent functions f (e.g. radiance) over the sphere using weighted sum of spherical harmonics
 - Weights (coefficients) f_l^m are parameters of representation
 - Using more coefficients (more spherical harmonics) enable higher frequencies

$$f(\theta, \phi) = \sum_{l=0}^k \sum_{m=-l}^l f_l^m Y_l^m(\theta, \phi)$$

- To represent functions with directional dependence, e.g. radiance $\mathbf{c}(x,y,z,\theta,\phi)$, store and optimize vector of coefficients f_l^m per grid point
 - Typically $k \leq 3$ (≤ 16 coefficients)

Plenoxels: Radiance Fields without Neural Networks

- CVPR 2022, <https://alexyu.net/plenoxels/>



	PSNR ↑	SSIM ↑	LPIPS ↓	Train Time
Ours	31.71	0.958	0.049	11 mins
NV [20]	26.05	0.893	0.160	>1 day
JAXNeRF [7, 26]	31.85	0.954	0.072	1.45 days
Ours	26.29	0.839	0.210	24 mins
LLFF [25]	24.13	0.798	0.212	—*
JAXNeRF [7, 26]	26.71	0.820	0.235	1.62 days
Ours	20.40	0.696	0.420	27 mins
NeRF++ [57]	20.49	0.648	0.478	~4 days

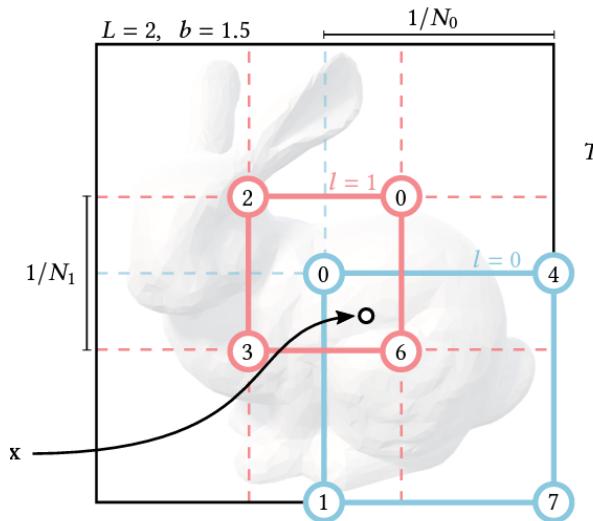
Hybrid: grid/neural network

- Store learnable feature vectors on grid
(instead of radiance or density)
- Use small neural network that takes feature vectors as input to produce desired output quantity (radiance, density, etc.)
- Potential advantages
 - Fast training due to smaller network
 - More accuracy due to feature vectors stored on spatial grid

Hybrid: grid/neural network

- Instant Neural Graphics Primitives (NGP)

<https://nvlabs.github.io/instant-ngp/>



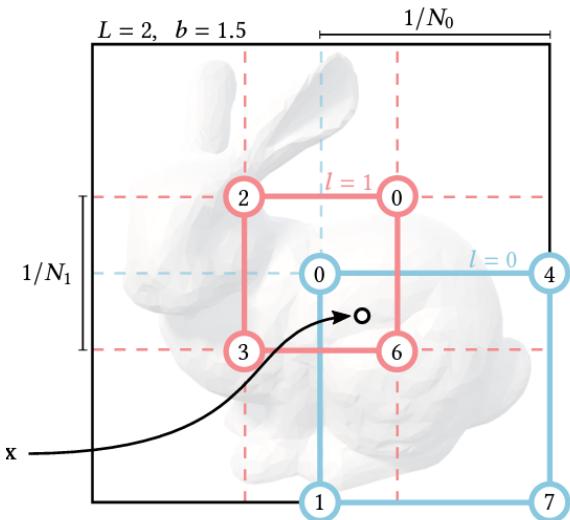
(1) Hashing of voxel vertices

Multi-resolution grid
storing learnable
feature vectors using
hashing (blue:
coarse grid; pink:
finer grid)

Hybrid: grid/neural network

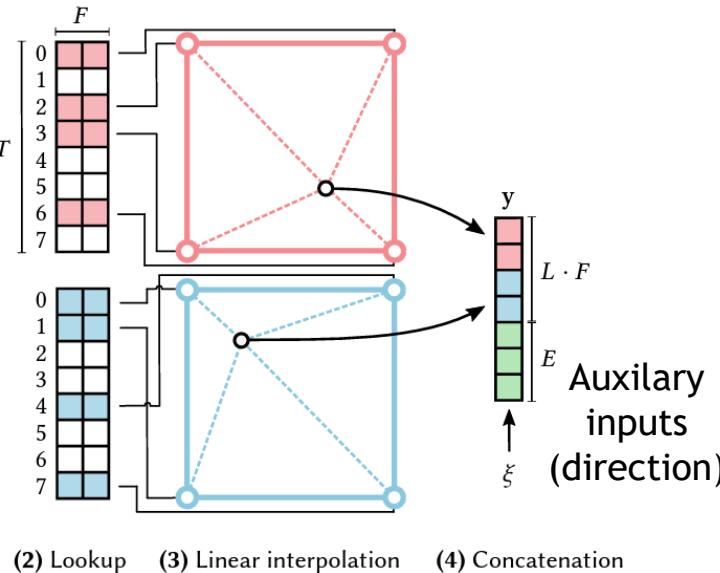
- Instant Neural Graphics Primitives (NGP)

<https://nvlabs.github.io/instant-ngp/>



(1) Hashing of voxel vertices

Multi-resolution grid
storing learnable
feature vectors using
hashing (blue:
coarse grid; pink:
finer grid)

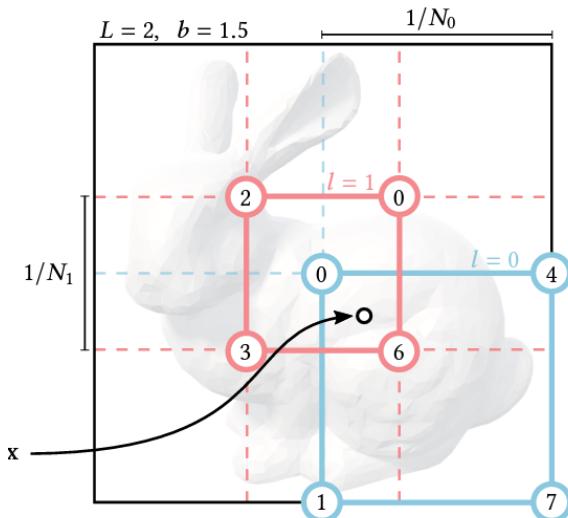


Sparse grid storage using
spatial hashing,
linear interpolation;
concatenation of feature
vectors over multiple
resolutions

Hybrid: grid/neural network

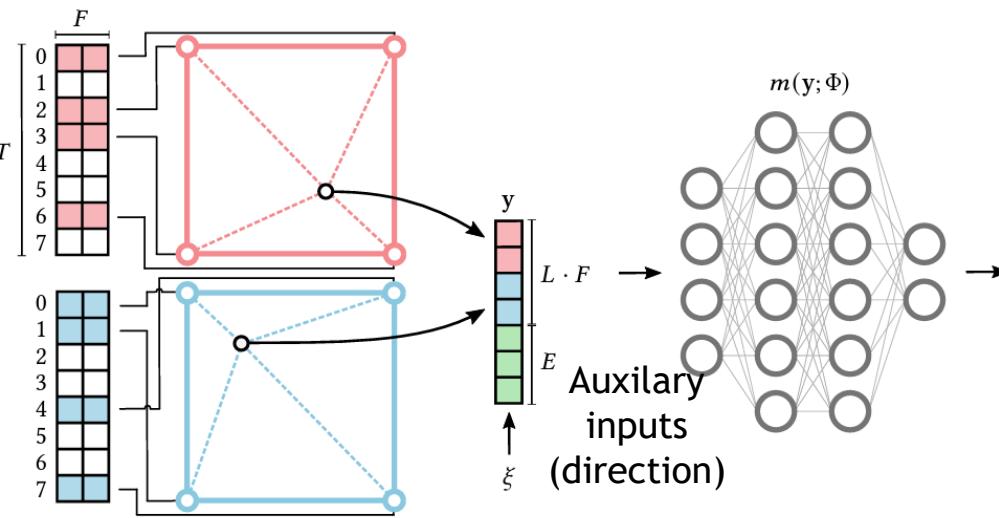
- Instant Neural Graphics Primitives (NGP)

<https://nvlabs.github.io/instant-ngp/>



(1) Hashing of voxel vertices

Multi-resolution grid storing learnable feature vectors using hashing (blue: coarse grid; pink: finer grid)



(2) Lookup

(3) Linear interpolation

(4) Concatenation

(5) Neural network

Sparse grid storage using spatial hashing, linear interpolation; concatenation of feature vectors over multiple resolutions

Small neural network to predict desired quantity (radiance, density)

Instant NGP

- Well-optimized implementation <https://nvlabs.github.io/instant-ngp/>
- Much faster than original NeRF
- No collision handling for spatial hashing, still works, see paper
- Hybrid sparse multiresolution hash grid/neural network better than spatial encoding (higher quality, faster)

	Mic	Ficus	Chair	Hotdog	Materials	Drums	Ship	Lego	avg.
Ours: Hash (1 s)	26.09	21.30	21.55	21.63	22.07	17.76	20.38	18.83	21.202
Ours: Hash (5 s)	32.60	30.35	30.77	33.42	26.60	23.84	26.38	30.13	29.261
Ours: Hash (15 s)	34.76	32.26	32.95	35.56	28.25	25.23	28.56	33.68	31.407
Ours: Hash (1 min)	35.92 ●	33.05 ●	34.34 ●	36.78	29.33	25.82 ○	30.20 ●	35.63 ●	32.635 ●
Ours: Hash (5 min)	36.22 ○	33.51 ●	35.00 ○	37.40 ○	29.78 ●	26.02 ●	31.10 ●	36.39 ●	33.176 ●
mip-NeRF (~hours)	36.51 ●	33.29 ○	35.14 ●	37.48 ●	30.71 ○	25.48 ●	30.41 ○	35.70 ○	33.090 ○
NSVF (~hours)	34.27	31.23	33.19	37.14 ●	32.68 ●	25.18	27.93	32.29	31.739
NeRF (~hours)	32.91	30.13	33.00	36.18	29.62	25.01	28.65	32.54	31.005
Ours: Frequency (5 min)	31.89	28.74	31.02	34.86	28.93	24.18	28.06	32.77	30.056
Ours: Frequency (1 min)	26.62	24.72	28.51	32.61	26.36	21.33	24.32	28.88	26.669

Image reconstruction quality, PSNR

Bottom two rows: their approach, but spatial encoding instead of hash grid

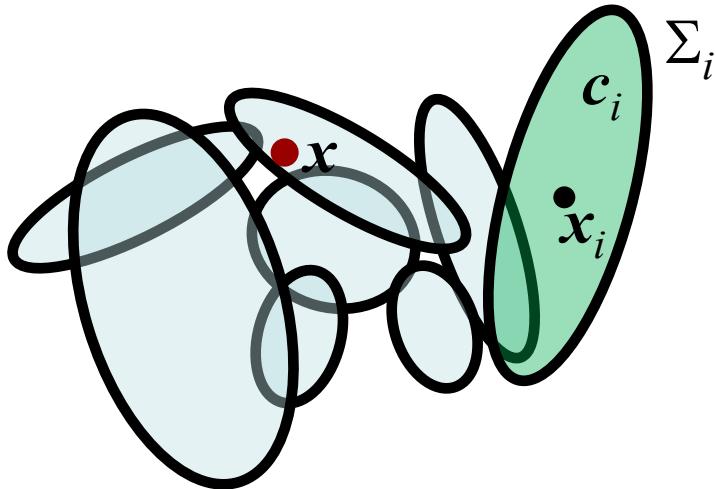
Slow training and rendering

- Training/evaluating neural network for radiance and density is slow
- Alternative representations for radiance, density that are faster to evaluate?
 1. Grid-based with interpolation function (e.g. linear)
 2. Radial (or elliptical) basis functions

Radial basis functions

- Function represented as weighted sum of simple (radially symmetric) basis functions centered at irregular locations \mathbf{x}_i https://en.wikipedia.org/wiki/Radial_basis_function
- Here: N Gaussian basis functions G , covariance Σ_i , centered at locations \mathbf{x}_i , to represent radiance $\mathbf{c}(\mathbf{x})$
- Parameters (to be optimized): coefficients c_i , covariances Σ_i , centers \mathbf{x}_i

$$\mathbf{c}(\mathbf{x}) = \sum_{i=1}^N c_i G_{\Sigma_i}(\mathbf{x} - \mathbf{x}_i)$$

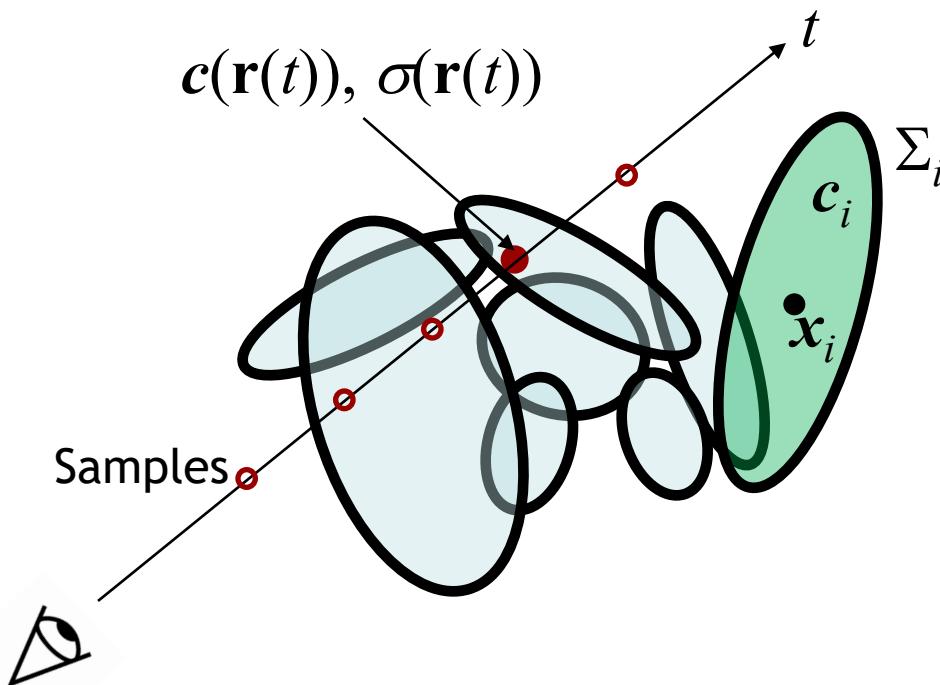


Volumetric rendering of Gaussians

- Same volume rendering model as before (NeRF)

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt \quad T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$$

- Color and density values $\mathbf{c}(\mathbf{r}(t)) = \sum_{i=1}^N \mathbf{c}_i G_{\Sigma_i}(\mathbf{r}(t) - \mathbf{x}_i)$



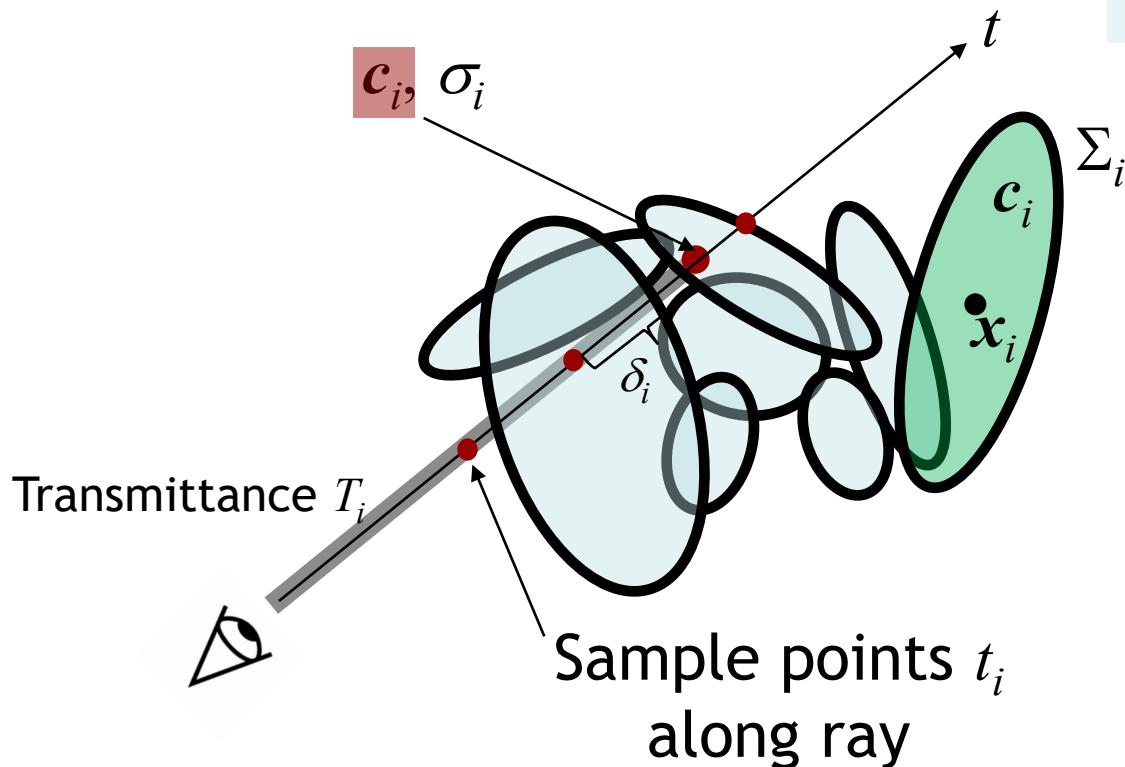
Volumetric rendering of Gaussians: ray marching

- Discretization: can do ray marching as before

$$\hat{C}(\mathbf{r}) = \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i \quad T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$
$$\delta_i = t_{i+1} - t_i \quad \text{Same as original NeRF}$$

$$\mathbf{c}_i = \sum_{i=1}^N \mathbf{c}_i G_{\Sigma_i}(\mathbf{r}(t_i) - \mathbf{x}_i)$$

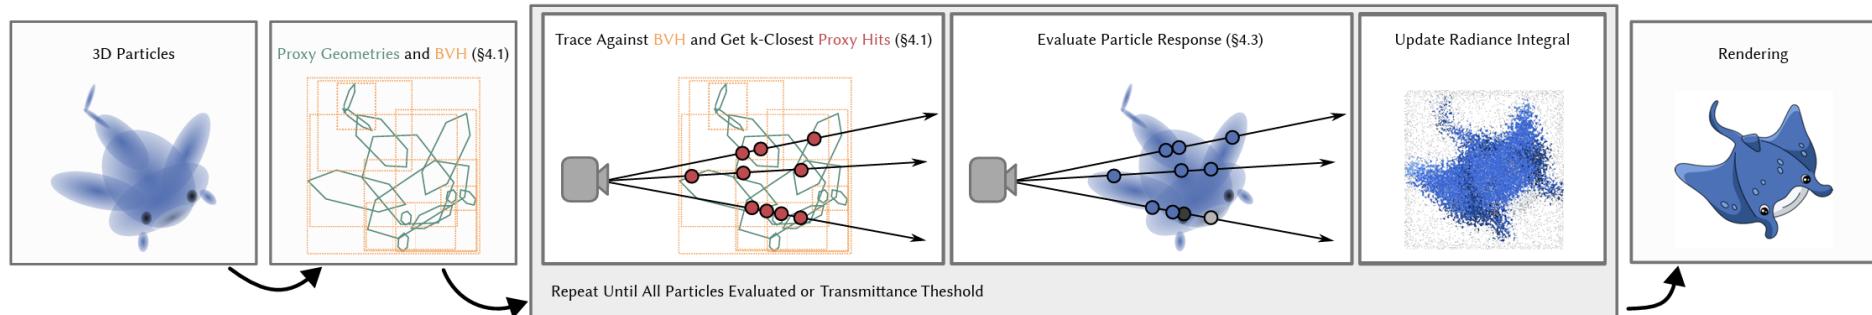
Gaussians instead of neural network, similar for density σ_i



For efficient evaluation, need spatial data structure to avoid summing over all N Gaussians at each sample point

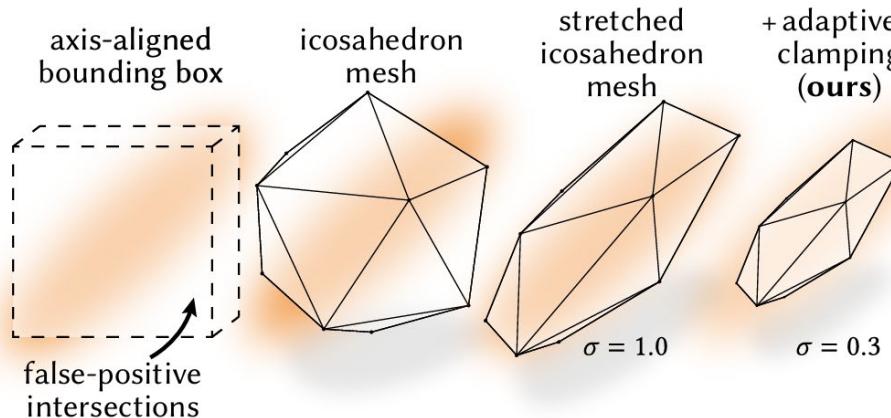
Volumetric rendering of Gaussians: ray marching

- Efficient ray marching requires spatial data structure (bounding volume hierarchy, BVH), proxy geometries for Gaussians



<https://gaussiantracer.github.io/>

- Potential proxy geometries for Gaussians



Volumetric rendering of Gaussians

- Implementation of ray tracing approach: 3D Gaussian Ray Tracing: Fast Tracing of Particle Scenes <https://gaussiantracer.github.io/>
- Advantage: light transport effects that require ray tracing (shadows, reflections, etc.)



Using ray tracing to include mirror reflection, reflection/refraction, depth-of-field blur, adding objects

Volumetric rendering of Gaussians: splatting

- 3D Gaussian splatting <https://repo-sam.inria.fr/fungraph/3d-gaussian-splatting/>
- Alternative rendering algorithm: instead of ray tracing/marching, use rasterization [Zwicker 2001]
- Loop over Gaussian in front to back order
 - Project 3D Gaussians one-by-one to 2D
 - Calculate integral along viewing rays per Gaussian, can be done analytically
 - Alpha-blend projected, integrated Gaussians based on densities
- Requires only **single loop** over all Gaussians per image (as opposed to per sample point loop over all Gaussians in naïve ray marching approach)
- Main benefit: fast rendering (“Ours”: 3D Gaussian splatting)

