

CMSC740

Advanced Computer Graphics

Matthias Zwicker
Fall 2024

Generative shape models without 3D data

- Given: database of images
- Goal: mechanism that produces random new 3D objects that, when rendered to 2D images, “look just like images from the database”
- Approach: differentiable rendering in combination with generative model for images



Training objective: generative model (GAN, diffusion model) to make output images look like images in database

GAN-based geometry synthesis without 3D data

- “Efficient Geometry-aware 3D Generative Adversarial Networks”, CVPR 2022
<https://github.com/NVlabs/eg3d>
- Contributions
 - Hybrid voxel grid-implicit 3D representation
 - GAN-based training of generative shape model without 3D supervision



Face images rendered via generative 3D model


<https://nvlabs.github.io/eg3d/>

Hybrid representation

- Goal: predict density, radiance at 3D locations as in NeRF

Original NeRF

“implicit” network (PE:
positional encoding, FC: fully
connected layer)




(a) NeRF (Implicit)

Hybrid representation

- Goal: predict density, radiance at 3D locations as in NeRF


Original NeRF

“implicit” network (PE:
positional encoding, FC: fully
connected layer)



(a) NeRF (Implicit)

3D grid with learned features




(b) Voxels (Explicit or Hybrid)

Hybrid representation


- Goal: predict density, radiance at 3D locations as in NeRF
- Benefits of hybrid voxel grid-implicit representation
 - Can use small network, faster training
 - Avoid storage overhead of full 3D grid, enables higher resolutions

Original NeRF


“implicit” network (PE:
positional encoding, FC: fully
connected layer)



3D grid with learned features



Tri-planes with learned feature grids



Hybrid approach:
Project 3D location
onto tri-planes; look up
and average three
feature vectors; use as
input for small neural
network to predict
density, color

(a) NeRF (Implicit)

(b) Voxels (Explicit or Hybrid)

(c) Ours (Hybrid)

Comparison on scene reconstruction

- Comparison using same scene reconstruction problem as in original NeRF



	MLP	Rel. Speed ↑	Rel. Mem. ↓
Mip-NeRF [2]	8×256	$1\times$	$1\times$
Voxels (hybrid)	4×128	$3.5\times$	$0.33\times$
Tri-plane (SSO)	4×128	$2.9\times$	$0.32\times$

Approx. 3x speedup and memory savings compared to original NeRF
MLP: nr of layers x nr of neurons per layer

Scene reconstruction on Tanks & Temples dataset


<https://www.tanksandtemples.org/>

3D GAN with differentiable rendering

3D shape generator

input: noise (latent space)


output: tri-plane features



3D GAN with differentiable rendering


3D shape generator
input: noise (latent space)
output: tri-plane features

Differentiable NeRF
rendering




3D GAN with differentiable rendering

3D shape generator
input: noise (latent space)
output: tri-plane features




Differentiable NeRF
rendering

Image
superresolution



3D GAN with differentiable rendering



- Image-based GAN discriminator, no 3D training data necessary
 - Camera parameters for real images estimated with standard methods
 - Camera parameters used as input to conditional generator and discriminator, and for volume rendering


Stylegan 2 generator

- Well-engineered architecture for GAN generators and discriminators to avoid empirically observed artifacts in previous methods

<https://github.com/NVlabs/stylegan2>

Basic convolutional generator

Noise (latent space)




Stylegan 2 generator


- Well-engineered architecture for GAN generators and discriminators to avoid empirically observed artifacts in previous methods

<https://github.com/NVlabs/stylegan2>

Basic convolutional generator



Stylegan 2 generator



3D shape generator

- Stylegan 2 architecture (“works well”)

- Learned convolution weights w_{ijk} (input channel i , output channel j , spatial location in convolution kernel k); learned biases b

- **Latent noise** vector (512-dim.) fed to convolutional layers via mapping network A

- Provides scaling factors s_i for input feature channel i


- **Modulation**

$$w'_{ijk} = s_i \cdot w_{ijk}$$


- **Demodulation** (small constant ϵ to avoid division by zero)

$$w''_{ijk} = w'_{ijk} / \sqrt{\sum_{i,k} {w'_{ijk}}^2 + \epsilon}$$

- 3D GAN: 256x256x96 output split into tri-planes at 256x256x32




Differentiable NeRF rendering




- Produces 32 feature channels, 3 interpreted as RGB colors and fed to discriminator

Superresolution network



- Upsampling from 128x128x32 to 512x512x3
- Using Stylegan 2 convolutional layers, same mapping network for modulation

Discriminator



- Convolutional layers
- “Dual discriminator” using final image I^+_{RGB} and upscaled 128x128 image, both at 512x512x3, concatenated to 512x512x6
- Conditional discriminator using camera parameters

Results



Trained with FFHQ and AFHQv2 Cats

<https://github.com/NVlabs/ffhq-dataset>

<https://www.v7labs.com/open-datasets/afhq>

Comparisons

GIRAFFE



Lift. SG



π -GAN



Ours



GIRAFFE, CVPR 2021 best paper <https://m-niemeyer.github.io/project-pages/giraffe/index.html>

π -GAN, CVPR 2021 <https://marcoamonteiro.github.io/pi-GAN-website/>

Lifting Stylegan <https://github.com/seasonSH/LiftedGAN>

Quantitative evaluation

- Challenge: quantitative evaluation of generative models requires comparison of generated and true densities
 - Difficult because of high dimensionality of data (images, video, etc.)
- Idea
 - Estimate densities in lower dimensional space that is relevant to human perception
 - Compare densities by fitting simple model (Gaussian distribution) to lower dimensional data points
- FID: Fréchet Inception Distance (FID)
https://en.wikipedia.org/wiki/Fr%C3%A9chet_inception_distance
 - Use deepest feature layer of pre-trained Inception V3 network as lower dimensional space (2048 dimensions)
 - Fit Gaussians (mean μ , variance Σ) to Inception V3 feature vectors of true, generated data
 - Compare Gaussians of true and generated data using Fréchet distance d_F

$$d_F(\mathcal{N}(\mu, \Sigma), \mathcal{N}(\mu', \Sigma'))^2 = \|\mu - \mu'\|_2^2 + \text{tr} \left(\Sigma + \Sigma' - 2 \left(\Sigma^{\frac{1}{2}} \cdot \Sigma' \cdot \Sigma^{\frac{1}{2}} \right)^{\frac{1}{2}} \right)$$

Comparisons

- FID: Fréchet Inception Distance using 50k generated images, all real data
- ID: consistency of face identity over different viewpoints using Arcface cosine similarity score
<https://arxiv.org/abs/1801.07698>
- Depth/pose: consistency of NeRF geometry/poses with “pseudo ground truth” geometry/poses reconstructed from rendered multiview images


	FFHQ				Cats
	FID↓	ID↑	Depth↓	Pose↓	FID↓
GIRAFFE 256 ²	31.5	0.64	0.94	.089	16.1
π -GAN 128 ²	29.9	0.67	0.44	.021	16.0
Lift. SG 256 ²	29.8	0.58	0.40	.023	—
Ours 256 ²	4.8	0.76	0.31	.005	3.88
Ours 512 ²	4.7	0.77	0.39	.005	2.77[†]

Text-based 3D synthesis

- “DreamFusion: Text-to-3D using 2D Diffusion”,
<https://dreamfusion3d.github.io/>
- Goal: conditional generation of 3D objects based on text inputs
 - “Create random 3D objects that are consistent with given text description”
- Challenge
 - Requires large collection of text-3D object pairs, which don’t exist, to train conditional generative model
- Approach: leverage existing, diffusion-based text-to-image conditional generative models in combination with NeRF-based differentiable rendering


Diffusion-based text-to-image

- Training: given text-image pairs, train U-net conditioned on text embedding trained to predict (synthetic) noise added to image
 - HereImagen model <https://Imagen.research.google/> <https://arxiv.org/pdf/2205.11487.pdf>
 - Transformer-based text embedding from pre-trained large language model (here T5-XXL, see also <https://github.com/google-research/t5x>)




U-net architecture

<https://en.wikipedia.org/wiki/U-Net>




Diffusion-based text-to-image

- Sampling (image synthesis): in each step, subtract (noise prediction – noise) from previous image (equivalent to subtracting noise prediction from (previous image + noise), see pseudocode from last time); add new noise; iterate




Diffusion-based text-to-image

- Sampling (image synthesis): in each step, subtract (noise prediction – noise) from previous image (equivalent to subtracting noise prediction from (previous image + noise), see pseudocode from last time); add new noise; iterate




Diffusion-based text-to-image

- Sampling (image synthesis): in each step, subtract (noise prediction – noise) from previous image (equivalent to subtracting noise prediction from (previous image + noise), see pseudocode from last time); add new noise; iterate




3D synthesis

- Image rendered using NeRF, starting from random initialization



3D synthesis

- Image rendered using NeRF, starting from random initialization
- Update to rendered image computed using pretrained diffusion model



3D synthesis

- Image rendered using NeRF, starting from random initialization
- Update to rendered image computed using pretrained diffusion model
- Desired update backpropagated to NeRF parameters




Image update given by
noise prediction - noise

Note

- NeRF parameters optimized for each 3D object that is generated (“sampling the generative model via optimization”)
 - Random viewpoints for rendering; text prompts augmented with “front”, “side”, “back” based on viewpoint location
 - 1.5 hours (15,000 iterations) per scene using TPUv4 machine with 4 chips; each chip rendering separate view and evaluating diffusion U-Net with per-device batch size of 1
 - More theory in DreamFusion paper <https://arxiv.org/abs/2209.14988>
- Text-to-image diffusion model is frozen during 3D generation; not involved in gradient computation for NeRF parameters
 - DreamFusion uses ImaGen model <https://Imagen.research.google/>

NeRF rendering


- Predict diffuse BRDF (albedo), instead of radiance
- Compute surface normals using gradient of density field
- Shading using randomly positioned point light, ambient light
 - Using shading model helps reconstruct geometry (empirically shown in ablation studies)
- NeRF only evaluated within bounding box; background color outside bounding box predicted using separate MLP

Evaluation

- “R-Precision”: accuracy with which CLIP retrieves correct image caption among a set of distractors given scene rendering
 - CLIP (“contrastive Language-Image Pre-training”) trained to predict which images were paired with which texts in large image/text training dataset <https://openai.com/blog/clip/>

Method	R-Precision ↑					
	CLIP B/32		CLIP B/16		CLIP L/14	
	Color	Geo	Color	Geo	Color	Geo
GT Images	77.1	–	79.1	–	–	–
Dream Fields	68.3	–	74.2	–	–	–
(reimpl.)	78.6	1.3	(99.9)	(0.8)	82.9	1.4
CLIP-Mesh	67.8	–	75.8	–	74.5 [†]	–
DreamFusion	75.1	42.5	77.5	46.6	79.7	58.5

R-precision using different CLIP models



Follow up work

- Better ways to backpropagate image update to rendering model (NeRF, or something else depending on application)
- Variational score distillation (NeurIPS 2023)
<https://github.com/thu-ml/prolificdreamer>
- Classifier score distillation (ICLR 2024)
<https://github.com/CVMI-Lab/Classifier-Score-Distillation>
- Other applications

- Scene texturing (CVPR 2024)

<https://daveredrum.github.io/SceneTex/>